
Dan's Cheat Sheets Documentation

Release 1

Dan Poirier

Dec 18, 2018

Contents

1	Android information	3
2	Ansible	7
3	AWS	33
4	Bootstrap	35
5	Debian	37
6	Diet	39
7	Django	45
8	Docker	101
9	Elasticsearch	103
10	Elixir	109
11	Git	115
12	Google APIs	121
13	Gulp Tasks	123
14	HAProxy	125
15	Haskell	127
16	i3	131
17	IPv6	133
18	Javascript	135
19	Kobo ereader	143
20	Linux Notes, Misc.	145

21 LIRC: linux infrared remote control	147
22 LVM	149
23 LXDE	151
24 Logitech Harmony	153
25 MPD	175
26 MySQL with Django	177
27 Nginx	181
28 NPM	185
29 OpenSSL	187
30 Org mode (Emacs)	193
31 Postfix	195
32 Postgres	197
33 Python	203
34 Raspberry Pi	225
35 Release - how to	229
36 reStructuredText	231
37 Salt Stack	237
38 Socat	239
39 Tmux	241
40 Travis CI	243
41 Video	245
42 X11 Window System	247
43 Virtualbox	249
44 Vue	251
45 YAML	255
46 Indices and tables	257

Contents:

Contents:

1.1 Nexus 6P

This information is specific to the Google/Huawei Nexus 6P. See also [more universal android info](#).

Code name: angler

There's lots of information about the Nexus 6P that is nicely gathered in [this thread](#) on xda-developers.

SEE ALSO: Android universal instructions

1.1.1 Booting into different modes

Enter safe mode

(this will boot with 3rd-party apps disabled):

1. Have Your Nexus 6P Powered on and at the Home Screen
2. Press and Hold the Power Button until You See the 'Power Off' Dialog Screen Pop Up,
3. Let Go of the Power Button
4. Then Tap and Hold Your Finger on the "Power Off" Option. After Holding for a Few Seconds You'll be Asked if You Want to Reboot Into Safe Mode.
5. Simply Tap the 'OK' Option to Reboot Your Nexus 6P Into Safe Mode. Then Wait for the Nexus 6P to Reboot

Enter fastboot aka bootloader mode

(<https://www.androidexplained.com/nexus-6p-fastboot-mode/>):

This lets you execute certain ADB and Fastboot commands from your computer, get to recovery mode from here.

1. Power Down the Nexus 6P
2. Once the Device is Off, Press and Hold the Power Button and Volume Down Button at the Same Time
3. Continue Holding These Two Buttons Until you See the Special Fastboot Mode Screen
4. When You are in Fastboot Mode, You Can Let Go of These Two Buttons

Enter recovery mode

This lets you do a factory reset, wipe the cache partition or sideload an OTA update, etc. Also, “reboot bootloader” goes back to fastboot/bootloader mode.

1. Boot the Nexus 6P into Fastboot Mode
2. Once in Fastboot Mode, Press the volume Down Button Twice. This Should Highlight the ‘Recovery’ Option
3. Press the Power Button to Select This Option. This Will Take You to a Black Screen with a Green Android Laying Down
4. Press and Hold the the Power Button, Then Press the Volume Up Button, You’ll Immediately be Taken to the Recovery Mode Menu

Here’s a kernel? <https://forum.xda-developers.com/nexus-6p/development/kernel-purez-kernel-v2-0-purezandroid-t3636909>
<https://github.com/purezandroid/purez-kernel-angler>

1.1.2 Installing/replacing stuff

Custom recovery

This is a pre-req to install custom ROM.

The following information is not specific to any Android device. See the links just above for device-specific info.

Android universal instructions (custom roms etc)

See device-specific pages for how to get into recovery mode, fastboot mode, etc etc

1.2 Installing adb and fastboot on a computer

(not the android, a linux or windows or something)

<https://developer.android.com/studio/releases/platform-tools.html>

Just download the zip, unpack, and put the binaries on your PATH or something.

1.3 Custom recovery

(Pre-req to install custom ROM)

<https://www.xda-developers.com/how-to-install-twrp/>

Download from <https://twrp.me/Devices/>

Before:

- Install adb and fastboot on a computer
- Enable USB debugging on the android device

To install TWRP:

1. Boot into bootloader mode (e.g. fastboot mode)
2. “fastboot flash recovery twrp-2.8.x.x-xxx.img”
3. “fastboot reboot”

ALTERNATIVELY if device is already rooted, you can [install and use TWRP manager](#)

1.4 Custom ROM

<http://www.android.gs/install-custom-roms-android-devices-guide/>

THIS WIPES ALL YOUR DATA

1. Copy zipfiles to phone file system
2. Boot into recovery mode
3. select “wipe data factory reset”, “wipe cache partition” and “wipe dalvick cache”.
4. Choose “install zip from SD card” and “choose zip from SD card”. Pick the update file and flash the same.
5. Optional: repeat this operation for applying the Google Apps package.
6. Reboot

1.5 Kernel

(as opposed to the whole ROM)

1.5.1 With fastboot

1. Plug in the phone and boot to fastboot mode (back+power or camera+power). Wait until the screen with skating ‘droids appears,
2. press the back button (the center bar should say “Fastboot” or “Fastboot USB”)
3. On the computer:

```
fastboot flash boot boot.img
fastboot reboot
```


This is growing into a minimal Ansible reference of sorts, since Ansible’s own docs have nothing like a reference.

- [Ansible](#).
- [list of keys that common playbook objects can take](#).
- [Release tarballs](#)
- [Ansible documentation for older releases](#)

Quickies:

To check the ubuntu version, `ansible_distribution_version|float < 18` (ansible_distribution_version is e.g. “16.04”)

2.1 Running Ansible tasks in the background

Example:

```
- name: start collectstatic in the background
  command: "{{ install_root }}/env/bin/python manage.py collectstatic --noinput -v 0"
  args:
    chdir: "{{ install_root }}/webapp"
  async: 1000
  poll: 0
  register: collectstatic_bg

#####
#
# PUT TASKS HERE THAT DON'T NEED TO BE RUN BEFORE COLLECTSTATIC CAN START,
# AND THAT WON'T AFFECT THE BACKGROUND COLLECTSTATIC.
#
#####
- name: clean up local tarball
```

(continues on next page)

(continued from previous page)

```

delegate_to: 127.0.0.1    # Run on localhost
run_once: yes             # only once
become: no                # Don't need sudo
file:
  state: absent
  path: "{{ tarball }}"

- name: migrate
  command: "{{ install_root }}/env/bin/python manage.py migrate --noinput"
  args:
    chdir: "{{ install_root }}/webapp"

- name: install tasks
  command: "{{ install_root }}/env/bin/python manage.py installtasks --traceback"
  args:
    chdir: "{{ install_root }}/webapp"

#####
#
# Check every 'delay' seconds, up to 'retries' times, until collectstatic is done
#
#####
- name: wait for collectstatic to finish
  async_status: jid={{ collectstatic_bg.ansible_job_id }}
  register: job_result
  until: job_result.finished
  retries: 80
  delay: 15

#####
#
# PUT TASKS AFTER THIS THAT CAN'T RUN UNTIL COLLECTSTATIC IS DONE
#
#####

```

2.2 Blocks

- [Blocks doc](#)
- [A blog post about blocks](#)
- [Blog post with examples](#)
- [Complete list of possible keywords](#)

Blocks can be used anywhere a task can (mostly?). They allow applying task keys to a group of tasks without having to repeat them over and over. They also provide a form of error handling.

Syntax:

```

block:
  - <task1>
  - <task2>
when: <condition>
become: true
become_user: <username>

```

(continues on next page)

(continued from previous page)

```
....
[rescue:
  - debug: msg="This task runs if there's an error in the block"
  - <task2>
  ...
always:
  - debug: msg="I always run"
  ... more tasks ..
]
```

2.3 Conditionals, return values, and registered variables

Ansible conditionals doc

2.3.1 Conditional tasks

See *Task*.

2.4 Return Values and registered variables

To create a variable with results from a task, use `register`:

```
- name: any task
  command: whatever
  register: varname
```

Then in later tasks, you can use `varname` in conditionals like `when`.

The variable is actually an object with lots of useful items in it. Some of them:

- `.changed` - boolean, whether the task changed anything
- `.failed` - boolean, true if the task failed
- `.skipped` - boolean, true if the task was skipped
- `.result` - ? (depends on the task?)
- `.results` - If this key exists, it indicates that a loop was present for the task and that it contains a list of the normal module 'result' per item.

more on common return values.

There are also useful filters:

- `is succeeded` - boolean, true if task succeeded
- `is failed` - boolean, true if task failed
- `is skipped` = boolean, true if the task was skipped

`succeeded` is probably the most useful here - the others just duplicate attributes.

2.5 Configuration

2.5.1 Configuration file

Syntax .ini file

See [The Ansible Configuration File doc](#).

Ansible uses the first config file it finds on this list:

- ANSIBLE_CONFIG (an environment variable)
- ansible.cfg (in the current directory)
- .ansible.cfg (in the home directory)
- /etc/ansible/ansible.cfg

Some useful vars in the [defaults] section:

any_errors_fatal

If true, stop immediately if any task fails. Default value of False only stops for the host that failed. The playbook invocation will eventually report failure, but the error itself might be thousands of lines back in the output. Recommend changing this to True.

display_args_to_stdout

If true, more information displayed as tasks execute. Default: False.

error_on_undefined_vars

If true, task fails if any undefined vars are encountered, which seems like it ought to be the default behavior, but it's not.

hostfile

This is the default location of the inventory file, script, or directory that Ansible will use to determine what hosts it has available to talk to:

hostfile = /etc/ansible/hosts

use_persistent_connections

Default False. Whether to use persistent connections. (Yes, this is in the *defaults* section.)

private_role_vars

Makes role variables inaccessible from other roles. This was introduced as a way to reset role variables to default values if a role is used more than once in a playbook. Default: False

retry_files_enabled

Default true. If True, create .retry files on failure. These are generally useless so I change this to False to not clutter up my file system with them.

roles_path

The roles path indicate additional directories beyond the 'roles/' subdirectory of a playbook project to search to find Ansible roles. For instance, if there was a source control repository of common roles and a different repository of playbooks, you might choose to establish a convention to checkout roles in /opt/mysite/roles like so:

```
roles_path = /opt/mysite/roles
```

Additional paths can be provided separated by colon characters, in the same way as other pathstrings:

```
roles_path = /opt/mysite/roles:/opt/othersite/roles
```

Roles will be first searched for in the playbook directory. Should a role not be found, it will indicate all the possible paths that were searched.

Some useful vars in the [inventory] section:

any_unparsed_is_failed

Default false. If 'true', it is a fatal error when any given inventory source cannot be successfully parsed by any available inventory plugin; otherwise, this situation only attracts a warning.

Some useful vars in the [ssh_connection] section:

pipelining

Pipelining, if supported by the connection plugin, reduces the number of network operations required to execute a module on the remote server, by executing many Ansible modules without actual file transfer. This can result in a very significant performance improvement when enabled. However this conflicts with privilege escalation (become). For example, when using 'sudo:' operations you must first disable 'requiretty' in /etc/sudoers on all managed hosts, which is why it is disabled by default.

scp_if_ssh

Preferred method to use when transferring files over ssh. When set to smart, Ansible will try them until one succeeds or they all fail. If set to True, it will force 'scp', if False it will use 'sftp'.

ssh_args

If set, this will override the Ansible default ssh arguments. In particular, users may wish to raise the ControlPersist time to encourage performance. A value of 30 minutes may be appropriate. Be aware that if -o ControlPath is set in ssh_args, the control path setting is not used.

Warning: If you set this, the default setting is completely overridden, so you should include it (possibly edited):
`-C -o ControlMaster=auto -o ControlPersist=60s`

Example:

```
ssh_args = -C -o ControlMaster=auto -o ControlPersist=300s -o ForwardAgent=yes -o   
ControlPath=./ansible_ssh_conn_%h
```

2.6 Host Patterns

See [ansible host patterns doc](#) for host patterns.

<hosts>:

“all” = all hosts in inventory file “*” = all hosts in inventory file “~<regex>” = use regex syntax for this pattern “<pattern1>:<pattern2>” = include all hosts that match pattern1 OR pattern2 “<pattern1>:&<pattern2>” = include all hosts that match pattern1 AND pattern2 “<pattern1>:!<pattern2>” = include all hosts that match pattern1 but NOT pattern2

2.7 Inventory

2.7.1 Inventory directory

Whatever directory the *Inventory file* is in.

2.7.2 Inventory file

Default /etc/ansible/hosts

Change set ANSIBLE_HOSTS in environment

```
ansible-playbook -i <inventoryfile> ...
```

set *hostfile* in configuration

Syntax .ini file, except initial lines don't need to be in a section

The inventory file is basically a list of hostnames or IP addresses, one per line. Can include port with hostname:port or address:port.

Ranges: Including [m:n] in a line will repeat the line for every value from m through n. m and n can be numbers or letters:

```
[mygroup]
host[1:25]
```

Host *Pre-defined variables*: Can specify per-host options after hostname on the same line. E.g.:

```
jumper ansible_ssh_port=5555 ansible_ssh_host=192.168.1.50
```

See also *Variables files*.

Group *Pre-defined variables*: add [groupname:vars] section and put var definitions in it, one per line. Example:


```
[all:vars]
project_name="PeterPan"
environment_name=staging
```

See also *Variables files*.

Groups of groups: add [newgroupname:children] and put other group names in it, one per line:

```
[group3]
host13
host14

[group3:children]
group1
group2

[group3:vars]
group3_var1=27
group3_var2="Hello, World"
```

2.8 Invoking

2.8.1 Ad-hoc

To run an ad-hoc command, use *ansible*. (But you almost always will want to run a playbook; see below.)

Examples of ad-hoc commands:

```
$ ansible all -m ping
# as bruce
$ ansible all -m ping -u bruce
# as bruce, sudoing to root
$ ansible all -m ping -u bruce --sudo
# as bruce, sudoing to batman
$ ansible all -m ping -u bruce --sdo --sudo-user batman
$ ansible all -a "/bin/echo hello"
```

Help:

Usage: ansible <host-pattern> [options]

Options:

```
-a MODULE_ARGS, --args=MODULE_ARGS
                        module arguments
-k, --ask-pass         ask for SSH password
--ask-su-pass          ask for su password
-K, --ask-sudo-pass    ask for sudo password
--ask-vault-pass       ask for vault password
-B SECONDS, --background=SECONDS
                        run asynchronously, failing after X seconds
                        (default=N/A)
-C, --check            don't make any changes; instead, try to predict some
                        of the changes that may occur
-c CONNECTION, --connection=CONNECTION
                        connection type to use (default=smart)
```

(continues on next page)

(continued from previous page)

```
-f FORKS, --forks=FORKS
                        specify number of parallel processes to use
                        (default=5)
-h, --help              show this help message and exit
-i INVENTORY, --inventory-file=INVENTORY
                        specify inventory host file
                        (default=/etc/ansible/hosts)
-l SUBSET, --limit=SUBSET
                        further limit selected hosts to an additional pattern
--list-hosts            outputs a list of matching hosts; does not execute
                        anything else
-m MODULE_NAME, --module-name=MODULE_NAME
                        module name to execute (default=command)
-M MODULE_PATH, --module-path=MODULE_PATH
                        specify path(s) to module library
                        (default=/usr/share/ansible/)
-o, --one-line          condense output
-P POLL_INTERVAL, --poll=POLL_INTERVAL
                        set the poll interval if using -B (default=15)
--private-key=PRIVATE_KEY_FILE
                        use this file to authenticate the connection
-S, --su                run operations with su
-R SU_USER, --su-user=SU_USER
                        run operations with su as this user (default=root)
-s, --sudo              run operations with sudo (nopasswd)
-U SUDO_USER, --sudo-user=SUDO_USER
                        desired sudo user (default=root)
-T TIMEOUT, --timeout=TIMEOUT
                        override the SSH timeout in seconds (default=10)
-t TREE, --tree=TREE    log output to this directory
-u REMOTE_USER, --user=REMOTE_USER
                        connect as this user (default=poirier)
--vault-password-file=VAULT_PASSWORD_FILE
                        vault password file
-v, --verbose           verbose mode (-vvv for more, -vvvv to enable
                        connection debugging)
--version               show program's version number and exit
```

2.8.2 Playbooks

To run a playbook, use `ansible-playbook`. Here's the help from 2.0.1.0:

Usage: `ansible-playbook playbook.yml`

Options:

```
--ask-become-pass      ask for privilege escalation password
-k, --ask-pass         ask for connection password
--ask-su-pass          ask for su password (deprecated, use become)
-K, --ask-sudo-pass    ask for sudo password (deprecated, use become)
--ask-vault-pass       ask for vault password
-b, --become           run operations with become (nopasswd implied)
--become-method=BECOME_METHOD
                        privilege escalation method to use (default=sudo),
                        valid choices: [ sudo | su | pbrun | pfexec | runas |
                        doas ]
```

(continues on next page)

(continued from previous page)

```

--become-user=BECOME_USER
    run operations as this user (default=root)
-C, --check
    don't make any changes; instead, try to predict some
    of the changes that may occur
-c CONNECTION, --connection=CONNECTION
    connection type to use (default=smart)
-D, --diff
    when changing (small) files and templates, show the
    differences in those files; works great with --check
-e EXTRA_VARS, --extra-vars=EXTRA_VARS
    set additional variables as key=value or YAML/JSON
--flush-cache
    clear the fact cache
--force-handlers
    run handlers even if a task fails
-f FORKS, --forks=FORKS
    specify number of parallel processes to use
    (default=5)
-h, --help
    show this help message and exit
-i INVENTORY, --inventory-file=INVENTORY
    specify inventory host path
    (default=/etc/ansible/hosts) or comma separated host
    list.
-l SUBSET, --limit=SUBSET
    further limit selected hosts to an additional pattern
--list-hosts
    outputs a list of matching hosts; does not execute
    anything else
--list-tags
    list all available tags
--list-tasks
    list all tasks that would be executed
-M MODULE_PATH, --module-path=MODULE_PATH
    specify path(s) to module library (default=None)
--new-vault-password-file=NEW_VAULT_PASSWORD_FILE
    new vault password file for rekey
--output=OUTPUT_FILE
    output file name for encrypt or decrypt; use - for
    stdout
--private-key=PRIVATE_KEY_FILE, --key-file=PRIVATE_KEY_FILE
    use this file to authenticate the connection
--scp-extra-args=SCP_EXTRA_ARGS
    specify extra arguments to pass to scp only (e.g. -l)
--sftp-extra-args=SFTP_EXTRA_ARGS
    specify extra arguments to pass to sftp only (e.g. -f,
    -l)
--skip-tags=SKIP_TAGS
    only run plays and tasks whose tags do not match these
    values
--ssh-common-args=SSH_COMMON_ARGS
    specify common arguments to pass to sftp/scp/ssh (e.g.
    ProxyCommand)
--ssh-extra-args=SSH_EXTRA_ARGS
    specify extra arguments to pass to ssh only (e.g. -R)
--start-at-task=START_AT_TASK
    start the playbook at the task matching this name
--step
    one-step-at-a-time: confirm each task before running
-S, --su
    run operations with su (deprecated, use become)
-R SU_USER, --su-user=SU_USER
    run operations with su as this user (default=root)
    (deprecated, use become)
-s, --sudo
    run operations with sudo (nopasswd) (deprecated, use
    become)
-U SUDO_USER, --sudo-user=SUDO_USER

```

(continues on next page)

(continued from previous page)

```

                                desired sudo user (default=root) (deprecated, use
                                become)
--syntax-check                  perform a syntax check on the playbook, but do not
                                execute it
-t TAGS, --tags=TAGS           only run plays and tasks tagged with these values
-T TIMEOUT, --timeout=TIMEOUT  override the connection timeout in seconds
                                (default=10)
-u REMOTE_USER, --user=REMOTE_USER
                                connect as this user (default=None)
--vault-password-file=VAULT_PASSWORD_FILE
                                vault password file
-v, --verbose                  verbose mode (-vvv for more, -vvvv to enable
                                connection debugging)
--version                      show program's version number and exit

```

2.8.3 Hosts pulling config

Ansible-pull ([ansible-pull doc](#)) is a small script that will checkout a repo of configuration instructions from git, and then run ansible-playbook against that content.

Assuming you load balance your checkout location, ansible-pull scales essentially infinitely.

Help from ansible-pull 2.0.1.0:

```

Usage: ansible-pull -U <repository> [options]

Options:
--accept-host-key              adds the hostkey for the repo url if not already added
--ask-become-pass             ask for privilege escalation password
-k, --ask-pass                ask for connection password
--ask-su-pass                 ask for su password (deprecated, use become)
-K, --ask-sudo-pass           ask for sudo password (deprecated, use become)
--ask-vault-pass              ask for vault password
-C CHECKOUT, --checkout=CHECKOUT
                                branch/tag/commit to checkout. Defaults to behavior
                                of repository module.
-c CONNECTION, --connection=CONNECTION
                                connection type to use (default=smart)
-d DEST, --directory=DEST     directory to checkout repository to
-e EXTRA_VARS, --extra-vars=EXTRA_VARS
                                set additional variables as key=value or YAML/JSON
-f, --force                   run the playbook even if the repository could not be
                                updated
--full                        Do a full clone, instead of a shallow one.
-h, --help                    show this help message and exit
-i INVENTORY, --inventory-file=INVENTORY
                                specify inventory host path
                                (default=/etc/ansible/hosts) or comma separated host
                                list.
-l SUBSET, --limit=SUBSET     further limit selected hosts to an additional pattern
--list-hosts                  outputs a list of matching hosts; does not execute
                                anything else
-m MODULE_NAME, --module-name=MODULE_NAME

```

(continues on next page)

(continued from previous page)

```

Repository module name, which ansible will use to
check out the repo. Default is git.
-M MODULE_PATH, --module-path=MODULE_PATH
    specify path(s) to module library (default=None)
--new-vault-password-file=NEW_VAULT_PASSWORD_FILE
    new vault password file for rekey
-o, --only-if-changed
    only run the playbook if the repository has been
    updated
--output=OUTPUT_FILE
    output file name for encrypt or decrypt; use - for
    stdout
--private-key=PRIVATE_KEY_FILE, --key-file=PRIVATE_KEY_FILE
    use this file to authenticate the connection
--purge
    purge checkout after playbook run
--scp-extra-args=SCP_EXTRA_ARGS
    specify extra arguments to pass to scp only (e.g. -l)
--sftp-extra-args=SFTP_EXTRA_ARGS
    specify extra arguments to pass to sftp only (e.g. -f,
    -l)
--skip-tags=SKIP_TAGS
    only run plays and tasks whose tags do not match these
    values
-s SLEEP, --sleep=SLEEP
    sleep for random interval (between 0 and n number of
    seconds) before starting. This is a useful way to
    disperse git requests
--ssh-common-args=SSH_COMMON_ARGS
    specify common arguments to pass to sftp/scp/ssh (e.g.
    ProxyCommand)
--ssh-extra-args=SSH_EXTRA_ARGS
    specify extra arguments to pass to ssh only (e.g. -R)
-t TAGS, --tags=TAGS
    only run plays and tasks tagged with these values
-T TIMEOUT, --timeout=TIMEOUT
    override the connection timeout in seconds
    (default=10)
-U URL, --url=URL
    URL of the playbook repository
-u REMOTE_USER, --user=REMOTE_USER
    connect as this user (default=None)
--vault-password-file=VAULT_PASSWORD_FILE
    vault password file
-v, --verbose
    verbose mode (-vvv for more, -vvvv to enable
    connection debugging)
--verify-commit
    verify GPG signature of checked out commit, if it
    fails abort running the playbook. This needs the
    corresponding VCS module to support such an operation
--version
    show program's version number and exit

```

2.9 Loops

See http://docs.ansible.com/playbooks_loops.html

2.9.1 Iterating with nested loops

Write a task:

```
- module: args
  with_subelements:
    - thelist
    - fieldname
```

Then Ansible will essentially do this:

```
for thing in thelist:
    item.0 = thing
    for fieldvalue in get(thing, fieldname):
        item.1 = fieldvalue
    EXECUTE (module, args)
```

In other words, it'll iterate over the first value as a list, call it `item.0`, then get the list from that value's field named 'fieldname', and iterate over that as well, calling it `item.1`.

Presumably you could nest this deeper.

Example from the docs. With variables:

```
---
users:
  - name: alice
    authorized:
      - /tmp/alice/onekey.pub
      - /tmp/alice/twokey.pub
  - name: bob
    authorized:
      - /tmp/bob/id_rsa.pub
```

You can write tasks:

```
- user: name={{ item.name }} state=present generate_ssh_key=yes
  with_items: "{{users}}"

- authorized_key: "user={{ item.0.name }} key='{{ lookup('file', item.1) }}'"
  with_subelements:
    - users
    - authorized
```

2.10 Playbook

2.10.1 Playbook directory

Default current dir

2.10.2 Playbook

Syntax A YAML file defining a list of *Play*s and *Playbook include*s:

```
- <play>
- <play>
- include: <path to playbook>
- include: <path to playbook>
```

Templating A playbook is rendered as a Jinja2 template (using [variables in playbooks doc](#)) before processing it, but playbooks should not use loops and conditionals.

2.10.3 Playbook include

A playbook can include other playbooks:

```
- include: <path to playbook>
```

Note that, unlike *Task include* s, playbook includes cannot set variables.

2.10.4 Play

[Complete list of possible keys](#)

A dictionary:

```
hosts:  hosta:pattern1:pattern2    # required
vars:
  var1: value1
  var2: value2
roles:
  - <rolename1>
  - {role: <rolename2>, var1: value1, tags: ['tag1', 'tag2']}
  - role: <rolename3>
    var1: value1
    var2: value2
    tags:
      - tag1
      - tag2
tags:
  - <tag1>
  - <tag2>
remote_user: username
sudo: yes|no
sudo_user: username
tasks:
  - <task>
  - include: <taskfile>
  - include: <taskfile2>
    tags: [tag1, tag2]
  - <task>
handlers:
  - <task>
  - include: <taskfile>
  - <task>
notify:
  - <handler name>
  - <handler name>
vars_files:
  - <path to external variables file>
  - [<path1>, <path2>, ...]    (ansible loads the first one found)
  - <path to external variables file>
strategy: linear|free
serial: <number>| "<number>%"
```

Required keys:

hosts A string, containing one or more *Host Patterns* s separated by colons

Optional keys:

handlers list of *Handler* s and *Task include* s.

pre_tasks list of *Task* s and *Task include* s. These are executed before roles.

roles list of names of *Role* s to include in the play. You can add parameters, tags, and conditionals:

```
roles:
- common
- { role: foo_app_instance, dir: '/opt/a', tags: ["bar", "baz"] }
- { role: foo_app_instance, dir: '/opt/b', when: "ansible_os_family == 'RedHat'
  ↪ " }
```

serial Set how many hosts at a time to run at a time. The default is to run tasks on all of a play's machines at once. See also *strategy*.

strategy How plays are run on multiple hosts. The default is “linear”, where each task is run on up to *serial* hosts in parallel, and then Ansible waits for them all to complete before starting the next task on all the hosts.

“free” lets each host run independently, starting its next task as soon as it finishes the previous one, regardless of how far other hosts have gotten.

tags see *Tags*.

tasks list of *Task* s and *Task include* s. These are executed after the *roles*.

post_tasks list of *Task* s and *Task include* s. These are executed after the *tasks*.

notify list of names of *Handler* s to trigger when done, but only if something changed

vars A dictionary defining additional *Pre-defined variables*

remote_user user to login as remotely

sudo yes/no

sudo_user user to sudo to remotely

2.10.5 Running a playbook

ansible-playbook <filepath of playbook> [options]

ansible-playbook playbook.yml --start-at="install packages" The above will start executing your playbook at a task named “install packages”.

ansible-playbook playbook.yml --step This will cause ansible to stop on each task, and ask if it should execute that task.

2.11 Roles

2.11.1 Role

A role is a directory with specified contents. The role directory must be in one of the directories on the *roles_path* and its name is used to refer to the role elsewhere.

[Complete list of possible keywords](#)

Inside the role's top-level directory, you might see a tree like this (most of this is optional).

defaults/main.yml variables within will be defined at the lowest priority (can be overridden by variables declared anywhere else, even inventory variables)

files/ Any copy tasks can reference files in roles/x/files/ without having to path them relatively or absolutely

Any script tasks can reference scripts in roles/x/files/ without having to path them relatively or absolutely

handlers/main.yml handlers listed therein will be added to the play

library/ modules here (directories) will be available in the role, and any roles called after this role

meta/main.yml *Role dependencies file*

tasks/main.yml *Tasks file*

Any include tasks can reference files in roles/x/tasks/ without having to path them relatively or absolutely

templates/ Any template tasks can reference files in roles/x/templates/ without having to path them relatively or absolutely

vars/main.yml variables listed therein will be added to the play. These override almost any other variables except those on the command line, so this is really better for the role's "constants" than variables :-)

2.11.2 Role dependencies file

Syntax YAML file

Templating Jinja2

Contents A dictionary

The role dependencies file defines what other roles this role depends on.

Keys:

dependencies A list of *Dependency dictionary* s

allow-duplicates yes/no

Defaults to no, preventing the same role from being listed as a dependency more than once. Set to yes if you want to list the same role with different variables.

Example:

```
---
dependencies:
  - role: role1
  - role: role2
    varname: value
```

2.11.3 Dependency dictionary

Required keys:

role

name of role, or quoted path to role file, or quoted repo URL:

```
role: postgres
role: '/path/to/common/roles/foo'
role: 'git+http://git.example.com/repos/role-foo,v1.1,foo'
role: '/path/to/tar/file.tgz,,friendly-name'
```

Optional keys: any parameters for the role - these define *Pre-defined variables*

2.11.4 Embedding modules in roles

2.12 Secrets

Ansible handles secrets using a feature called **Vault**.

Vault lets you encrypt any of your `.yaml` files, but typically you'd apply it to files containing variable definitions, then use the variables' values as needed elsewhere.

Vault provides subcommands that let you encrypt a file in place, decrypt a file in place, edit a file that's encrypted in one step, etc.

When ansible is running your playbook or whatever, any time it comes across a `.yaml` file that appears to be encrypted, it will decrypt it (in memory) and use the decrypted contents, fairly transparently. You can have as many of your files encrypted as you want.

However, all the encrypted files have to use the same password.

2.12.1 Providing the password to Ansible

1. Have Ansible prompt for it by passing `--ask-vault-pass`. Most secure, but inconvenient.
2. Put it plaintext in a well-protected file, and pass `--vault-password-file <filename>`. Most insecure, but more convenient than the prompt.
3. Write a script or program that outputs the password on stdout, mark it executable, and pass that: `--vault-password-file <path-to-program>`. This makes it possible to use a local system key-chain or something, which might be more secure than the other options. Or worse...

2.12.2 Ways to use it

One approach I've used is to have a single encrypted `secrets.yaml` file in my base directory containing all my secret variables, and another file with very restrictive permissions (and outside of source control) containing my password, then add these arguments when running ansible:

```
--extra-vars @secrets.yaml --vault-password-file path/to/passfile
```

The advantage of that is that if I don't need the secrets, I can leave all that off and Ansible will run fine. (As opposed to having the encrypted file automatically read by Ansible every time.)

I'm not sure if that will scale, though.

2.12.3 Limitations

- This is symmetric encryption. In other words, anyone with the password can encrypt and decrypt a file.
- All the encrypted files must be encrypted using the same password.
- That means you have to protect the decrypting password (the only password) very carefully, and makes providing it to Ansible awkward.

Links:

- [Vault](#)
- [Not logging secrets](#)
- [How to upload encrypted file using ansible vault?](#)
- [Managing Secrets with Ansible Vault – The Missing Guide \(Part 1 of 2\)](#)
- [Managing Secrets with Ansible Vault – The Missing Guide \(Part 2 of 2\)](#)

2.13 synchronize

The Ansible synchronize module gets its own page because it is a bitch.

(Update: apparently some of these bad behaviors were bugs in Ansible 2.0.0.x, but I'm keeping this page around for history.)

Let me count the ways:

- By default, it tries to become *locally* the user you've specified using the `become_user` variable that you have said you want to become *remotely*. [Apparently that was a bug in 2.0.0.x and works correctly in 1.9.x and 2.0.1+.]
- Then it does *not* try to *remotely* become the user you've specified; you have to hack it by setting `rsync_path`: `"sudo rsync"`. [I have not tried this again with 2.0.1+.]
- Unlike every other Ansible module, the `owner` and `group` options are *booleans*, not the names or numbers of users and groups. If true, it'll try to copy the owner of the local files, but if you want to specify the ownership of the target files yourself, you'll have to fix it afterward.

Here's a working example:

```
- name: sync source from local directory
  synchronize:
    dest: "{{ source_dir }}"
    src: "{{ local_project_dir }}"
    delete: yes
    rsync_path: "sudo rsync" # Use sudo on the remote system
    recursive: true
    rsync_opts:
      - "--exclude=.git"
      - "--exclude=*.pyc"
    become: no # stops synchronize trying to sudo locally
```

NOTE: Ansible 2.0.1 fixed numerous bugs in synchronize:

- Fixes a major compatibility break in the synchronize module shipped with 2.0.0.x. That version of synchronize ran `sudo` on the controller prior to running `rsync`. In 1.9.x and previous, `sudo` was run on the host that `rsync` connected to. 2.0.1 restores the 1.9.x behaviour.

- Additionally, several other problems with where synchronize chose to run when combined with `delegate_to` were fixed. In particular, if a playbook targetted localhost and then delegated_to a remote host the prior behavior (in 1.9.x and 2.0.0.x) was to copy files between the src and destination directories on the delegated host. This has now been fixed to copy between localhost and the delegated host.
- Fix a regression where synchronize was unable to deal with unicode paths.
- Fix a regression where synchronize deals with inventory hosts that use localhost but with an alternate port.

2.14 Tags

When you apply tags to things, you can then control whether they're executed by adding command line options.

2.14.1 How to tag things

Plays and tasks have optional `tags` attributes where you can specify a list of tags. Here are some tagged *Task* s:

```
tasks:
  - module: parm1=a parm2=b
    tags:
      - packages

  - module2: parm1=x parm2=y
    tags:
      - configuration
```

And here's a playbook with some tagged *Play* s:

```
- hosts: all
  tags:
    - foo
    - bar
  roles:
    - role1
    - role2
```

You can also apply tags when invoking a role from a playbook:

```
roles:
  - { role: webserver, port: 5000, tags: [ 'web', 'foo' ] }
```

and when including tasks:

```
- include: foo.yml
  tags: [web,foo]
```

2.14.2 What tags do

Adding a tag to a play or task says that *if* ansible is invoked with `--tags=x,y,z`, that the tagged play or task will only be executed if at least one of its tags is included in the list of tags from the command line.

Specifying `--tags=all` is equivalent to the default behavior, where all playbooks and tasks are run regardless of their tags.

Specifying `--tags=tagged` runs only things that have *some* tag, while `--tags=untagged` runs only things that have *no* tag.

You could alternatively invoke ansible with `--skip-tags=a,b,c` and it will execute all plays and tasks that are *not* tagged with a, b, or c.

Presumably `--skip-tags=tagged` does the opposite of `--tags=tagged`, and `--skip-tags=untagged` does the opposite of `--tags=untagged`.

If a play or task is tagged `always`, then it will be executed *unless* ansible is invoked with `skip-tags=always`.

2.15 Task

2.15.1 Tasks file

Syntax YAML FILE

Templating Jinja2

Content A list of task definitions, task includes, and *Blocks*.

2.15.2 Task include

Anywhere there can be a task definition, you can also use a task include:

```
- include: <path to tasks file> [options]
```

The path is relative to the *Playbook directory*, or the file is also searched for in the tasks directory of a role.

[options] is an optional list of additional variable settings, e.g.:

```
- include: tasks/footasks.yml vara=1 varb=2 varc=3
```

You can use an expanded syntax with a vars setting to set more complicated values:

```
- include: wordpress.yml
  vars:
    wp_user: timmy
    some_list_variable:
      - alpha
      - beta
      - gamma
```

Or use this more compact but apparently equivalent syntax:

```
- { include: wordpress.yml, wp_user: timmy, ssh_keys: [ 'keys/one.txt', 'keys/two.txt' ] }
```

2.15.3 Task

[ansible tasks doc](#), complete list of possible keywords

A dictionary:

```
name: string      # optional but highly recommended
module: args      # required; the "action"
environment: dictionary
remote_user: username
sudo: yes|no
sudo_user: username
otheroption: othervalue # depending on module
tags:
- <tag1>
- <tag2>
```

Required keys:

name text

modulename options

Optional keys that can be used on any task:

environment dictionary (in YAML, or variable containing dictionary)

ignore_errors if true, continue even if task fails

register <varname> store result of task in <varname>. See also [when](#) for some ways to use.

remote_user user to login as remotely

sudo yes|no

sudo_user user to sudo to remotely

tags list of tags to associate with the task

when expression controls whether task is executed [ansible when doc](#):

```
when: <varname>
when: not <varname>
```

Special filters for checking result of a prior task:

```
when: <varname>|failed
when: <varname>|skipped
when: <varname>|success
```

Additional keys might be required and optional depending on the module being used.

2.15.4 Handler

Same syntax as a [Task](#), it just gets triggered under different circumstances.

2.16 Variables

2.16.1 Pre-defined variables

Ansible defines some variables for you

These are not mentioned when you list Facts (see below) - go figure.

`inventory_hostname` is the name of the current host as you've configured it in your Ansible inventory file, regardless of the system's actual hostname.

If you have a long FQDN, `inventory_hostname_short` also contains the part up to the first period, without the rest of the domain.

2.16.2 Variables

Some variables alter the behavior of ansible (see http://docs.ansible.com/intro_inventory.html#list-of-behavioral-inventory-parameters for a list). You can set some of these using environment variables ([ansible variables doc](#)).

CORRECTION: Use `ansible_ssh_user`, not `ansible_user`.

Any of them can be used anywhere Jinja2 templating is in effect.

Places to define variables:

- inventory
- playbook
- included files and roles
- local facts
- ansible command line (`--extra-vars "foo=1 bar=2"` or `--extra-vars @filepath.json` or `--extra-vars @filepath.yml`)

See also “Variable Precedence”, a little farther down...

2.16.3 Variables file

A variables file ([doc](#)) is a file that defines values of *Pre-defined variables*.

Syntax YAML defining a single dictionary

Templating The file does not appear to undergo template expansion, but the values of variables do??

2.16.4 Variables files

Ansible will look in *Inventory directory* and *Playbook directory* for directories named `host_vars` or `group_vars`. Inside those directories, you can put a single *Variables file* with the same name as a host or group (respectively) and Ansible will use those *Pre-defined variables* definitions.

Or a file named `all` that will be used for all hosts or groups.

Or you can create a directory with the same name as a host or group and Ansible will use all the files in that directory as *Variables files*.

You can also include vars files from a *Play* ([ansible variable files doc](#)).

2.16.5 Variable precedence

[docs](#)

From 2.0 on, from lowest priority to highest - in other words, if a variable is defined in two places, the place that's farther down in this list takes precedence.

- role defaults [1]
- inventory file or script group vars [2]
- inventory group_vars/all [3]
- playbook group_vars/all [3]
- inventory group_vars/* [3]
- playbook group_vars/* [3]
- inventory file or script host vars [2]
- inventory host_vars/*
- playbook host_vars/*
- host facts / cached set_facts [4]
- inventory host_vars/* [3]
- playbook host_vars/* [3]
- host facts
- play vars
- play vars_prompt
- play vars_files
- role vars (defined in role/vars/main.yml)
- block vars (only for tasks in block)
- task vars (only for the task)
- include_vars
- set_facts / registered vars
- role (and include_role) params
- include params
- extra vars (defined on command line with `-e`, always win precedence)

[1] Tasks in each role will see their own role's defaults. Tasks defined outside of a role will see the last role's defaults.

[2] (1, 2) Variables defined in inventory file or provided by dynamic inventory. [3] (1, 2, 3, 4, 5, 6) Includes vars added by 'vars plugins' as well as host_vars and group_vars which are added by the default vars plugin shipped with Ansible.

[4] When created with set_facts's cacheable option, variables will have the high precedence in the play, but will be the same as a host facts precedence when they come from the cache.

2.16.6 Facts

Ansible automatically defines a whole bunch of variables with information about the system that it's running on (the system the plays and tasks are running on, not the system you're controlling ansible from).

You can add to the facts with config files called local facts ([ansible local facts doc](#)) though I don't know how that's any better than putting variables in all the other places you can set them...

To see a list of all of the facts that are available about a machine, you can run the "setup" module as an ad-hoc action:

```
ansible -m setup hostname
```


This will print out a dictionary of all of the facts that are available for that particular host.

The Ansible docs used to show an example of this output, but apparently they've removed or moved that. And here's [an example](#).

The top of the output will look like:

```
staging-web2 | SUCCESS => {
  "ansible_facts": {
    "ansible_all_ipv4_addresses": [
      "10.132.77.14",
      "138.197.111.207",
      "10.17.0.12"
    ],
    "ansible_all_ipv6_addresses": [
```

Ignore the "ansible_facts" part of that. To reference any of these variable, start with the next level. E.g. {{ ansible_all_ipv4_addresses[1] }}.

ALTERNATIVELY, you can access the same variables as items in the `ansible_facts` dictionary, only without the individual keys prefixed by `ansible_` (or so the docs say https://docs.ansible.com/ansible/latest/reference_appendices/config.html#inject-facts-as-vars) and this should work even if `INJECT_FACTS_AS_VARS` has been set `False`).

2.17 Ansible Galaxy

2.17.1 Links

- [Galaxy doc](#)
- [popular galaxy roles and recent activity](#)
- [search galaxy](#)

2.17.2 Role specification

Format when installing roles from galaxy:

- `username.rolename[,version]`
- `scm+repo_url[,version]`
- `tarball_url`

Versions represent tags in the role's source repository.

E.g.:

```
user2.role2
user1.role1,v1.0.0
user1.role2,master
git+http://bitbucket.org/willthames/git-ansible-galaxy
https://some.webserver.example.com/files/master.tar.gz
```

2.17.3 Ways of installing

Command-line

List roles on the command line:

```
ansible-galaxy install user2.role2 user1.role1,v1.0.9
```

Simple file

List roles in a file, one per line. Example file:

```
# file: roles.txt
user2.role2
user1.role1,v1.0.0
```

And install with `-r`:

```
ansible-galaxy install -r roles.txt
```

YAML file

Use a YAML file to provide more control. The YAML file should contain a list of dictionaries. Each dictionary specifies a role to install. Keys can include:

src (*required*) a role specification as above. (Since there's a separate dictionary key for `version`, I don't know whether you can include version here, or if you're required to list it separately as `version`.)

path Where to install (directory, can be relative)

version version to install. e.g. `master` or `v1.4`.

name install as a different name

scm default `git` but could say `hg` and then in `src` provide a URL to a mercurial repository.

Example:

```
# install_roles.yml

# from galaxy
- src: yatesr.timezone

# from github
- src: https://github.com/bennojoy/nginx

# from github installing to a relative path
- src: https://github.com/bennojoy/nginx
  path: vagrant/roles/

# from github, overriding the name and specifying a specific tag
- src: https://github.com/bennojoy/nginx
  version: master
  name: nginx_role

# from a webserver, where the role is packaged in a tar.gz
```

(continues on next page)

(continued from previous page)

```
- src: https://some.webserver.example.com/files/master.tar.gz
  name: http-role

# from bitbucket, if bitbucket happens to be operational right now :)
- src: git+http://bitbucket.org/willthames/git-ansible-galaxy
  version: v1.4

# from bitbucket, alternative syntax and caveats
- src: http://bitbucket.org/willthames/hg-ansible-galaxy
  scm: hg
```

And again install with `-r`:

```
ansible-galaxy install -r install_roles.yml
```

Misc. stuff I need to file somewhere:

2.18 Ad-hoc command

ansible *Host Patterns* -m <module> [options]

e.g.

```
$ ansible all -m ping --ask-pass
```

Shortcut to run a command:

```
$ ansible all -a "/bin/echo hello"
```

options: see output of “ansible --help” for now

See [ansible ad-hoc commands doc](#) for ad-hoc commands.

Contents:

3.1 S3

3.1.1 Access control

- How S3 evaluates access control
- Guidelines for Using the Available Access Policy Options

“The only recommended use case for the bucket ACL is to grant write permission to the Amazon S3 Log Delivery group”...

“In general, you can use either a user policy or a bucket policy to manage permissions.”

Here’s a bucket policy to grant some IAM user complete access to a bucket:

```
{
  "Statement": [
    {
      "Sid": "PublicReadForGetBucketObjects",
      "Effect": "Allow",
      "Principal": {
        "AWS": "*"
      },
      "Action": ["s3:GetObject"],
      "Resource": ["arn:aws:s3:::BUCKET-NAME/*"]
    },
    {
      "Action": "s3:*",
      "Effect": "Allow",
      "Resource": [
```

(continues on next page)

(continued from previous page)

```
        "arn:aws:s3::BUCKET-NAME",
        "arn:aws:s3::BUCKET-NAME/*"
    ],
    "Principal": {
        "AWS": [
            "USER-ARN"
        ]
    }
}
]
```

What about read-only access? Let's see...

seems like s3auth.com used this example:

```
{
  "Statement": [
    {
      "Effect": "Allow",
      "Action": ["s3:GetObject", "s3:GetBucketWebsite"],
      "Resource": [
        "arn:aws:s3::bucket-1.example.com/*",
        "arn:aws:s3::bucket-2.example.com/*"
      ]
    }
  ]
}
```

3.1.2 Updating metadata to improve response headers for caching

Install `s3cmd`, then do it like this:

```
s3cmd --recursive modify \
  --add-header="Expires: Thu, 31 Dec 2099 20:00:00 GMT" \
  --add-header="Cache-Control: max-age=94608000" \
  s3://cactus-website-production-2015/media/community_logos
```

You can use `s3cmd ls` to get a list of the buckets you can access.

CHAPTER 4

Bootstrap

Warning: THIS IS NOT DONE AND PROBABLY WRONG.

The grid.

SO what does a class *col-SIZE-N* mean?

Each SIZE has a BREAKPOINT:

xs: -1 sm: 750px md: 970px lg: 1170px

Call the window width WIDTH.

For a single class col-SIZE-N:

```
if WIDTH >= BREAKPOINT(SIZE), then
    ELEMENT-WIDTH=WIDTH*N/12
    display INLINE (same line as previous element if possible)
else
    ELEMENT-WIDTH=100%
    display BLOCK (element gets its own line)
```

What if we have col-SIZE1-N1 and col-SIZE2-N2, with BREAKPOINT(SIZE1) < BREAKPOINT(SIZE2)?:

```
IF WIDTH >= BREAKPOINT(SIZE2), then
    ELEMENT_WIDTH = WIDTH * N2 / 12
    INLINE
elif WIDTH >= BREAKPOINT(SIZE1), then
    ELEMENT_WIDTH = WIDTH * N1 / 12
    INLINE
else:
    BLOCK display
```

and so forth - just look at the class with the largest size

NOTE: Since all widths are \geq the breakpoint of XS, then if XS is present, the element will ALWAYS be laid out inline. Though col-xs-12 is pretty much equivalent to not having an XS class, right??????????/

5.1 Services

update-rc.d:

- Make a service run in its default runlevels:

```
update-rc.d <service> defaults
```

or:

```
update-rc.d <service> enable
```

- Make a service not run in any runlevel:

```
update-rc.d <service> disable
```

Making a new init script:

- Read `/etc/init.d/README`, which will point to other docs
- Copy `/etc/init.d/skeleton` and edit it.

5.2 Packages

- List packages that match a pattern: `dpkg -l <pattern>`
- List contents of a package: `dpkg -L packagename`
- Show packages that installed files matching pattern: `dpkg -S pattern`
- Show info about an installed package: `dpkg-query -s packagename`
- show info about a package that is known: `apt-cache showpkg packagename`
- Reconfigure a package: `dpkg-reconfigure packagename`

- Change alternatives: `update-alternatives ...`

5.3 Alternatives

Change 'alternatives' default browser or editor:

```
sudo update-alternatives --set x-www-browser /usr/bin/chromium-browser
sudo update-alternatives --set editor /usr/bin/emacs24
```

Be prompted for which alternative you prefer for a link group:

```
sudo update-alternatives --config x-www-browser
```

Find out what the top-level link groups are:

```
sudo update-alternatives --get-selections
```

Set xdg program to open/browse a directory (DOES NOT WORK) (do NOT use sudo):

```
xdg-mime default /usr/share/applications/Thunar.desktop x-directory/normal
```

Change 'xdg' default browser (for user):

```
xdg-settings get default-web-browser
xdg-settings set default-web-browser google-chrome.desktop
xdg-settings set default-web-browser firefox.desktop
```

Install without any prompts (<http://askubuntu.com/questions/146921/how-do-i-apt-get-y-dist-upgrade-without-a-grub-config-prompt>):

```
sudo DEBIAN_FRONTEND=noninteractive apt-get -y \
-o Dpkg::Options::="--force-confdef" -o Dpkg::Options::="--force-confold" \
<COMMAND>
```

5.4 Desktop applications

Put your own .desktop files in `~/.local/share/applications`.

[Archlinux on desktop entries](#)

[Desktop file spec](#)

To let the system know about new or changed desktop files:

```
update-desktop-database [directory]
```

Launch the application from command line that has a `<name>.desktop` file somewhere:

```
gtk-launch <name>
```

6.1 Foods with low glycemic index

- Breads
- Dense wholegrain breads
- White corn tortillas
- Grain and seed breads
- Fruit Loaf such as Raisin
- Multigrain breads (look for breads where you can see lots of grains)
- Authentic Sourdough bread
- Breakfast Cereals
- Traditional porridge oats
- Muesli*
- Bircher Muesli
- Wholegrain high fibre cereals
- Vegetables
- Sweetcorn
- Silverbeet
- Carrots
- Zucchini
- Peas, frozen or fresh
- Snowpeas
- Carisma™ Potatoes*

- Green Beans
- Broccoli
- Eggplant
- Cauliflower
- Squash
- Capsicum
- Salad Vegetables
- Celery
- Leeks
- Tomatoes
- Mushrooms – very low carb or no GI rating
- Butternut Pumpkin (lower GI)
- Avocadoes
- Drinks
- Milo®
- Skim Latte
- Sustagen®
- Soy Drinks
- Fruit Smoothies
- Fruit Juice
- Snacks
- Grain & Fruit bars
- Wholegrain crackers
- Nut & Seed bars
- Dried fruit and nuts
- Legumes
- Split Peas; Green or red Lentils
- Baked Beans
- Canned & Dried beans – kidney, cannellini, butter, borlotti, chickpeas
- Spreads
- Fruit Spreads
- Hummus
- Nut butters
- Main Meal Carbs
- Doongara Low GI White rice
- Fresh Noodles – Hokkein, Udon, Rice

- Low GI Brown rice*
- Soba Noodles
- Basmati rice (lower GI)
- Buckwheat
- Pasta, cooked al dente*
- Vermicelli
- Pearl Couscous*
- Bulgur
- Quinoa*
- Semolina
- Pearl Barley
- Cracked Wheat
- Fruit
- Apples*
- Pears*
- Bananas
- Kiwi Fruit
- Grapes*
- Mango
- Strawberries
- Oranges
- Peaches
- Grapefruits
- Apricots
- Berries, fresh or frozen
- Plums
- Dried fruits such as prunes, raisins, sultanas, apricots
- Canned Fruit in natural juice
- Dairy Foods
- Reduced fat milk
- Reduced fat custard
- Reduced fat yoghurt, plain or fruit flavoured
- Low fat ice-cream*

6.2 For lowering triglycerides

- Decrease or eliminate:
 - Sweets
 - Alcohol
 - Refined carbohydrates:
 - White rice
 - bread and pasta made from white flour or semolina
 - Saturated fats and fried foods:
 - high fat meats
 - skin on poultry
 - sauces and spreads
 - Trans fatty acids and hidden fats:
 - hydrogenated vegetable oil
 - regular fat meats
 - lunchmeats
 - hot dogs
 - fatty snack foods
- Eat more:
 - omega 3 fatty acids:
 - fatty fish
 - salmon
 - mackerel
 - sardines
 - tuna
 - trout
 - ground flax seed
 - flaxseed oil
 - soy products
 - legumes
 - walnuts
 - dark leafy green vegetables
 - high fiber foods:
 - beans
 - whole grains
 - ground flaxseed

- pumpkin seeds
- rice bran
- oat bran
- fruits and vegetables
- Eat more plant foods: Vegetable proteins such as
- dried beans,
- peas, and
- soy products;
- White poultry, prepared without the skin, is also a good source of protein without a lot of fat content.

CHAPTER 7

Django

These are just things I always find myself looking up, so I try to make some notes of the most important parts that I can refer back to.

Contents:

7.1 Admin

7.1.1 URLs

List `{{ app_label }}` `_{{ model_name }}` `_changelist` Change `{{ app_label }}` `_{{ model_name }}` `_change` `object_id`
<https://docs.djangoproject.com/en/stable/ref/contrib/admin/#reversing-admin-urls>

7.1.2 Customize top-right corner of admin pages

Create your own `templates/admin/base_site.html` that comes ahead of the admin's default one in the templates path.

At least in Django 1.8+, this gives you a “View site” link for free:

```
% extends "admin/base.html" %}

{% block title %}{{ title }} | {{ site_title|default:_('Django site admin') }}{% _
↪endblock %}

{% block branding %}
<h1 id="site-name"><a href="{% url 'admin:index' %}">{{ site_header|default:_('Django_
↪administration') }}</a></h1>
{% endblock %}

{% block userlinks %}
  <a href="{% url 'clear-cache' %}">Clear cache</a> /
```

(continues on next page)

(continued from previous page)

```
{{ block.super }}
{% endblock userlinks %}
```

Prior to Django 1.8:

```
{% extends "admin/base.html" %}

{% block title %}{{ title }} | Caktus Admin{% endblock %}

{% block branding %}<h1 id="site-name">Caktus Admin</h1>{% endblock %}

{% block nav-global %}
    <div style='display:block; padding:0 1em 0.5em 1em; float:right;'>
        <a href='{% url "home" %}'>Return to Caktus Home</a>
        | <a href='{% url "clear-cache" %}'>Clear cache</a>
    </div>
{% endblock %}
```

7.2 Applications

<https://docs.djangoproject.com/en/stable/ref/applications/#django.apps.AppConfig>

In `__init__.py`:

```
# programs/__init__.py

default_app_config = 'programs.apps.ProgramsConfig'
```

In `apps.py`:

```
# programs/apps.py

from django.apps import AppConfig

class ProgramsConfig(AppConfig):
    name = 'programs' # required: must be the Full dotted path to the app
    label = 'programs' # optional: app label, must be unique in Django project
    verbose_name = "Rock 'n' roll" # optional

    def ready():
        """
        This runs after all models have been loaded, but you may not
        modify the database in here.

        Here's a trick to run something after each migration, which is often
        good enough.
        """
        from django.db.models.signals import post_migrate

        post_migrate.connect(`callable`)
```

7.3 Celery

(Yes, I know Celery isn't Django-specific.)

<http://docs.celeryproject.org/en/latest/>

7.3.1 Useful settings

<http://docs.celeryproject.org/en/latest/configuration.html>

CELERY_ALWAYS_EAGER: If this is True, all tasks will be executed locally by blocking until the task returns. `apply_async()` and `Task.delay()` will return an `EagerResult` instance, which emulates the API and behavior of `AsyncResult`, except the result is already evaluated.

That is, tasks will be executed locally instead of being sent to the queue.

CELERY_EAGER_PROPAGATES_EXCEPTIONS: If this is True, eagerly executed tasks (applied by `task.apply()`, or when the **CELERY_ALWAYS_EAGER** setting is enabled), will propagate exceptions.

It's the same as always running `apply()` with `throw=True`.

CELERY_IGNORE_RESULT: Whether to store the task return values or not (tombstones). If you still want to store errors, just not successful return values, you can set **CELERY_STORE_ERRORS_EVEN_IF_IGNORED**.

CELERYD_HIJACK_ROOT_LOGGER: By default any previously configured handlers on the root logger will be removed. If you want to customize your own logging handlers, then you can disable this behavior by setting **CELERYD_HIJACK_ROOT_LOGGER** = False.

CELERYBEAT_SCHEDULE: In each task, you can add an 'options' dictionary and set 'expires' to a number of seconds. If the task doesn't run within that time, it'll be discarded rather than run when it finally gets to a worker. This can help a lot with periodic tasks when workers or the queue gets hung up for a while and then unjammed - without this, the workers will have to work through a huge backlog of the same periodic tasks over and over, for no reason.

Example:

```
CELERYBEAT_SCHEDULE = {
    'process_new_scans': {
        'task': 'tasks.process_new_scans',
        'schedule': timedelta(minutes=15),
        'options': {
            'expires': 10*60, # 10 minutes
        }
    },
}
```

CELERY_DEFAULT_QUEUE: In the absence of more complicated configuration, celery will use this queue name for everything. Handy when multiple instances of a site are sharing a queue manager:

```
CELERY_DEFAULT_QUEUE = 'queue_%s' % INSTANCE
```

7.4 django-compressor

[django-compressor docs](#)

Warning: much of the documentation is casual about saying things that are only true in some scenarios, without making clear that that's the case.

7.4.1 ACTUALLY USING

Here are some practical scenarios for using django-compressor.

For what to put in your templates, you can go by the django-compressor documentation, and be sure to use `{% static %}` and not `STATIC_URL`.

For what to put in your settings... it's a lot more complicated. Set the compressor filters and precompilers however you want. For the rest, keep reading.

Scenario: Development using runserver, DEBUG, not offline

If `DEBUG` is `True`, then compressor won't even do anything and so everything should just work.

Scenario: Running using local files, not offline

This is the typical small server situation. You unpack your project on the server, run `collectstatic`, point nginx or some other server at `STATIC_ROOT` and go.

Example settings:

```
# Django settings
DEBUG = False
STATIC_ROOT = '/var/www/static/'
STATIC_URL = '/static/'
# set compressor filters and precompilers as desired.
# leave other compressor settings at defaults.
```

```
# nginx settings
location /static {
    alias /var/www/static;
}
```

Scenario: running using local files, with offline

Like the previous scenario, but you want compressor to do all its work at deploy time so the results are cached and ready to go immediately when you start your server.

```
# Django settings like before, plus:
COMPRESS_OFFLINE = True
```

Now at deploy time you have more steps:

```
$ python manage.py collectstatic
$ python manage.py compress
```

Run `compress` *after* `collectstatic` so that compressor can find its input files. It'll write its output files under `{STATIC_ROOT}/CACHE`, and get them from there at runtime.

Scenario: running with storage on the network, with offline

In this scenario, you're putting your static files somewhere off of the server where you're running Django. For example, S3. Or just your own static file server somewhere. Whatever.

Let's start with how this would be setup without django-compressor, then we can modify it to add django-compressor.

```
# settings/no_compressor.py
STATIC_ROOT = None # Unused
STATIC_URL = None # Unused
STATIC_FILE_STORAGE = 'package.of.FileStorageClass'
```

At deploy time you can just run collectstatic, and all your static files will be pushed to the network:

```
$ python manage.py collectstatic
```

And at runtime, `{% static %}` will ask your file storage class to come up with a URL for each file, which will turn out to be on your other server, or S3, or whatever.

Now, suppose we want to add compressor with offline processing (not using offline makes no sense with network storage). Here are the settings you can use at runtime for that, assuming things have been prepared correctly:

```
# settings/deployed.py
# Django settings we'll use in production
STATIC_ROOT = None # Unused
STATIC_URL = None # Unused
STATIC_FILE_STORAGE = 'path.to.network.filestorage'
COMPRESS_ENABLED = True
COMPRESS_OFFLINE = True
```

The preparation is the tricky part. It turns out that for compress to work, a copy of the static files must be gathered in a local directory first. Most of the tools we might use to compile, compress, etc. are going to read local files and write local output.

To gather the static files into a local directory, we might, for example, use a different settings file that uses the default file storage class, and run collectstatic. E.g.:

```
# settings/gather.py
# Django settings when first running collectstatic
from .deployed import *

# Override a few settings to make storage local
STATIC_ROOT = '/path/to/tmp/dir'
STATIC_URL = None # Unused
STATIC_FILE_STORAGE = 'django.core.files.storage.FileSystemStorage'
```

```
$ python manage.py collectstatic --settings=settings.gather
```

After running collectstatic with these settings, all your source static files will be gathered under `/path/to/tmp/dir`.

Now you could run compress, and the resulting files would be added under `/path/to/tmp/dir`. There's an important *gotcha* that will cause problems, though - for compressor to match up the output it makes now with what it'll be looking for later, the contents of each `{% compress %}` tag must be identical now to what it'll be then, which means the URLs must point at the production file server. We can accomplish this by setting `STATIC_URL` before running the compress:

```
# settings/compress.py
# Django settings when running compress
from .deployed import *

# Override a few settings to make storage local, but URLs look remote
STATIC_ROOT = '/path/to/tmp/dir'
```

(continues on next page)

(continued from previous page)

```
STATIC_URL = 'https://something.s3.somewhere/static/' # URL prefix for runtime
STATIC_FILE_STORAGE = 'django.core.files.storage.FileSystemStorage'
```

```
$ python manage.py compress --settings=settings.compress
```

The problem now is to get all these files onto the remote server. You could just use `rsync` or `s3cmd` or something, which will work fine. But for maximum flexibility, let's figure out a way to do it using Django. Our approach will be to tell Django that our SOURCE static files are in `"/path/to/tmp/dir"`, and we want them collected using our production file storage class, which will put them where we want them.

```
# Django settings when running collectstatic again after compress,
# to copy the resulting files to the network
# settings/copy_upstream.py
from .deployed import * # Set up for network file storage
# Tell collectstatic to use the files we collected and compressed
STATICFILES_FINDERS = ['django.contrib.staticfiles.finders.FileSystemFinder']
STATICFILES_DIRS = ['/path/to/tmp/dir']
```

```
$ python manage.py collectstatic --settings=settings.copy_upstream
```

That should copy things to the network. Then if you run using the 'deployed' settings, things should work!

TODO: TEST THAT!!!!!!!!!!!!!!!!!!!!!!

Other approaches

The compressor docs suggest a different approach – hack the storage class you're using so when you run `collectstatic`, it saves a copy of each file into a local directory in addition to pushing it upstream. Then you can use the same storage class for `collectstatic`, `compress`, and runtime.

7.4.2 More detailed notes

Cache

For some things, compressor uses the cache named by `COMPRESS_CACHE_BACKEND`, which defaults to `None`, which gives us the default Django cache.

Principles of compression

Whether compressor is processing templates offline ahead of time or at runtime, there are some common principles.

First, if `COMPRESS_ENABLED` is `False`, the `{% compress %}` tag will simply render as its contents; compressor won't change anything.

Otherwise, compressor will

1. parse the contents of the tag and figure out which css and javascript files would be included
2. fetch those files (See "accessing the files to be compressed")
3. run those files through any configured preprocessors
4. concatenate the result and save it using `COMPRESS_STORAGE`

5. at rendering, the tag and contents will be replaced with one or two HTML elements that will load the compressed file instead of the original ones.

Offline

If `COMPRESS_OFFLINE` is `True`, compressor expects all uses of `{% compress ... %}` in templates to have been pre-processed by running `manage.py compress` ahead of time, which puts the results in compressor's *offline* cache. If anything it needs at run-time is not found there, things break/throw errors/render wrong etc.

Note: If `COMPRESS_OFFLINE` is `True` and files have not been pre-compressed, compressor will *not* compress them at runtime. Things will break.

The offline cache manifest is a json file, stored using `COMPRESS_STORAGE`, in the subdirectory `COMPRESS_OUTPUT_DIR` (default: `CACHE`), using the filename `COMPRESS_OFFLINE_MANIFEST` (default: `manifest.json`).

The keys in the offline cache manifest are generated from *the template content inside each compress tag*, *not* the contents of the compressed files. So, you must arrange to re-run the offline compression anytime your content files might have changed, or it'll be serving up compressed files generated from the old file contents.

Note: It sounds like you must *also* be *sure* the contents of the compress tags don't change between precompressing and runtime, for example by changing the URL prefix!

The values in the offline cache manifest are paths of the compressed files in `COMPRESS_STORAGE`.

Note: RECOMMENDATION FROM DOCS: make `COMPRESS_OFFLINE_MANIFEST` change depending on the current code revision, so that during deploys, servers running different versions of the code will each use the manifest appropriate for the version of the code they're running. Otherwise, servers might use the wrong manifest and strange things could happen.

Not offline

If `COMPRESS_OFFLINE` is `False`, compressor will look in `COMPRESS_STORAGE` for previously processed results, but if not found, will create them on the fly and save them to use again.

Storage

Compressor uses a [Django storage class](#) for some of its operations, controlled by the setting `COMPRESS_STORAGE`.

The default storage class is `compressor.storage.CompressorFileStorage`, which is a subclass of the standard filesystem storage class. It uses `COMPRESS_ROOT` as the base directory in the local filesystem to store files in, and builds URLs by prefixing file paths within the storage with `COMPRESS_URL`.

If you change `COMPRESS_STORAGE`, then *ignore* anything in the docs about `COMPRESS_ROOT` and `COMPRESS_URL` as they won't apply anymore (except in a few cases... see exceptions noted as they come up, below).

Accessing the files to be compressed

For each file to be compressed, compressor starts with the URL from the rendered original content inside the compress tag. For example, if part of the content is `<script src="http://example.com/foo.js"></script>`, then it extracts `"http://example.com/foo.js"` as the URL.

It checks that the URL starts with `COMPRESS_STORAGE`'s `base_url`, or if accessing that fails (quite possible since `base_url` is not a standard part of the file storage class API), uses `COMPRESS_URL`.

Note: This is a place where compressor can use `COMPRESS_URL` even if it's not using its default storage.

If the URL doesn't start with that string, compressor throws a possibly misleading error, `"'%s' isn't accessible via COMPRESS_URL ('%s') and can't be compressed"`.

Otherwise, compressor tries to come up with a local filepath to access the file, as follows:

- Try to get a local filepath from `COMPRESS_STORAGE` using `.path()`.
- If that's not implemented (for example, for remote storages), it tries again using `compressor.storage.CompressorFileStorage` (regardless of what `COMPRESS_STORAGE` is set to), so basically it's going to look for it under `COMPRESS_ROOT`.
- If it still can't get a local filepath, throws an error: `"'%s' could not be found in the COMPRESS_ROOT '%s'%s"` which is very misleading if you're not using a storage class that looks at `COMPRESS_ROOT`.

7.5 Data fixtures

Export/dump data to use as a fixture:

```
python manage.py dumpdata --format=yaml --natural app.model >data.yaml
```

Load it again:

```
python manage.py loaddata data.yaml
```

7.5.1 Natural keys

<https://docs.djangoproject.com/en/stable/topics/serialization/#natural-keys>

```
from django.db import models

class PersonManager(models.Manager):
    def get_by_natural_key(self, first_name, last_name):
        return self.get(first_name=first_name, last_name=last_name)

class Person(models.Model):
    objects = PersonManager()
    ...

    def natural_key(self):
        return (self.first_name, self.last_name)

class Meta:
    unique_together = (('first_name', 'last_name'),)
```


Dependencies

If part of the natural key is a reference to another model, then that model needs to be deserialized first:

```
class Book(models.Model):
    name = models.CharField(max_length=100)
    author = models.ForeignKey(Person)

    def natural_key(self):
        return (self.name,) + self.author.natural_key()
    natural_key.dependencies = ['example_app.person']
```

7.6 Databases

7.6.1 Performance

From Django 1.6 on, always add `CONN_MAX_AGE` to database settings to enable persistent connections. `300` is a good starting value (5 minutes). `None` will keep them open indefinitely.

BUT - keep in mind that every open connection to Postgres consumes database server resources. So you might want *instead* to run pgbouncer locally.

7.7 Django Debug Toolbar

Install/config

Install:

```
pip install django-debug-toolbar
```

settings.py:

```
DEBUG = True
INTERNAL_IPS = ['127.0.0.1']
INSTALLED_APPS += [
    'debug_toolbar',
]
# The order of MIDDLEWARE and MIDDLEWARE_CLASSES is important. You should include
# the Debug Toolbar middleware as early as possible in the list. However, it must
# come after any other middleware that encodes the response's content, such as
# GZipMiddleware.
MIDDLEWARE = [
    'debug_toolbar.middleware.DebugToolbarMiddleware',
] + MIDDLEWARE
```

urls.py:

```
from django.conf import settings
from django.conf.urls import include, url

if settings.DEBUG:
    import debug_toolbar
    urlpatterns += [
```

(continues on next page)

(continued from previous page)

```
url(r'^__debug__/', include(debug_toolbar.urls)),  
]
```

7.8 Dokku

Readying a Django project for deploying to [Dokku](#).

This lists the things to add or change to easily deploy a Django application to Dokku.

It started out not trying to cover all of setting up a site on Dokku, only the parts relevant to a Django project – but it has grown. Still, you should read the [Dokku getting started docs](#), then use this as a cheatsheet to quickly enable existing Django projects to deploy on Dokku.

Start with the pages in this list, in order, then come back to this page and continue reading:

7.8.1 Dokku server administration

This page has information for those who have to set up and maintain a Dokku server. If you're just using one, you can ignore this.

Initial install

The Dokku docs recommend setting up a new OS install of a supported operating system, then running the Dokku install script.

Experience suggests that that approach is more likely to work than trying to install Dokku on a system that has already had some configuration done for other things.

Simple hostnames

The simple way to set up hostnames is:

- Pick a hostname you can control, e.g. `dokku.me`.
- During initial setup of Dokku, configure that as the server's name.
- Create a DNS A record pointing `dokku.me` at the server's IP address.
- Add a wildcard entry for `*.dokku.me` at the same address.
- For each app you put on that server, give the app the same name you want to use for its subdomain. For example, an app named `foo` would be accessible on the internet at `foo.dokku.me`, without having to make any more changes to your DNS settings.

Managing users

In other words, who can mess with the apps on a dokku server?

The way this currently works is that everybody ends up sshing to the server as the `dokku` user to do things. To let them do that, we want to add a public key for them to the dokku config, by doing this (from any system):

```
$ cat /path/to/ssh_keyfile.pub | ssh dokku ssh-keys:add <KEYNAME>
```

The <KEYNAME> is just to identify the different keys. I suggest using the person's typical username. Just remember there will not be a user of that name on the dokku server.

When it's time to revoke someone's access:

```
$ ssh dokku ssh_keys:remove <KEYNAME>
```

and now you see why the <KEYNAME> is useful.

For now, there's not a simple way to limit particular users to particular apps or commands.

7.8.2 Files

Setting up files in a Django project for deploying it to Dokku

requirements.txt

There needs to be a `requirements.txt` file at the top level. If you prefer to keep your requirements somewhere else, the top-level one can just look like:

```
-r path/to/real/requirements.txt
```

Wherever your requirements are, add the latest versions of:

```
dj-database-url
unicorn
whitenoise
```

settings

Add a `.../deploy.py` settings file, e.g. `<appname>/settings/deploy.py`.

It can start out looking like this (edit the top line if your main settings file isn't `base.py`):

```
# Settings when deployed to Dokku
from .base import * # noqa
import dj_database_url

# Disable Django's own staticfiles handling in favour of WhiteNoise, for
# greater consistency between gunicorn and `./manage.py runserver`. See:
# http://whitenoise.evans.io/en/stable/django.html#using-whitenoise-in-development
INSTALLED_APPS.remove('django.contrib.staticfiles')
INSTALLED_APPS.extend([
    'whitenoise.runserver_nostatic',
    'django.contrib.staticfiles',
])

MIDDLEWARE.remove('django.middleware.security.SecurityMiddleware')
MIDDLEWARE = [
    'django.middleware.security.SecurityMiddleware',
    'whitenoise.middleware.WhiteNoiseMiddleware',
] + MIDDLEWARE

# Update database configuration with $DATABASE_URL.
db_from_env = dj_database_url.config(conn_max_age=500)
```

(continues on next page)

(continued from previous page)

```
DATABASES['default'].update(db_from_env)

# Honor the 'X-Forwarded-Proto' header for request.is_secure()
SECURE_PROXY_SSL_HEADER = ('HTTP_X_FORWARDED_PROTO', 'https')

# Allow all host headers (feel free to make this more specific)
ALLOWED_HOSTS = ['*']

# Simplified static file serving.
# https://warehouse.python.org/project/whitenoise/
STATICFILES_STORAGE = 'whitenoise.storage.CompressedManifestStaticFilesStorage'
```

wsgi.py

Find your `wsgi.py` file.

1. Edit to change the default settings module to `<appname>.settings.deploy` (the path to the new settings file you created above).
2. Add to the end:

```
from whitenoise.django import DjangoWhiteNoise
application = DjangoWhiteNoise(application)
```

Procfile

Create [Procfile](#) (more on [dokku Procfile](#)) in the top directory. For our simple case, it can just contain one line, starting with `web:` and containing the command to start gunicorn for our site:

```
web: gunicorn {{ project_name }}.wsgi
```

See also the section on running Celery and other processes.

runtime.txt

Create `runtime.txt` in the top directory. It only needs one line, e.g.:

```
python-3.6.1
```

This *has* to be specific. E.g. `python-3.5.2` or `python-3.6.1` might work if the dokku server supports it, but `python-3.5` or `python-3.6` probably won't.

app.json

Create `app.json` in the top-level project directory. You might see examples on the Interwebs with lots of things in `app.json` (because Heroku uses `app.json` for lots of things), but as of this writing, dokku ignores everything but `scripts.dokku.predeploy` and `scripts.dokku.postdeploy`. Example:

```
{
  "scripts": {
    "dokku": {
```

(continues on next page)

(continued from previous page)

```

    "predeploy": "python manage.py migrate --noinput"
  }
}
}

```

Note: Dokku automatically runs `collectstatic` for you, so you don't need to do that from `app.json`.

buildpacks

If your app is not pure Python - e.g. if it uses node - you'll need to [override](#) the automatic buildpack detection, because it only works for a single application type.

Do this by adding a top-level `.buildpacks` file, containing links to the buildpacks to use:

```

https://github.com/heroku/heroku-buildpack-nodejs.git
https://github.com/heroku/heroku-buildpack-python.git
https://github.com/heroku/heroku-buildpack-apt

```

Heroku maintains a [list of buildpacks](#).

7.8.3 Postgres with Dokku

There's nothing Django-specific about this, but I'm including it just because we probably want to do it on every single Django deploy.

To install the [postgres plugin](#), inside your server run (because plugins must be installed as root):

```
$ sudo dokku plugin:install https://github.com/dokku/dokku-postgres.git
```

Now you need to create a database, and link the database to the app. You can do this from your own system:

```
$ ssh dokku postgres:create example-database
$ ssh dokku postgres:link example-database django-tutorial

```

Now when dokku runs your app, it'll set an env var to tell it where its DB is, e.g.:

```
DATABASE_URL=postgres://user:pass@host:port/db
```

For Django, install the tiny `dj_database_url` package, then in `settings.py`:

```

import dj_database_url
db_from_env = dj_database_url.config(conn_max_age=500)
DATABASES['default'].update(db_from_env)

```

There are built-in commands making it easy to backup and restore databases:

```

$ ssh dokku postgres:export [db_name] > [db_name].dump
$ ssh dokku postgres:import [db_name] < [db_name].dump

```

7.8.4 SSL for Dokku (Letsencrypt etc.)

Letsencrypt

Note: Get the site up and running, and accessible from the Internet, first. Let's Encrypt will not be able to get you a certificate until then.

There's nothing Django-specific about this part, but I'm including it just because we probably want to do it on every single Django deploy.

To add SSL with the [Let's Encrypt plugin \(more\)](#), first install the plugin by running on the dokku server (plugins must be installed as root):

```
$ sudo dokku plugin:install https://github.com/dokku/dokku-letsencrypt.git
```

Then on your system, to configure your app and tell letsencrypt to manage its certs and renew them periodically:

```
$ ssh dokku config:set --no-restart <appname> DOKKU_LETSENCRYPT_EMAIL=your@email.tld
$ ssh dokku letsencrypt myapp
$ ssh dokku letsencrypt:cron-job --add <appname>
```

Forcing SSL from Django

If we don't want to figure out how to override the default nginx config to redirect non-SSL requests to SSL, we can have Django do it with a few settings.

First, set `SECURE_SSL_REDIRECT` to `True`. This will tell Django to do the redirect.

```
SECURE_SSL_REDIRECT = True
```

Commit, deploy, and make sure the site still works.

Second, set `SECURE_HSTS_SECONDS` to a relatively small number of seconds.

```
SECURE_HSTS_SECONDS = 1800
```

This adds a header to all responses, telling any browser that receives it that this site should only be accessed via SSL, for that many seconds.

Commit, deploy, and make sure the site still works.

If everything still seems good, bite the bullet, increase `SECURE_HSTS_SECONDS` to a large number (e.g. 31536000 seconds, or 1 year), commit, deploy, and test again.

Additional info (move to their own files as they're ready)...

7.8.5 Environment variables

I can't find docs on what environment variables Dokku sets globally when running apps. But just about any Django app is going to be linked to a database, so if you need to detect whether the app is running under Dokku or Heroku, looking for `DATABASE_URL` should be good enough.

7.8.6 Static files

We use `whitenoise` to serve static files from Django. If the site gets incredible amounts of traffic, throw a CDN in front, but honestly, very few sites actually need that. (If you have a philosophical objection to serving static files from

Django, you can customize the nginx config through Dokku and probably manage to get nginx to do the static file serving, but I haven't bothered figuring it out myself.)

Or put your static files on S3.

7.8.7 Django media

If your site requires uploaded files to be persisted, remember that the container running your code is ephemeral and any changes made to files inside it will vanish at the next deploy.

First, you can use S3 for your media files.

Or you can use persistent storage on Dokku, which is a way of mounting directories from the dokku server inside the running container where your site code can store and read files that will continue to exist past the lifetime of that particular container.

7.8.8 Simple hostnames

If the server is set up properly, assuming the server's domain name is `dokkuserver.domain`, creating an app named `foo` will automatically `foo.dokkuserver.domain` resolve to the server's address, and requests with host `foo.dokkuserver.domain` to be routed to that app.

If you want to get fancier, <http://dokku.viewdocs.io/dokku/configuration/domains/>.

Also, note that any requests simply addressed to `dokku.me` will get routed to the alphabetically first app on the server, but you can change that: <http://dokku.viewdocs.io/dokku/configuration/domains/> or just set up a "00default" app.

7.8.9 Zero downtime deploys

WRITEME - see <http://dokku.viewdocs.io/dokku/deployment/zero-downtime-deploys/>

7.8.10 Behind a load balancer

If requests are being terminated at a load balancer and then proxied to our dokku server, some nginx config customization will be needed so your app can see the actual origin of the requests: <http://dokku.viewdocs.io/dokku/configuration/ssl/#running-behind-a-load-balancer>

7.8.11 Run a command

Suppose you want to run something like `python manage.py createsuperuser` in the app environment?

```
$ ssh dokku run <appname> python manage.py createsuperuser
```

will do it.

7.8.12 Running other daemons (like Celery)

Suppose you need to run another instance of your app in another way, for example to run `celery beat` and `celery worker`.

Use the Procfile to tell Dokku what processes to run. E.g. if your Procfile is:

```
web: gunicorn path/to/file.wsgi
```

try editing it to:

```
web: gunicorn path/to/file.wsgi
beat: celery beat -A appname -linfo
worker: celery worker -A appname -linfo
```

With just that, the extra processes won't run automatically. You can run them by telling Dokku to [scale them up](#), e.g.:

```
$ ssh dokku ps:scale <appname> beat=1 worker=4
```

You can check the current scaling settings:

```
$ ssh dokku ps:scale <appname>
-----> Scaling for <appname>
-----> proctype          qty
-----> -----          ---
-----> web                1
-----> beat                1
-----> worker              4
```

and see what's actually running (example from another app that only has one process):

```
$ ssh dokku ps:report <appname>
=====> <appname>
Processes:          1
Deployed:           true
Running:            true
Restore:            true
Restart policy:     on-failure:10
Status web.1       true      (CID: 03ea8977f37e)
```

Since we probably don't want to have to remember to manually scale these things up and check that they're running, we can add a `DOKKU_SCALE` file to our repo:

```
web=1
beat=1
worker=4
```

which is equivalent to running `ps:scale web=1 beat=1 worker=4`

7.8.13 Secrets

First, if possible, use Dokku plugin integrations for things like databases, Redis, cache, etc. They automatically set environment variables in each app that your settings can read, so you don't have to manage different settings for each environment. See each plugin's doc, of course, for more on that.

The way to handle other secrets for each environment is to set them as config on the app, which will add them as environment variables, again so your settings can read them at runtime.

The downside is that the secrets aren't being managed/versioned etc for us... but I think we can handle the few we'll need by manually keeping them in Lastpass or something.

[Here are the docs for setting config.](#)

Before setting any config, note that by default, making any change to the config triggers a new deploy. If you're not ready for that, include `--no-restart` in the command, as these examples will do.

To set config variables:

```
$ ssh dokku config:set <appname> --no-restart VAR1=VAL1 VAR2=VAL2 ...
```

To remove a variable:

```
$ ssh dokku config>unset <appname> --no-restart VAR1
```

Check the value of some variable:

```
$ ssh dokku config:get <appname> VAR1
VAL1
```

Get all the settings in a single line, handy for use in shell commands or to set on another app:

```
$ ssh dokku config <appname> --shell
VAR1=VAL1 VAR2=VAL2
$ export $(ssh dokku config <appname> --shell)
$ ssh dokku config:set <appname1> $(ssh dokku config <appname2> --shell)
```

Get all the settings in a format handy to save in a file for later sourcing in a local shell:

```
$ ssh dokku config <appname> --export
export VAR1=VAL1
export VAR2=VAL2
$ ssh dokku config <appname> --export >appname.env
$ . appname.env
```

Note: you can also set config vals globally - just change <appname> to --global in any of these commands.

7.8.14 Logs

Dokku collects stdout from your processes and you can view it with the `dokku logs` command.

`nginx` logs are similarly stored on the server and can be accessed using `dokku nginx:access-log <appname>` or `dokku nginx:error-log <appname>`.

Dokku event logging can be enabled with `dokku events:on` and then viewed with `dokku events`. This shows things like deploy steps.

7.8.15 Deploying from private git repos

Note: this doesn't apply to your main project repo. That can be private and Dokku doesn't care, because you're pushing it directly from your development system to the Dokku server.

But if your requirements include references to private git repos, then you'll need to arrange for Dokku to get access to those repos when it's `pip` installing your requirements.

Dokku docs, such as they are...

I think the upshot is:

- Create a new ssh key for deploying
- Add it to the repo on Github (or whatever) as an authorized deploy key (TODO: Link to github docs on that)
- Drop a copy of the public key file into `/home/dokku/.ssh/` on the Dokku system (with appropriate permissions)

7.8.16 Deploying non-master branch

The docs

By default, dokku deploys when the app's master branch on the dokku server is updated. There are (at least) two ways to deploy a branch other than master.

1. Push your non-master local branch to the master branch on the server:

```
$ git push dokku <local branch>:master
```

but that requires you to always remember that if you have apps that are always supposed to use a different branch than master.

2. Configure your app so the default branch is different, by using the `git:set` command:

```
$ ssh dokku git:set appname deploy-branch SOME_BRANCH_NAME
```

This seems like a more useful approach. Just “`git push dokku some-branch`” every time you want to deploy your app.

7.8.17 Developing with multiple remote apps

Suppose you have local development and sometimes you want to push to staging and sometimes to production. Maybe you have other apps too.

The key is to set up a different git remote for each remote app. E.g.:

```
$ ssh dokku app:create staging
$ ssh dokku git:set staging deploy-branch develop

$ git remote add staging dokku@my-dokku-server.com:staging

$ ssh dokku app:create production
$ ssh dokku git:set production deploy-branch master

$ git remote add production dokku@my-dokku-server.com:production
```

Then to deploy, push the appropriate branch to the appropriate branch:

```
$ git push staging develop
$ git push production master
```

7.8.18 Customizing nginx config

If you need to completely override the nginx config, you'll need to [provide an nginx config file template](#).

Luckily, much customization can be done just by [providing snippets](#) of configuration for nginx to include after it's base config file.

To do this, arrange for the snippets to get copied to `/home/dokku/<appname>/nginx.conf.d/` during deployment, probably in a pre- or post-deploy script.

7.8.19 Logging to papertrail

Use the [logspout](#) plugin.

7.8.20 Adding Sentry service

<https://github.com/darklow/dokku-sentry>

7.9 Elastic Beanstalk with Django

SSH into a random instance. This assumes that you have copied the SSH private key into your `$HOME/.ssh` directory:

```
(bringfido)$ eb ssh staging
```

Open the AWS Elasticbeanstalk web console:

```
(bringfido)$ eb console staging
```

Scale the application to N web instances:

```
(bringfido)$ eb scale <N> staging
```

Check the overall status of the environment, or detailed info about each instance:

```
(bringfido)$ eb status -v staging
(bringfido)$ eb health staging
```

If you need to work with Django on a server, after ssh'ing in:

```
$ . /opt/python/current/env
$ cd /opt/python/current/app
$ python manage.py ....
```

7.10 Email

<https://docs.djangoproject.com/en/stable/topics/email/>

7.10.1 API

Send one email:

```
send_mail(subject, message, from_email, recipient_list, fail_silently=False, auth_
↪user=None, auth_password=None, connection=None)
```

7.10.2 Attachments

```
msg = EmailMessage(...)
msg.attach(
    filename="any string",
    content=b"the contents",
    mimetype="application/sunshine"
)
```

or

```
msg.attach(instance of MIMEBase)
```

7.10.3 Email backends/handlers

<https://docs.djangoproject.com/en/stable/topics/email/#email-backends>

For development:

```
EMAIL_BACKEND = 'django.core.mail.backends.console.EmailBackend'
```

For real:

```
EMAIL_BACKEND = 'django.core.mail.backends.smtp.EmailBackend'
```

In-memory backend - The 'locmem' backend stores messages in a special attribute of the django.core.mail module. The outbox attribute is created when the first message is sent. It's a list with an EmailMessage instance for each message that would be sent:

```
EMAIL_BACKEND = 'django.core.mail.backends.locmem.EmailBackend'
```

This backend is not intended for use in production – it is provided as a convenience that can be used during development and testing.

7.10.4 Settings for email addresses

ADMINS A tuple that lists people who get code error notifications:

```
((('John', 'john@example.com'), ('Mary', 'mary@example.com')))
```

MANAGERS Not needed

DEFAULT_FROM_EMAIL Default email address to use for various automated correspondence from the site manager(s). This doesn't include error messages sent to ADMINS and MANAGERS; for that, see SERVER_EMAIL.

SERVER_EMAIL The email address that error messages come from, such as those sent to ADMINS and MANAGERS.

7.11 Fabric

Sample tasks:

```
@task
def db_backup():
    """
    Backup the database to S3 just like the nightly cron job
    """
    require('environment')
    require('instance', provided_by='instance')
    manage_run("dbbackup --encrypt")

def db_exists(dbname):
```

(continues on next page)

(continued from previous page)

```

"""
Return True if a db named DBNAME exists on the remote host.
"""
require('environment', provided_by=SERVER_ENVIRONMENTS)
output = sudo('psql -l --pset=format=unaligned', user='postgres')
dbnames = [line.split('|')[0] for line in output.splitlines()]
return dbname in dbnames

@task
def db_dump(file):
    """
    Dump an instance's database to a remote file.

    Example:

    `fab staging instance:iraq db_dump:/tmp/staging_iraq.dump`

    dumps to staging_iraq.dump
    """
    require('environment', provided_by=SERVER_ENVIRONMENTS)
    require('instance', provided_by='instance')
    remote_file = file

    if files.exists(file):
        if not confirm("Remote file {file} exists and will be overwritten. Okay?"
                        .format(file=remote_file)):
            abort("ERROR: aborting")

    # Don't need remote DB user and password because we're going to run pg_dump as_
    ↪ user postgres
    sudo('pg_dump --format=custom --file={outputfile} {dbname}'
        .format(dbname=env.db_name, outputfile=remote_file),
        user='postgres')
    print("Database from {environment} {instance} has been dumped to remote file
    ↪ {file}"
        .format(environment=env.environment, instance=env.instance, file=remote_
    ↪ file))

@task
def local_restore(file):
    """
    Restore a local dump file to the local instance's database.
    :param file:
    :return:
    """
    # Find out the local DB settings
    import sys
    sys.path[0:0] = ['.']
    from cts.settings.local import DATABASES
    DB = DATABASES['default']
    assert DB['ENGINE'] == 'django.contrib.gis.db.backends.postgis'
    dbname = DB['NAME']
    owner = DB['USER'] or os.getenv('USER')
    local('dropdb {dbname} || true'.format(dbname=dbname), shell="/bin/sh")
    local('createdb --encoding UTF8 --lc-collate=en_US.UTF-8 --lc-ctype=en_US.UTF-8 --
    ↪ template=template0 --owner {owner} {dbname}'.format(owner=owner, dbname=dbname))

```

(continues on next page)

(continued from previous page)

```

    local('sudo -u postgres pg_restore -Ox -j4 --dbname={dbname} {file}').
↪format(dbname=dbname, file=file))

@task
def db_restore(file):
    """
    Restore a remote DB dump file to a remote instance's database.

    This will rename the existing database to {previous_name}_bak
    and create a completely new database with what's in the dump.

    If there's already a backup database, the restore will fail.

    Example:

        `fab staging instance:iraq db_restore:/tmp/staging_iraq.dump`

    :param file: The remote file to restore.
    """
    require('environment', provided_by=SERVER_ENVIRONMENTS)
    require('instance', provided_by='instance')

    renamed = False
    restored = False

    if not files.exists(file):
        abort("Remote file {file} does not exist".format(file=file))

    try:
        if db_exists(env.db_name):
            # Rename existing DB to backup
            db_backup = '{dbname}_bak'.format(dbname=env.db_name)
            if db_exists(db_backup):
                if confirm("There's already a database named {db_backup}. Replace_
↪with new backup?")
                    .format(db_backup=db_backup)):
                    sudo('dropdb {db_backup}'.format(db_backup=db_backup),
                        user='postgres')
                else:
                    abort("ERROR: There's already a database named {db_backup}. "
                        "Restoring would clobber it."
                        .format(db_backup=db_backup))
            sudo('psql -c "ALTER DATABASE {dbname} RENAME TO {db_backup}"'
                .format(dbname=env.db_name, db_backup=db_backup),
                user='postgres')
            renamed = True
            print("Renamed {dbname} to {db_backup}".format(dbname=env.db_name, db_
↪backup=db_backup))

            remote_file = file

            # Create new, very empty database.
            # * We can't use --create on the pg_restore because that will always restore_
↪to whatever
            # db name was saved in the dump file, and we don't want to be restricted_
↪that way.

```

(continues on next page)

(continued from previous page)

```

# * Any extensions the backed-up database had will be included in the restore,
↳ so we
#   don't need to enable them now.

# If these parameters change, also change the parameters in conf/salt/project/
↳ db/init.sls
# (TODO: we could use the output of psql -l to copy most of these settings_
↳ from the
# existing database.)
sudo('createdb --encoding UTF8 --lc-collate=en_US.UTF-8 '
    '--lc-ctype=en_US.UTF-8 --template=template0 --owner {owner} {dbname}'
    .format(dbname=env.db_name, owner=env.db_owner),
    user='postgres')

# Don't need remote DB user and password because we're going to
# run pg_restore as user postgres
sudo('pg_restore -l --dbname={dbname} {filename}'
    .format(dbname=env.db_name, filename=remote_file),
    user='postgres')
restored = True

# Run ANALYZE on the db to help Postgres optimize how it accesses it
sudo('psql {dbname} -c ANALYZE' .format(dbname=env.db_name),
    user='postgres')

print("Database for {environment} {instance} has been restored from remote_
↳ file {file}")
    .format(environment=env.environment, instance=env.instance, file=remote_
↳ file))
finally:
    if renamed and not restored:
        print("Error occurred after renaming current database, trying to rename_
↳ it back")
        if db_exists(env.db_name):
            # We already created the new db, but restore failed; delete it
            sudo('dropdb {dbname}' .format(dbname=env.dbname), user='postgres')
            sudo('psql -c "ALTER DATABASE {db_backup} RENAME TO {dbname}"'
                .format(dbname=env.db_name, db_backup=db_backup),
                user='postgres')
            print("Successfully put back the original database.")

```

7.12 Filtering and Pagination with Django

If you want to build a list page that allows filtering and pagination, you have to get a few separate things to work together. Django provides some tools for pagination, but the documentation doesn't tell us how to make that work with anything else. Similarly, `django_filter` makes it relatively easy to add filters to a view, but doesn't tell you how to add pagination (or other things) without breaking the filtering.

The heart of the problem is that both features use query parameters, and we need to find a way to let each feature control its own query parameters without breaking the other one.

7.12.1 Filters

Let's start with a review of filtering, with an example of how you might subclass `ListView` to add filtering. To make it filter the way you want, you need to create a subclass of `FilterSet` and set `filterset_class` to that class. (See that link for how to write a `filterset`.)

```
class FilteredListView(ListView):
    filterset_class = None

    def get_queryset(self):
        # Get the queryset however you usually would. For example:
        queryset = super().get_queryset()
        # Then use the query parameters and the queryset to
        # instantiate a filterset and save it as an attribute
        # on the view instance for later.
        self.filterset = self.filterset_class(self.request.GET, queryset=queryset)
        # Return the filtered queryset
        return self.filterset.qs.distinct()

    def get_context_data(self, **kwargs):
        context = super().get_context_data(**kwargs)
        # Pass the filterset to the template - it provides the form.
        context['filterset'] = self.filterset
        return context
```

Here's an example of how you might create a concrete view to use it:

```
class BookListView(FilteredListView):
    filterset_class = BookFilterset
```

And here's part of the template that uses a form created by the `filterset` to let the user control the filtering.

```
<h1>Books</h1>
<form action="" method="get">
    {{ filterset.form.as_p }}
    <input type="submit" />
</form>

<ul>
    {% for object in object_list %}
        <li>{{ object }}</li>
    {% endfor %}
</ul>
```

`filterset.form` is a form that controls the filtering, so we just render that however we want and add a way to submit it.

That's all you need to make a simple filtered view.

7.12.2 Default values for filters

I'm going to digress slightly here, and show a way to give filters default values, so when a user loads a page initially, for example, the items will be sorted with the most recent first. I couldn't find anything about this in the `django_filter` documentation, and it took me a while to figure out a good solution.

To do this, I override `__init__` on my filter set and add default values to the data being passed:


```
class BookFilterSet(django_filters.FilterSet):
    def __init__(self, data, *args, **kwargs):
        data = data.copy()
        data.setdefault('format', 'paperback')
        data.setdefault('order', '-added')
        super().__init__(data, *args, **kwargs)
```

I tried some other approaches, but this seemed to work out the simplest, in that it didn't break or complicate things anywhere else.

7.12.3 Pagination

Now let's review pagination in Django.

Django's `ListView` has some built-in support for pagination, which is easy enough to enable:

```
class BookListView(FilteredListView):
    paginate_by = 50
```

Once `paginate_by` is set to the number of items you want per page, `object_list` will contain only the items on the current page, and there will be some additional items in the context:

paginator A `Paginator` object

page_obj A `Page` object

is_paginated True if there are pages

We need to update the template so the user can control the pages.

Let's start our template updates by just telling the user where we are:

```
{% if is_paginated %}
Page {{ page_obj.number }} of {{ paginator.num_pages }}
{% endif %}
```

To tell the view which page to display, we want to add a query parameter named `page` whose value is a page number. In the simplest case, we can just make a link with `?page=N`, e.g.:

```
<a href="?page=2">Goto page 2</a>
```

You can use the `page_obj` and `paginator` objects to build a full set of pagination links, but there's a problem we should solve first.

7.12.4 Combining filtering and pagination

Unfortunately, linking to pages as described above breaks filtering. More specifically, whenever you follow one of those links, the view will forget whatever filtering the user has applied, because that filtering is also controlled by query parameters, and these links don't include the filter's parameters.

So if you're on a page `https://example.com/objectlist/?type=paperback` and then follow a page link, you'll end up at `https://example.com/objectlist/?page=3` when you wanted to be at `https://example.com/objectlist/?type=paperback&page=3`.

It would be nice if Django helped out with a way to build links that set one query parameter without losing the existing ones, but I found a nice example of a template tag on [StackOverflow](#) and modified it slightly into this custom template tag that helps with that:

```
# <app>/templatetags/my_tags.py
from django import template

register = template.Library()

@register.simple_tag(takes_context=True)
def param_replace(context, **kwargs):
    """
    Return encoded URL parameters that are the same as the current
    request's parameters, only with the specified GET parameters added or changed.

    It also removes any empty parameters to keep things neat,
    so you can remove a parm by setting it to ``""`.

    For example, if you're on the page ``/things/?with_frosting=true&page=5``,
    then

    <a href="/things/?{% param_replace page=3 %}">Page 3</a>

    would expand to

    <a href="/things/?with_frosting=true&page=3">Page 3</a>

    Based on
    https://stackoverflow.com/questions/22734695/next-and-before-links-for-a-django-
    paginated-query/22735278#22735278
    """
    d = context['request'].GET.copy()
    for k, v in kwargs.items():
        d[k] = v
    for k in [k for k, v in d.items() if not v]:
        del d[k]
    return d.urlencode()
```

Here's how you can use that template tag to build pagination links that preserve other query parameters used for things like filtering:

```
{% load my_tags %}

{% if is_paginated %}
    {% if page_obj.has_previous %}
        <a href="?{% param_replace page=1 %}">First</a>
        {% if page_obj.previous_page_number != 1 %}
            <a href="?{% param_replace page=page_obj.previous_page_number %}">Previous</a>
        {% endif %}
    {% endif %}

    Page {{ page_obj.number }} of {{ paginator.num_pages }}

    {% if page_obj.has_next %}
        {% if page_obj.next_page_number != paginator.num_pages %}
            <a href="?{% param_replace page=page_obj.next_page_number %}">Next</a>
        {% endif %}
        <a href="?{% param_replace page=paginator.num_pages %}">Last</a>
    {% endif %}
```

(continues on next page)

(continued from previous page)

```
<p>Objects {{ page_obj.start_index }}&mdash;{{ page_obj.end_index }}</p>
{% endif %}
```

Now, if you're on a page like <https://example.com/objectlist/?type=paperback&page=3>, the links will look like [?type=paperback&page=2](#), [?type=paperback&page=4](#), etc.

7.12.5 Useful links

- [django_filter](#)
- [Django pagination](#)
- [param_replace](#) template tag

I haven't tried it, but if you need something more sophisticated for building these kinds of links, [django-qurl-templatetag](#) might be worth looking at.

7.13 Forms

- <https://docs.djangoproject.com/en/stable/topics/forms/>
- <https://docs.djangoproject.com/en/stable/ref/forms/api/>

A basic form:

```
from django import forms

class ContactForm(forms.Form):
    subject = forms.CharField(max_length=100)
```

Processing a form in a view function:

```
def contact(request):
    if request.method == 'POST': # If the form has been submitted...
        form = ContactForm(request.POST) # A form bound to the POST data
        if form.is_valid(): # All validation rules pass
            # Process the data in form.cleaned_data
            # ...
            return HttpResponseRedirect('/thanks/') # Redirect after POST
    else:
        form = ContactForm() # An unbound form

    return render_to_response('contact.html', {
        'form': form,
    })
```

<https://docs.djangoproject.com/en/stable/topics/forms/modelforms/>

A model form:

```
from django.forms import ModelForm, ValidationError

# Create the form class.
class ArticleForm(ModelForm):
    class Meta:
```

(continues on next page)

(continued from previous page)

```

model = Article
fields = ('name', 'title')
# or
exclude = ('birth_date')

def clean_fieldname(self):
    if 'fieldname' in self.cleaned_data:
        data = self.cleaned_data['fieldname']
        if not /valid/:
            raise ValidationError("msg")
        return data

def clean(self):
    data = self.cleaned_data
    if not /valid/:
        raise ValidationError("msg")
    return data

# Creating a form to add an article.
form = ArticleForm()
...
new_article = form.save()

# Creating a form to change an existing article.
article = Article.objects.get(pk=1)
form = ArticleForm(instance=article)
...
form.save()

# Create a form to edit an existing Article, but use POST data to populate the form.
a = Article.objects.get(pk=1)
f = ArticleForm(request.POST, instance=a)
f.save()

```

Render form in template:

```

<!-- Using table - avoid that part - but this does show how to render the fields_
->individually -->
    <form {% if form.is_multipart %}enctype="multipart/form-data"{% endif %} action=
->" " method="post" class="uniForm">{% csrf_token %}
    <table>
        <fieldset>
            {% if form.non_field_errors %}
                <tr><td colspan="2">{{ form.non_field_errors }}</td></tr>
            {% endif %}
            {% for field in form %}
                <tr{% if field.field.required %} class="required"{% endif %}>
                    <th style="text-align: left"><label for="{{ field.id_for_label }}">{{
->field.label }}:</label></th>
                    <td>{% if field.errors %}
                        {{ field.errors }}<br/>
                    {% endif %}

                    {{ field }}

                    or even

```

(continues on next page)

(continued from previous page)

```

        <input id="{{ field.id_for_label }}"
              name="{{ field.html_name }}"
              value="{{ field.value }}"
              {% if field.field.max_length != None %}
                maxlength="{{ field.field.max_length }}"
              {% endif %}
              {% if field.field.min_length != None %}
                minlength="{{ field.field.min_length }}"
              {% endif %}
              >
        {% if field.help_text %}
        <br/><span class="helptext">{{ field.help_text }}</span>
        {% endif %}
    </td>
</tr>
{% endfor %}
</fieldset>
</table>
<div class="ctrlHolder buttonHolder">
    <button type="submit" class="primaryAction" name="submit_changes">Submit_
↪ changes</button>
</div>
</form>

<!-- Using a list, which is preferred -->

    <form {% if form.is_multipart %}enctype="multipart/form-data"{% endif %} action="
↪ " method="post" class="uniForm">{% csrf_token %}
        <fieldset>
            <ul>
                {{ form.as_ul }}
                <li>
                    <div class="ctrlHolder buttonHolder">
                        <button type="submit" class="primaryAction" name="submit_
↪ changes">Submit changes</button>
                    </div>
                </li>
            </ul>
        </fieldset>
    </form>

```

7.14 Read-only form

Call this on the form:

```

def make_form_readonly(form):
    """
    Set some attributes on a form's fields that, IN COMBINATION WITH TEMPLATE CHANGES,
    allow us to display it as read-only.
    """

    # Note that a new BoundField is constructed on the fly when you access
    # form[name], so any data we want to persist long enough for the template
    # to access needs to be on the "real" field. We just use the BoundField

```

(continues on next page)

(continued from previous page)

```
# to get at the field value.

for name in form.fields:
    field = form.fields[name]
    bound_field = form[name]
    if hasattr(field.widget, 'choices'):
        try:
            display_value = dict(field.widget.choices)[bound_field.value()]
        except KeyError:
            display_value = ''
    else:
        display_value = bound_field.value()

    field.readonly = True
    field.display_value = display_value
```

Do things like this in the templates:

```
{# Date field #}
{% if field.field.readonly %}
    <span class="form-control">{{ field.value|date:'c' }}</span>
{% else %}
    <input type="date" class="form-control" id="{{ field.id_for_label }}" name="{{ field.html_name }}" value="{{ field.value|date:'c' }}">
{% endif %}

{# input fields #}
{% if field.field.readonly %}
    <span class="form-control">{{ field.value }}</span>
{% else %}
    <input type="{{ block input_field_type }}" text{% endblock %}" class="form-control" id="{{ field.id_for_label }}" name="{{ field.html_name }}" value="{{ field.value }}" placeholder="{{ field.field.widget.attrs.placeholder }}" {% endif %} {% block input_attrs %}{% endblock %}>
{% endif %}

{# select fields #}
{% if field.field.readonly %}
    <span class="form-control">{{ field.field.display_value }}</span>
{% else %}
    <select class="form-control" id="{{ field.id_for_label }}" name="{{ field.html_name }}" placeholder="">
        {% for val, label in field.field.widget.choices %}
            <option value="{{ val }}" {% if field.value|stringformat:'s' == val|stringformat:'s' %} selected{% endif %}>{{ label }}</option>
        {% endfor %}
    </select>
{% endif %}
```

7.15 Quick And Dirty Home Page

In `urls.py`:

```
from django.views.generic import TemplateView

urlpatterns = [
    ...
    url(r'^$', TemplateView.as_view(template_name='home.html'), name='home'),
]
```

7.16 Logging

Some best practices for Django logging.

<https://docs.djangoproject.com/en/stable/topics/logging/#configuring-logging> <https://docs.djangoproject.com/en/stable/topics/logging/>

<https://docs.python.org/2/howto/logging.html>

<https://docs.python.org/2/library/logging.handlers.html#rotatingfilehandler> <https://docs.python.org/2/library/logging.handlers.html#timedrotatingfilehandler>

INFO level logging for celery is very verbose

If you have DEBUG on, Django logs all SQL queries

7.16.1 Default

Here's what Django uses (around 1.7, anyway) if you don't configure logging:

```
# Default logging for Django. This sends an email to the site admins on every
# HTTP 500 error. Depending on DEBUG, all other log records are either sent to
# the console (DEBUG=True) or discarded by mean of the NullHandler (DEBUG=False).
LOGGING = {
    'version': 1,
    'disable_existing_loggers': False,
    'filters': {
        'require_debug_false': {
            '()': 'django.utils.log.RequireDebugFalse',
        },
        'require_debug_true': {
            '()': 'django.utils.log.RequireDebugTrue',
        },
    },
    'handlers': {
        'console': {
            'level': 'INFO',
            'filters': ['require_debug_true'],
            'class': 'logging.StreamHandler',
        },
        'null': {
            'class': 'logging.NullHandler',
        },
        'mail_admins': {
            'level': 'ERROR',
            'filters': ['require_debug_false'],
            'class': 'django.utils.log.AdminEmailHandler'
        }
    }
}
```

(continues on next page)

(continued from previous page)

```
    },
    'loggers': {
        'django': {
            'handlers': ['console'],
        },
        'django.request': {
            'handlers': ['mail_admins'],
            'level': 'ERROR',
            'propagate': False,
        },
        'django.security': {
            'handlers': ['mail_admins'],
            'level': 'ERROR',
            'propagate': False,
        },
        'py.warnings': {
            'handlers': ['console'],
        },
    },
}
```

7.16.2 Console

Log errors to console in local.py:

```
LOGGING.setdefault('formatters', {})
LOGGING['formatters']['verbose'] = {
    'format': '[%(name)s] Message "%(message)s" from %(pathname)s:%(lineno)d in
    ↪ %(funcName)s'
}
LOGGING['handlers']['console'] = {
    'class': 'logging.StreamHandler',
    'formatter': 'verbose',
    'level': 'ERROR',
}
LOGGING['loggers']['django'] = {
    'handlers': ['console'],
    'level': 'ERROR',
    'propagate': True,
}
```

7.16.3 Development

For local development, we want lots of output saved to a log file in case we need to look back at a problem, but no emailing of exceptions and such.

In settings:

```
LOG_DIR = os.path.join(PROJECT_ROOT, '..', 'log')

LOGGING = {
    'version': 1,
    'disable_existing_loggers': False,
```

(continues on next page)

(continued from previous page)

```

'handlers': {
    'console': { # Log to stdout
        'level': 'INFO',
        'class': 'logging.StreamHandler',
    },
    'file': {
        'level': 'DEBUG',
        'class': 'logging.FileHandler',
        'filename': os.path.join(LOG_DIR, 'django_debug.log',
    },
},
'root': { # For dev, show errors + some info in the console
    'handlers': ['console'],
    'level': 'INFO',
},
'loggers': {
    'django.request': { # debug logging of things that break requests
        'handlers': ['file'],
        'level': 'DEBUG',
        'propagate': True,
    },
},
}

```

Or how about:

```

LOGGING = {
    'version': 1,
    'disable_existing_loggers': False,
    'formatters': {
        'simple': {
            'format': '%(name)-20s %(levelname)-8s %(message)s',
        },
    },
    'handlers': {
        'console': { # Log to stdout
            'level': 'INFO',
            'class': 'logging.StreamHandler',
            'formatter': 'simple',
        },
    },
    'root': { # For dev, show errors + some info in the console
        'handlers': ['console'],
        'level': 'INFO',
    },
}

```

7.16.4 Staging

FIXME: Add celery exceptions

@tobiasmcnulty also mentioned: “re: celery error emails, this is a good setting to have enabled: <http://celery.readthedocs.org/en/latest/configuration.html#celery-send-task-error-emails>”

On staging, we still want lots of info logged semi-permanently (to files), but we also want to be emailed about exceptions to make sure we find out about problems before we deploy them to production.

Emails should go to the devs, not the client or production site admins.

Like so:

```
ADMINS = (
    ('XXX DevTeam', 'xxx-dev-team@example.com'),
)

LOG_DIR = os.path.join(PROJECT_ROOT, '..', 'log')

LOGGING = {
    'version': 1,
    'disable_existing_loggers': False,
    'handlers': {
        'file': { # Rotate log file daily, only keep 1 backup
            'level': 'DEBUG',
            'class': 'logging.handlers.TimedRotatingFileHandler',
            'filename': os.path.join(LOG_DIR, 'django_debug.log',
            'when': 'd',
            'interval': 1,
            'backupCount': 1,
        },
        'mail_admins': {
            'level': 'ERROR',
            'class': 'django.utils.log.AdminEmailHandler'
        },
    },
    # EMAIL all errors (might not want this, but let's try it)
    'root': {
        'handlers': ['mail_admins'],
        'level': 'ERROR',
    },
    'loggers': {
        'django.request': {
            'handlers': ['file'],
            'level': 'INFO',
            'propagate': True,
        },
    },
}
```

7.16.5 Production

Mark says: for production I like to log to syslog which can then be shipped elsewhere without changing the application (<https://docs.python.org/2/library/logging.handlers.html#logging.handlers.SysLogHandler> ?)

@Scottm and I have been talking about making that more common: log to syslog, ship to Logstash, monitor via Kibana <http://www.elasticsearch.org/overview/kibana/>

getting Nginx to log to syslog is kind of a pain you basically have to get syslog to monitor the file and ship it Logstash + Kibana looks much easier to manage/configure than Graylog2

the plan was to add it to Ona but that isn't done yet (as of Aug 28, 2014) CCSR was/is using Graylog2 Minidam does syslog -> Loggly libya is using logstash -> graylog (in addition to sentry)

7.16.6 Example

Here's what we've got set up for Django logging on one project. This sends everything level INFO and higher to a local log file and a Graylog instance. Anything ERROR and higher is emailed to admins and sent to a Sentry instance, which can send more notifications.

In environment:

```
SENTRY_DSN: http://long_hex_string:long_hex_string@hostname:9000/3
```

Requirements:

```
raven==3.6.1
```

Settings:

```
INSTALLED_APPS = (
    ...
    'raven.contrib.django.raven_compat', # Sentry logging client
    ...
)

CELERY_SEND_TASK_ERROR_EMAILS = True

# Send ERRORS to email and sentry.
# Send a fair bit of info to graylog and a local log file
# (but not debug level messages, ordinarily).
LOGGING = {
    'version': 1,
    'disable_existing_loggers': True,
    'filters': {
        # This filter strips out request information from the message record
        # so it can be sent to Graylog (the request object is not picklable).
        'django_exc': {
            '()': 'our_filters.RequestFilter',
        },
        'require_debug_false': {
            '()': 'django.utils.log.RequireDebugFalse'
        },
        # This filter adds some identifying information to each message, to make
        # it easier to filter them further, e.g. in Graylog.
        'static_fields': {
            '()': 'our_filters.StaticFieldFilter',
            'fields': {
                'deployment': 'project_name',
                'environment': 'staging' # can be overridden, e.g. 'staging' or
↪ 'production'
            },
        },
    },
    'formatters': {
        'basic': {
            'format': '%(asctime)s %(name)-20s %(levelname)-8s %(message)s',
        },
    },
    'handlers': {
        'file': {
            'level': 'DEBUG', # Nothing here logs DEBUG level messages ordinarily

```

(continues on next page)

(continued from previous page)

```

        'class': 'logging.handlers.RotatingFileHandler',
        'formatter': 'basic',
        'filename': os.path.join(LOG_ROOT, 'django.log'),
        'maxBytes': 10 * 1024 * 1024, # 10 MB
        'backupCount': 10,
    },
    'graylog': {
        'level': 'INFO',
        'class': 'graypy.GELFHandler',
        'host': env_or_default('GRAYLOG_HOST', 'monitor.caktusgroup.com'),
        'port': 12201,
        'filters': ['static_fields', 'django_exc'],
    },
    'mail_admins': {
        'level': 'ERROR',
        'class': 'django.utils.log.AdminEmailHandler',
        'include_html': False,
        'filters': ['require_debug_false'],
    },
    'sentry': {
        'level': 'ERROR',
        'class': 'raven.contrib.django.raven_compat.handlers.SentryHandler',
    },
},
'root': {
    # graylog (or any handler using the 'django_exc' filter ) should be last
    # because it will alter the LogRecord by removing the `request` field
    'handlers': ['file', 'mail_admins', 'sentry', 'graylog'],
    'level': 'WARNING',
},
'loggers': {
    # These 2 loggers must be specified, otherwise they get disabled
    # because they are specified by django's DEFAULT_LOGGING and then
    # disabled by our 'disable_existing_loggers' setting above.
    # BEGIN required loggers #
    'django': {
        'handlers': [],
        'propagate': True,
    },
    'py.warnings': {
        'handlers': [],
        'propagate': True,
    },
    # END required loggers #
    # The root logger will log anything WARNING and higher, so there's
    # no reason to add loggers here except to add logging of lower-level
    →information.
    'libya_elections': {
        'handlers': ['file', 'graylog'],
        'level': 'INFO',
    },
    'nlid': {
        'handlers': ['file', 'graylog'],
        'level': 'INFO',
    },
    'register': {
        'handlers': ['file', 'graylog'],

```

(continues on next page)

(continued from previous page)

```

        'level': 'INFO',
    },
    'bulk_sms': {
        'handlers': ['file', 'graylog'],
        'level': 'INFO',
    },
}

#
# our_filters.py
#
import logging

class QuotelessStr(str):
    """
    Return the repr() of this string *without* quotes. This is a
    temporary fix until https://github.com/severb/graypy/pull/34 is resolved.
    """
    def __repr__(self):
        return self

class StaticFieldFilter(logging.Filter):
    """
    Python logging filter that adds the given static contextual information
    in the ``fields`` dictionary to all logging records.
    """
    def __init__(self, fields):
        self.static_fields = fields

    def filter(self, record):
        for k, v in self.static_fields.items():
            setattr(record, k, QuotelessStr(v))
        return True

class RequestFilter(logging.Filter):
    """
    Python logging filter that removes the (non-pickable) Django ``request``
    object from the logging record.
    """
    def filter(self, record):
        if hasattr(record, 'request'):
            del record.request
        return True

```

7.16.7 Including info like the emailed errors do

```

from django.views.debug import TECHNICAL_500_TEXT_TEMPLATE, get_safe_settings, \
    get_exception_reporter_filter
from django.views.decorators.debug import sensitive_post_parameters

t = Template(TECHNICAL_500_TEXT_TEMPLATE)

```

(continues on next page)

(continued from previous page)

```
filter = get_exception_reporter_filter(request)
r = t.render(Context({
    'request': request,
    'is_email': True,
    'filtered_POST': filter.get_post_parameters(request),
    'settings': get_safe_settings(),
    'server_time': timezone.now(),
    'django_version_info': get_version(),
}, autoescape=False))
logger.error(
    "Got CSRF failure, reason=%s. %s", reason, r,
)
```

7.17 Django login and logout

Django doc

- In settings.py, add:

```
LOGIN_URL = '/login/'
```

- In urls.py, add:

```
urlpatterns += patterns('',
    (r'^login/$', django.contrib.auth.views.login),
    (r'^logout/$', django.contrib.auth.views.logout),
)
```

- Create a template “registration/login.html”, copying from the [sample in the doc](#)
- Add a logout link to your base template:

```
<a href="/logout/">Logout</a>
```

- On each view function where users should be logged in before using, add the decorator:

```
@login_required
def myview(...)
```

7.18 Middleware

7.18.1 Middleware ordering

Middleware ordering docs

Re: whitenoise middleware: The WhiteNoise middleware should be placed directly after the Django SecurityMiddleware (if you are using it) and before all other middleware:

Mostly copy of the Django middleware docs (2.1):

Here are some hints about the ordering of various Django middleware classes:

1. `SecurityMiddleware`

It should go near the top of the list if you're going to turn on the SSL redirect as that avoids running through a bunch of other unnecessary middleware.

2. `UpdateCacheMiddleware`

Before those that modify the Vary header (`SessionMiddleware`, `GZipMiddleware`, `LocaleMiddleware`).

3. `GZipMiddleware`

Before any middleware that may change or use the response body.

After `UpdateCacheMiddleware`: Modifies Vary header.

4. `SessionMiddleware`

After `UpdateCacheMiddleware`: Modifies Vary header.

5. `ConditionalGetMiddleware`

Before any middleware that may change the response (it sets the ETag header).

After `GZipMiddleware` so it won't calculate an ETag header on gzipped contents.

6. `LocaleMiddleware`

One of the topmost, after `SessionMiddleware` (uses session data) and `UpdateCacheMiddleware` (modifies Vary header).

7. `CommonMiddleware`

Before any middleware that may change the response (it sets the Content-Length header). A middleware that appears before `CommonMiddleware` and changes the response must reset Content-Length.

Close to the top: it redirects when `APPEND_SLASH` or `PREPEND_WWW` are set to True.

8. `CsrfViewMiddleware`

Before any view middleware that assumes that CSRF attacks have been dealt with.

It must come after `SessionMiddleware` if you're using `CSRF_USE_SESSIONS`.

9. `AuthenticationMiddleware`

After `SessionMiddleware`: uses session storage.

10. `MessageMiddleware`

After `SessionMiddleware`: can use session-based storage.

11. `FetchFromCacheMiddleware`

After any middleware that modifies the Vary header: that header is used to pick a value for the cache hash-key.

12. `FlatpageFallbackMiddleware`

Should be near the bottom as it's a last-resort type of middleware.

13. `RedirectFallbackMiddleware`

Should be near the bottom as it's a last-resort type of middleware.

7.19 Migrations

Data migration:

```
def no_op(apps, schema_editor):
    pass

def create_types(apps, schema_editor):
    ServiceType = apps.get_model('services', 'ServiceType')
    db_alias = schema_editor.connection.alias
    # do stuff
    ServiceType.objects.using(db_alias)....

class Migration(migrations.Migration):
    ...
    operations = [
        migrations.RunPython(create_types, no_op),
    ]
```

7.19.1 Getting past bad migrations

For example, earlier migrations refer to models in apps that no longer exist.

The simplest thing is to start from an existing database so you don't have to migrate.

If you need to start from scratch, you should be able to:

```
syncdb --all
migrate --fake
```

7.20 NGINX

Some tips on Nginx for use with Django

7.20.1 Files

Just add a new file to `/etc/nginx/sites-enabled` for each site, making sure `server_name` is set correctly in each.

7.20.2 Redirecting to SSL

We usually want to force SSL:

```
server {
    listen *:80;
    listen [::]:80;
    server_name DOMAIN;
    access_log PATHo_access.log;
    error_log PATH_error.log;
    return 301 https://DOMAIN$request_uri;
}
```


7.20.3 Proxying to unicorn

Serve static and media files with nginx, and proxy everything else to Django:

```
upstream django {
    server unix:/tmp/PATH fail_timeout=0;
}

server {
    listen *:443 ssl;    # add spdy here too if you want
    listen [::]:443 ssl;
    server_name DOMAIN;
    ssl_certificate PATH.crt;
    ssl_certificate_key PATH.key;

    access_log PATH_access.log;
    error_log PATH_error.log;
    root PATH;
    location /media {
        alias PATH;
    }
    location /static {
        alias PATH;
    }
    location / {
        client_max_body_size 500M;
        proxy_set_header X-Real-IP $remote_addr;
        proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
        proxy_set_header X-Forwarded-Proto $scheme;
        proxy_set_header Host $host;
        proxy_redirect off;
        proxy_buffering on;
        proxy_intercept_errors on;
        proxy_pass http://django;
    }

    # See https://www.trevorparkar.com/hardening-ssl-in-nginx/
    ssl_protocols          TLSv1 TLSv1.1 TLSv1.2;
    ssl_prefer_server_ciphers on;
    ssl_ciphers             DHE-RSA-AES128-GCM-SHA256:DHE-DSS-AES128-GCM-SHA256:ECDE-
↪RSA-AES1\
28-GCM-SHA256:ECDE-ECDSA-AES128-GCM-SHA256:ECDE-RSA-AES256-GCM-SHA384:ECDE-ECDSA-
↪AES256-GCM\
-SHA384:kEDH+AESGCM:DHE-RSA-AES128-SHA256:DHE-RSA-AES128-SHA:DHE-DSS-AES128-
↪SHA256:DHE-RSA-AES\
256-SHA256:DHE-DSS-AES256-SHA:DHE-RSA-AES256-SHA:ECDE-RSA-AES128-SHA256:ECDE-ECDSA-
↪AES128-SH\
A256:ECDE-RSA-AES128-SHA:ECDE-ECDSA-AES128-SHA:ECDE-RSA-AES256-SHA384:ECDE-ECDSA-
↪AES256-SH\
A384:ECDE-RSA-AES256-SHA:ECDE-ECDSA-AES256-SHA:AES128-GCM-SHA256:AES256-GCM-
↪SHA384:AES128-SH\
A256:AES256-SHA256:AES128-SHA:AES256-SHA:AES:CAMELLIA:DES-CBC3-SHA:!aNULL:!eNULL:!
↪EXPORT:!DES:\
!RC4:!MD5:!PSK:!aECDH:!EDH-DSS-DES-CBC3-SHA:!EDH-RSA-DES-CBC3-SHA:!KRB5-DES-CBC3-SHA;
    ssl_session_timeout    5m;
    ssl_session_cache       shared:SSL:10m;

    add_header Strict-Transport-Security max-age=31536000;
```

(continues on next page)

(continued from previous page)

```
}
```

7.21 Permissions

(Note: this page is about authorization, not authentication.)

Link: <https://docs.djangoproject.com/en/stable/topics/auth/default/#topic-authorization>

User objects have two many-to-many fields: `groups` and `user_permissions`. User objects can access their related objects in the same way as any other Django model:

```
myuser.groups = [group_list]
myuser.groups.add(group, group, ...)
myuser.groups.remove(group, group, ...)
myuser.groups.clear()
myuser.user_permissions = [permission_list]
myuser.user_permissions.add(permission, permission, ...)
myuser.user_permissions.remove(permission, permission, ...)
myuser.user_permissions.clear()
```

Assuming you have an application with an `app_label` `foo` and a model named `Bar`, to test for basic permissions you should use:

- add: `user.has_perm('foo.add_bar')`
- change: `user.has_perm('foo.change_bar')`
- delete: `user.has_perm('foo.delete_bar')`

Note: There is no default permission for read access.

The Permission model is rarely accessed directly, but here it is:

```
permission = Permission.objects.create(codename='can_publish',
                                      name='Can Publish Posts', # verbose human-
↳ readable name
                                      content_type=content_type)
```

Permissions are more commonly referred to by a string, of the form “`app_label.codename`”. E.g., if `user.has_perm('myapp.codename')`: do something.

Confusingly, the `Permission` model has no `app_label` field. It uses `content_type__app_label` for that.

Warning: This means all permissions need a content type, whether it makes sense for that permission to be applied to a particular model or not.

To create a new permission programmatically:

```
content_type = ContentType.objects.get_for_model(BlogPost)
permission = Permission.objects.create(codename='can_publish',
                                      name='Can Publish Posts',
                                      content_type=content_type)
```

7.21.1 Default permissions

<https://docs.djangoproject.com/en/stable/topics/auth/default/#default-permissions>

For every model in an installed app, Django automatically creates three permissions: *app_label.add_modelname*, *app_label.change_modelname*, and *app_label.delete_modelname*, where the modelname is lowercased.

7.21.2 Adding model permissions

<https://docs.djangoproject.com/en/stable/ref/models/options/#permissions>

You can ask Django to create more permissions for a model:

```
class Meta:
    permissions = [
        ('codename', 'verbose name'),
    ]
```

When the table is created during `syncdb`, Django will create the additional Permission objects too.

In Django 1.11 (and probably earlier, but definitely not before 1.7), if you edit these permissions and `makemigrations`, Django will create a migration for you that when run, will add any missing migrations. (I don't know whether it'll update verbose names of existing permissions.)

You can programmatically force Django to create additional Permissions with code like:

```
from django.db.models import get_models, get_app
from django.contrib.auth.management import create_permissions

apps = set([get_app(model._meta.app_label) for model in get_models()])
for app in apps:
    create_permissions(app, None, 2)
```

7.21.3 Best practices

- Plan not to give users specific permissions, except when you have to make an exception to your usual policies.
- Design groups with useful sets of permissions.
- Plan to add users to the appropriate groups depending on their roles.
- Provide a way to ensure the groups continue to have the permissions you want.

Fixtures aren't a bad way to provide initial data, but setting them up for automatic loading is deprecated with Django 1.7 and will go away with Django 2.0. Instead, load them from a data migration. This is better in some ways anyway, because the migration will use the same version of the models that the fixtures were written for at the time. (Though, this doesn't matter so much for Permissions and Groups, which we don't really expect to change their schemas...)

Add utility methods like this, maybe in *accounts/utlis.py* or equivalent:

```
def permission_names_to_objects(names):
    """
    Given an iterable of permission names (e.g. 'app_label.add_model'),
    return an iterable of Permission objects for them. The permission
    must already exist, because a permission name is not enough information
    to create a new permission.
    """
```

(continues on next page)

(continued from previous page)

```

result = []
for name in names:
    app_label, codename = name.split(".", 1)
    # Is that enough to be unique? Hope so
    try:
        result.append(Permission.objects.get(content_type__app_label=app_label,
                                             codename=codename))
    except Permission.DoesNotExist:
        logger.exception("NO SUCH PERMISSION: %s, %s" % (app_label, codename))
        raise
return result

def get_all_perm_names_for_group(group):
    # Return the set of permission names that the group should contain

def create_or_update_groups():
    for group_name, perm_names in GROUP_PERMISSIONS.iteritems():
        group, created = Group.objects.get_or_create(name=group_name)
        perms_to_add = permission_names_to_objects(get_all_perm_names_for_
→group(group))
        group.permissions.add(*perms_to_add)
        if not created:
            # Group already existed - make sure it doesn't have any perms we didn't_
→want
            to_remove = set(group.permissions.all()) - set(perms_to_add)
            if to_remove:
                group.permissions.remove(*to_remove)

```

7.21.4 Checking permissions in templates

<https://docs.djangoproject.com/en/stable/topics/auth/default/#authentication-data-in-templates>

```

{% if user.is_authenticated %} {% if perms.applabel %} {%# user has any permissions in app applabel #}
{% if 'applabel' in perms %} {%# same as above %} {% if perms.applabel.change_thing %} {%# user has
'change_thing' permission in app applabel #} {% if 'applabel.change_thing' in perms %} {%# same as
above #}

```

7.22 Queries and Querysets

7.22.1 Field lookups

exact, iexact, contains, icontains, startswith, istartswith, endswith, iendswith, in, gt, gte, lt, lte, range, year, month, day, week_day, hour, minute, second, isnull, search, regex

7.22.2 Optimizing Django queries on big data

Suppose you have a query that needs to run against a table or tables with many millions of rows. Maybe you need to operate on a couple million of them. What are the do's and don't's of a Django query that will not pessimize performance (time and memory use)?

- Don't bother with `.iterator()`, it downloads the whole result and then iterates over it. It does not do what many of us think/thought it did (use a database cursor to pull down the results only as you work through them)
- Do limit the query (`[start:end]`) and run it repeatedly in reasonable sized batches, to avoid downloading too big a chunk
- Do use `.only()` and its kin to minimize how much of each record is downloaded to what you need
- Do not `order_by()` unless you must - it forces the DB to collect the results of the whole query first so it can sort them, even if you then limit the results you retrieve
- Same for `.distinct()`.

7.22.3 The model that a queryset is over

`queryset.model`

7.22.4 Combining querysets

Given two querysets over the same model, you can do things like this:

```
queryset = queryset1 & queryset2
queryset = queryset1 | queryset2
queryset = queryset1 & ~queryset2
```

(similar to `Q` objects)

7.22.5 Custom QuerySets

Calling custom `QuerySet` methods from the manager

Creating a manager with `QuerySet` methods¶:

```
class Person(models.Model):
    ...
    people = PersonQuerySet.as_manager()

class BaseManager(...):
    ....

class MyModel(models.Model):
    objects = BaseManager.from_queryset(CustomQuerySet)()
```

7.22.6 Custom Lookups

Adding to the kwargs you can pass to *filter* and *exclude* etc.

Custom Lookups

7.23 Security

7.23.1 Protect internal services

Use a VPN, or check out [oauth2_proxy](#) or similar services.

7.23.2 Django

(django-secure appears to be abandoned. Last change was in 2014, and it doesn't load under Django 1.11/Python 3.6.)

-Best practice: install `django-secure` and run `manage.py checksecure` to make sure all the right settings are enabled.-

See also [OWASP](#).

7.23.3 Admin

Don't leave it externally accessible, even with a password.

7.23.4 SSH

Two important settings in `/etc/sshd_config`:

- Disable root login:

```
PermitRootLogin no
```

- Disable password auth:

```
PasswordAuthentication no
```

Also consider changing to some port other than 22.

7.23.5 SSL

SEE ALSO [NGINX](#) and [Django docs on SSL and https](#).

Basically, make sure nginx is setting X-Forwarded-Proto, then add to settings:

```
SECURE_PROXY_SSL_HEADER = ('HTTP_X_FORWARDED_PROTO', 'https')
```

Django security

djangocon 2011 Paul McMillan

<http://subversivecode.com/talks/djangocon-us-2011>

HSTS

django needs better password hash (SHA1 not broken but very fast)

OpenID much more secure against password cracking (because cracker won't have passwords)

password reset strings can be eventually worked out with a timing attack (if you have a long time and a fast connection)

same for which userids exist on the site

you should do rate limiting:

`mod_evasive` (apache) `HttpLimitReqModule` (nginx)

do NOT use `random.Random()` for security functions, not cryptographically secure; use `random.SystemRandom()` instead e.g.:

```
from random import SystemRandom as random
xxxx random.choice(yyyy)...
```

Be very careful with pickle, it'll execute anything in the pickled data when you unpickle it

BOOK: The web application hacker's handbook (new version coming out soon (as of 9/8/2011))

SITE: lost.org? (not sure I heard that right)(no I didn't)

7.24 Sentry on a Django site

Requirements:

```
raven==6.6.0 # or whatever
```

Settings:

```
INSTALLED_APPS += ('raven.contrib.django.raven_compat',)

LOGGING[handlers]['sentry'] = {
    'level': 'ERROR', # To capture more than ERROR, change to WARNING, INFO, etc.
    'class': 'raven.contrib.django.raven_compat.handlers.SentryHandler',
    # 'tags': {'custom-tag': 'x'},
}

LOGGING['root']['handlers'].append('sentry')
# OR
LOGGING['root'] = {
    'level': 'WARNING',
    'handlers': ['sentry'],
}

RAVEN_CONFIG = {
    'dsn': '{{ RAVEN_DSN }}',
    'release': '{{ commit }}',
    'site': 'TypeCoach',
    'environment': '{{ env }}',
    'processors': [
        'raven.processors.SanitizePasswordsProcessor',
    ]
}
```

Base template:

```
{% load raven %}
<!doctype html>
<head>
    ...
    <script src="https://cdn.ravenjs.com/3.23.2/raven.min.js" crossorigin="anonymous">
    ↪</script>
    <script>Raven.config('{{ sentry_public_dsn }}').install()</script>
```

wsgi.py:

```
from raven.contrib.django.raven_compat.middleware.wsgi import Sentry
from django.core.wsgi import get_wsgi_application

application = Sentry(get_wsgi_application())
```

7.25 Settings

7.25.1 Using the right one

Baumgartner suggests symlinking the desired one (e.g. *dev.py* or *deploy.py*) to *local.py* and hard-coding that in *manage.py*.

Greenfelds suggest... (FILL THIS IN)

12-factor says there should only be one settings file, and any values that vary by deploy should be pulled from the environment. See *Env vars*.

7.25.2 Secret key

Generate a secret key:

```
from django.utils.crypto import get_random_string
chars = 'abcdefghijklmnopqrstuvwxyz0123456789!@#%^&*(_+=) '
SECRET_KEY = get_random_string(50, chars)
```

(<https://github.com/django/django/blob/master/django/core/management/commands/startproject.py#L26>)

7.25.3 Env vars

Suppose you have env vars in a .env file:

```
SECRET_KEY=jdfsdfsdf
PASSWORD=jsdkfjsdlkfjsdf
```

You can load them into Django using *dotenv*. Pop open *manage.py*. Add:

```
import dotenv
dotenv.read_dotenv()
```

Or in a settings file:

```
SECRET_KEY = os.environ.get("SECRET_KEY")
```

And if they're not all strings, use ast:

```
import ast, os

DEBUG = ast.literal_eval(os.environ.get("DEBUG", "True"))
TEMPLATE_DIRS = ast.literal_eval(os.environ.get("TEMPLATE_DIRS", "/path1,/path2"))
```

You can load them into a shell this way:


```
export $(cat .env | grep -v ^# | xargs)
```

7.26 Templates

7.26.1 Template tag to set variables

Example usage:

```
{% set foo="bar" a=1 %}
...
Hello everyone, foo is {{ foo }} and a is {{ a }}.
```

Here's the code:

```
from django import template

register = template.Library()

@register.simple_tag(takes_context=True)
def set(context, **kwargs):
    context.update(kwargs)
    return ''
```

7.26.2 Working block tag with arguments

Here's an example of a working block tag. Usage is `{% detail_link arg1=1 arg2=2 %}...{% end_detail_link %}` and what ends up in the output is `<a arg1="1" arg2="2" target="_blank">...`.

The code:

```
from django import template
from django.template.base import token_kwargs, TemplateSyntaxError
from django.utils.html import format_html, format_html_join
from django.utils.safestring import mark_safe

register = template.Library()

def do_detail_link(parser, token):
    """
    Block tag to help render links to detail pages consistently
    with an option to open in a new tab or window.

    {% detail_link href="xxxx" arg1="yyy" ... %} what to display {% end_detail_link %}

    is rendered as

    <a href="xxxx" arg1="yyy" ... target="_blank" %} what to display </a>

    This adds `target="_blank"` to open the link in a new tab or window.
    That's the main purpose of this block tag (and so we can disable that in
    one place, here, if we ever want to). But you can override it by specifying
```

(continues on next page)

(continued from previous page)

```

another value for `target` if you want.
"""
# This is called each time a detail_link tag is encountered while parsing
# a template.

# Split the contents of the tag itself
args = token.split_contents() # e.g. ["detail_link", "arg1='foo'", "arg2=bar"]
tag_name = args.pop(0)

# Parse out the arg1=foo arg2=bar ... arguments from the arg list into a
→dictionary.
# kwargs will have "arg1" etc as keys, while the values will be
# template thingies that can later be rendered using different contexts
# to get their value at different times.
kwargs = token_kwargs(args, parser)

# If there are any args left, we have a problem; this tag only
# accepts kwargs.
if args:
    raise TemplateSyntaxError("%r only accepts named kwargs" % tag_name)

# Open in new tab unless otherwise told (by setting target to something else).
if 'target' not in kwargs:
    kwargs['target'] = parser.compile_filter("_blank")

# Parse inside of block *until* we're looking at {% end_detail_link %},
# then we don't care about end_detail_link, so discard it.
# When we return, the parsing will then continue after our end tag.
odelist = parser.parse(('end_detail_link',))
parser.delete_first_token()

# Now return a node for the parsed template
return DetailLinkNode(odolist, tag_name, kwargs)

register.tag('detail_link', do_detail_link)

class DetailLinkNode(template.Node):
    """
    Stores info about one occurrence of detail_link in a template.

    See also `do_detail_link`.
    """
    def __init__(self, odolist, tag_name, kwargs):
        self.odolist = odolist
        self.tag_name = tag_name
        self.kwargs = kwargs

    def render(self, context):
        """Turn this node into text using the given context."""

        # Start with the part inside the block
        innerds = self.odolist.render(context)

        # Now work out the <a> wrapper.
        args = format_html_join(

```

(continues on next page)

(continued from previous page)

```

        ' ',
        ' {}=" {}"',
        ((name, value.resolve(context)) for name, value in self.kwargs.items())
    )
    result = format_html(
        mark_safe("<a {}>{}</a>"),
        args,
        mark_safe(innerds)
    )
    return result

```

7.26.3 Debugging template syntax errors during tests

The normal error message when a view fails rendering a template during testing gives no clue where the error is.

You can get a better idea by temporarily editing your local Django installation. Find the file `django/template/base.py`. Around line 194 (in Django 1.8.x), in the `__init__` method of the `Template` class, look for this code:

```
self.nodelist = engine.compile_string(template_string, origin)
```

and change it to:

```

try:
    self.nodelist = engine.compile_string(template_string, origin)
except TemplateSyntaxError:
    print("ERROR COMPILING %r" % origin.name)
    raise

```

TODO: would be nice to get a line number too (this just gives a filename, which is often enough in combination with the error message).

7.27 Testing

```

self.assertEqual(a, b, msg=None)

rsp = self.client.get(url, [follow=True])
rsp = self.client.post(url, data, [follow=True])

rsp.content is a byte string
rsp['HeaderName']
rsp.context['template_var']

assert self.client.login(**login_parms)
login(email='foo@example.com', password='cleartextpassword')

```

<http://docs.python.org/library/unittest.html> # <https://docs.djangoproject.com/en/stable/topics/testing/>

```

from django.test import TestCase
from django.contrib.auth.models import User

```

(continues on next page)

(continued from previous page)

```
class XxxTest(TestCase):
    def setUp(self):
        self.user = User.objects.create_user('tester', 'test@example.com', 'testpass')

    def test_something(self):
        response = self.client.get(show_timemap)
        self.assertEqual(response.status_code, 200)
        self.assertEqual(response.context['lat'], '123.123')
        self.assertEqual(response.context['lon'], '456.456')
        self.assertContains(response, "some text")
        self.assertNotContains(response, "other text")
        self.assertIsNone(val)
        self.assertIsNotNone(val)
        self.assertIn(thing, iterable) # Python >= 2.7
        self.assertNotIn(thing, iterable) # Python >= 2.7
```

Test uploading a file

(from <https://docs.djangoproject.com/en/stable/topics/testing/tools/#django.test.Client.post>)

Submitting files is a special case. To POST a file, you need only provide the file field name as a key, and a file handle to the file you wish to upload as a value. For example:

```
>>> c = Client()
>>> with open('wishlist.doc') as fp:
...     c.post('/customers/wishes/', {'name': 'fred',
...                                   'attachment': fp})
```

(The name attachment here is not relevant; use whatever name your file-processing code expects.)

Note that if you wish to use the same file handle for multiple post() calls then you will need to manually reset the file pointer between posts. The easiest way to do this is to manually close the file after it has been provided to post(), as demonstrated above.

You should also ensure that the file is opened in a way that allows the data to be read. If your file contains binary data such as an image, this means you will need to open the file in rb (read binary) mode.

7.27.1 Writing a test for a separately-distributed Django app

setup.py:

```
...
setup(
    ...
    test_suite="runtests.runtests",
    ...
)
```

runtests.py:

```
#!/usr/bin/env python
import os
import sys

from django.conf import settings
```

(continues on next page)

(continued from previous page)

```

if not settings.configured:
    settings.configure(
        DATABASES={
            'default': {
                'ENGINE': 'django.db.backends.sqlite3',
                'NAME': ':memory:',
            }
        },
        INSTALLED_APPS=(
            'selectable',
        ),
        SITE_ID=1,
        SECRET_KEY='super-secret',
        ROOT_URLCONF='selectable.tests.urls',
    )

from django.test.utils import get_runner

def runtests():
    TestRunner = get_runner(settings)
    test_runner = TestRunner(verbosity=1, interactive=True, failfast=False)
    args = sys.argv[1:] or ['selectable', ]
    failures = test_runner.run_tests(args)
    sys.exit(failures)

if __name__ == '__main__':
    runtests()

```

7.28 Translation

Switch context to a new language:

```

from django.utils import translation
translation.activate('en-us')

```

([link to Using translations outside views and templates](#))

7.29 URLs

7.29.1 reverse

```

from django.core.urlresolvers import reverse

reverse(viewname='myview',    # remaining args optional
        urlconf=None,
        args=(1, 2),

```

(continues on next page)

(continued from previous page)

```
kwargs={'pk': foo.pk},
current_app=None)
```

7.29.2 redirect

```
from django.shortcuts import redirect
```

redirect (*to* [, *permanent=False*], **args*, ***kwargs*)

Returns an `HttpResponseRedirect` to the appropriate URL for the arguments passed.

The arguments could be:

- A model: the model's `get_absolute_url()` function will be called.
- A view name, possibly with arguments: `urlresolvers.reverse()` will be used to reverse-resolve the name.
- A URL, which will be used as-is for the redirect location.

By default issues a temporary redirect; pass `permanent=True` to issue a permanent redirect

Examples

You can use the `redirect()` function in a number of ways.

1. By passing some object; that object's `get_absolute_url()` method will be called to figure out the redirect URL:

```
def my_view(request):
    ...
    object = MyModel.objects.get(...)
    return redirect(object)
```

2. By passing the name of a view and optionally some positional or keyword arguments; the URL will be reverse resolved using the `reverse()` method:

```
def my_view(request):
    ...
    return redirect('some-view-name', foo='bar')
```

3. By passing a hardcoded URL to redirect to:

```
def my_view(request):
    ...
    return redirect('/some/url/')
```

This also works with full URLs:

```
def my_view(request):
    ...
    return redirect('http://example.com/')
```

By default, `redirect()` returns a temporary redirect. All of the above forms accept a `permanent` argument; if set to `True` a permanent redirect will be returned:

```
def my_view(request):
    ...
    object = MyModel.objects.get(...)
    return redirect(object, permanent=True)
```

<https://docs.djangoproject.com/en/stable/topics/http/urls/>

```
from django.conf.urls.defaults import patterns, include, url
from django.contrib import admin admin.autodiscover()

urlpatterns = patterns('',
    (r'^polls/', include('polls.urls')),
    url(r'^admin/', include(admin.site.urls)),
)
urlpatterns = patterns('polls.views',
    (r'^$', 'index'),
    (r'^(?P<poll_id>\d+)/', 'detail'),
    (r'^(?P<poll_id>\d+)/results/', 'results'),
    (r'^(?P<poll_id>\d+)/vote/', 'vote'),
    url(r'^(?P<poll_id>\d+)/foo/', 'fooview', name='app-viewname'),
)
```

7.30 Django with Vue

Also Dokku...

If you want to serve your Django app, your Django app's static files, and the files constituting your vue app, all from your Django process, using [whitenoise](#), this seems to work.

Note: this doesn't work as well locally; see end.

7.30.1 Set up Django to gather static files from dist

In settings:

```
STATICFILES_DIRS = [
    ...
    os.path.join(BASE_DIR, 'dist'),
    ...
]
```

7.30.2 Configure whitenoise

Do the basic Django configuration, but add a few things:

```
WHITENOISE_INDEX_FILE = True
WHITENOISE_ROOT = os.path.join(STATIC_ROOT, 'vue')
```

`WHITENOISE_INDEX_FILE` tells whitenoise to serve `index.html` for `/` when it seems appropriate.

`WHITENOISE_ROOT` tells whitenoise to serve files at the root URL if it finds them in `WHITENOISE_ROOT`. E.g. if a request comes in for `/app.js` and there's a file `<WHITENOISE_ROOT>/app.js`, then Whitenoise will handle the request and return that file.

7.30.3 At deploy

Build your vue app for production and have the resulting files put under dist/vue:

```
mkdir -p dist/vue  
yarn run build --dest=dist/vue
```

(that runs vue-cli-service build).

Run collectstatic. It'll gather the files from dist along with other static files:

```
python manage.py collectstatic --noinput
```

Now 'runserver' or whatever.

7.30.4 Running locally

This prebuilds the Vue stuff, so if you're doing development and would like things to rebuild when you edit files, this won't do that.

Misc to file:

Avoid circular imports:

```
from django.db.models import get_model  
MyModel = get_model('applabel', 'mymodelname'.lower())
```


8.1 Terminology

Container: A container is used to execute an image. See `docker run`, `docker-compose up`.

Image: A kind of snapshot of a computer system that can be run in a container. Images are built from Dockerfiles. See `docker build`, `docker-compose build`.

8.2 Common commands

Build an image from Dockerfile in current directory:

```
docker build -t imagetag .
```

Start a container from an image, run a command, and when the command exits, stop the container without saving any changes:

```
docker run --rm -it imagetag bash
```

8.3 Cleanup

8.3.1 delete volumes

see: <https://github.com/chadoe/docker-cleanup-volumes>

```
$ docker volume rm $(docker volume ls -qf dangling=true)
$ docker volume ls -qf dangling=true | xargs -r docker volume rm
```

8.3.2 delete networks

```
$ docker network ls $ docker network ls | grep "bridge" $ docker network rm $(docker network ls | grep "bridge" | awk '/ / { print $1 }')
```

8.3.3 remove docker images

see: <http://stackoverflow.com/questions/32723111/how-to-remove-old-and-unused-docker-images>

```
$ docker images
$ docker rmi $(docker images --filter "dangling=true" -q --no-trunc)

$ docker images | grep "none"
$ docker rmi $(docker images | grep "none" | awk '/ / { print $3 }')
```

8.3.4 remove docker containers

see: <http://stackoverflow.com/questions/32723111/how-to-remove-old-and-unused-docker-images>

```
$ docker ps
$ docker ps -a
$ docker rm $(docker ps -qa --no-trunc --filter "status=exited")
```

8.3.5 Resize disk space for docker vm

```
$ docker-machine create --driver virtualbox --virtualbox-disk-size "40000" default
```

9.1 Documentation

- [Elasticsearch Reference](#)
- [Elasticsearch Definitive Guide](#)

Contents:

9.1.1 Indices

Create an index

[Create Index API](#)

Example args:

```
{
  "settings" : {
    "number_of_shards" : 3,
    "number_of_replicas" : 2
  }
}

{
  "settings" : {
    "number_of_shards" : 1
  },
  "mappings" : {
    "type1" : {
      "_source" : { "enabled" : false },
      "properties" : {
        "field1" : { "type" : "string", "index" : "not_analyzed" }
      }
    }
  }
}
```

(continues on next page)

(continued from previous page)

```

    }
  }
}

```

Index settings

TBD

dynamic [Dynamic mapping](#) Values `true`, `false`, `"strict"`.

9.1.2 Mappings

[Mappings API](#)

[Mappings reference](#)

[Intro to Mappings in the Guide](#)

[Types and Mappings in the Guide](#)

Example mapping for a `tweet` doctype:

```

{
  "tweet" : {
    "properties" : {
      "message" : {
        "type" : "string",
        "store" : true,
        "index" : "analyzed",
        "null_value" : "na"
      },
      "user" : {
        "type" : "string",
        "index" : "not_analyzed",
        "norms" : {
          "enabled" : false
        }
      },
      "postDate" : {"type" : "date"},
      "priority" : {"type" : "integer"},
      "rank" : {"type" : "float"}
    }
  }
}

```

The string type

The string type is analyzed as full text by default.

[String type reference](#) includes all the possible attributes.

The number type

[Number type reference](#)

The date type

[Date type reference](#)

The boolean type

[Boolean type reference](#) and it's boolean not Boolean.

Multi fields

[Multi fields reference](#)

You can store the same source field in several index fields, analyzed differently.

Object type

[Object Type Ref](#)

You can define a *field* to contain other fields.

Root object type

[Root object type ref](#)

[Root object in guide](#) - better

The top-level doc type in a mapping is also an object type, but has some special characteristics. For example, you can set the default analyzers that fields without an explicit analyzer will use.

You can also turn off dynamic mapping for a doctype, though the Root object type ref doesn't mention this. See [Dynamic mapping](#). You can even turn it back on for one field. Example:

```
{
  "my_type": {
    "dynamic": "strict",
    "properties": {
      ...
      "stash": {
        "type": "object",
        "dynamic": True
      }
    }
  }
}
```

9.1.3 Analysis

[Analysis reference](#)

[Analysis and Analyzers in the Guide](#)

An analyzer consists of:

- Zero or more `CharFilters`
- One `Tokenizer`
- Zero or more `TokenFilters`

Analysis can be configured when creating an index with the top-level `analysis` key in the API argument. Example:

```
analysis :
  analyzer :
    standard :
      type : standard
      stopwords : [stop1, stop2]
    myAnalyzer1 :
      type : standard
      stopwords : [stop1, stop2, stop3]
      max_token_length : 500
      # configure a custom analyzer which is
      # exactly like the default standard analyzer
    myAnalyzer2 :
      tokenizer : standard
      filter : [standard, lowercase, stop]
  tokenizer :
    myTokenizer1 :
      type : standard
      max_token_length : 900
    myTokenizer2 :
      type : keyword
      buffer_size : 512
  filter :
    myTokenFilter1 :
      type : stop
      stopwords : [stop1, stop2, stop3, stop4]
    myTokenFilter2 :
      type : length
      min : 0
      max : 2000
```

Built-in analyzers

[Built-in analyzers in the Guide](#)

Custom analyzers

You can define custom analyzers.

[Custom Analyzers in the Guide](#)

Example:

```
analysis :
  analyzer :
    myAnalyzer2 :
      type : custom
      tokenizer : myTokenizer1
      filter : [myTokenFilter1, myTokenFilter2]
      char_filter : [my_html]
      position_offset_gap: 256
  tokenizer :
    myTokenizer1 :
      type : standard
      max_token_length : 900
  filter :
    myTokenFilter1 :
      type : stop
      stopwords : [stop1, stop2, stop3, stop4]
    myTokenFilter2 :
      type : length
      min : 0
      max : 2000
  char_filter :
    my_html :
      type : html_strip
      escaped_tags : [xxx, yyy]
      read_ahead : 1024
```


10.1 Basic types

Examples:

```
iex> 1           # integer
iex> 0x1F        # integer (31)
iex> 0x1010      # integer (10)
iex> 0o777       # integer (511)
iex> 1.0         # float
iex> 1.0e-10     # float
iex> true        # boolean
iex> false       # boolean
iex> :atom       # atom / symbol
iex> "elixir"    # string
iex> [1, 2, 3]   # list
iex> {1, 2, 3}   # tuple
iex> is_atom(false) # true
iex> is_atom(:false) # true
iex> is_boolean(true) # true
iex> is_boolean(:true) # true
iex> is_integer(1.0) # false
iex> is_float("foo")
iex> is_number(1.0) # true
```

10.2 Arithmetic

Infix: `40 + 2` is 42. `/` on integers returns a float, e.g. `10 / 2` is 5.0:

```
10 / 2      5.0
div(10, 2)   5
```

(continues on next page)

(continued from previous page)

```
div 10, 2      5
rem 10, 3      1
round(3.58)    4
trunc(3.58)    3
```

10.3 Boolean expressions

```
true == false is false.
```

```
true != false is true.
```

You can also use `or`, `and`, and `not`. `or` and `and` are short-circuiting operators. All three of these require Boolean arguments.

Elixir also has `||`, `&&`, and `!`, which accept values of any type, and where any value except `false` and `nil` are truthy.

For comparison: `==`, `!=`, `===`, `!==`, `<=`, `>=`, `<`, and `>`. The only difference between `==` and `===` is that `1 == 1.0` is true but `1 === 1.0` is not true.

Elixir allows applying the comparison operators to values of any type or combination of types, for simplicity when doing things like sorting. There is an ordering among types, e.g. `<number> < <atom>`.

10.4 Strings

Double quoted literals are strings (single quoted literals are char lists, not strings).

Use `<>` to concatenate, e.g. `"hello" <> " world"`

You can interpolate:

```
iex> "hellö #{:world}"
"hellö world"
```

The typical `\x` char codes work, e.g. `"hello\nworld"`

Internally strings are binary, sequences of bytes (UTF-8):

```
iex> String.length("hellö")
5
iex> is_binary("hellö")
true
iex> byte_size("hellö")
6
```

The `String` module has lots of useful methods like `upcase`.

You can print a string using `IO.puts/1`:

```
iex> IO.puts "hello\nworld"
hello
world
:ok
```

Note that the return value of `IO.puts` is `:ok`.

10.5 Functions

Note: Functions in Elixir are identified by name and by number of arguments (i.e. arity). Therefore, `is_boolean/1` identifies a function named `is_boolean` that takes 1 argument. `is_boolean/2` identifies a different (nonexistent) function with the same name but different arity.:

```
iex> is_boolean(true)
true
iex> is_boolean(1)
false
```

You can get help on a function with `h` and its name/arity:

```
iex> h is_boolean
iex> h is_boolean/1
iex> h ==/2
```

BUT you can't call a function using its full name and arity, you have to leave off the arity:

```
iex> is_boolean/1(1)
** (SyntaxError) iex:2: syntax error before: '('
```

10.6 Anonymous functions

Define anonymous functions with `fn`, `->`, and `end`:

```
iex> add = fn a, b -> a + b end
...
iex> is_function(add)
true
iex> is_function(add, 2)
true
iex> is_function(add, 1)
false
```

Anonymous functions require a dot `.` to invoke:

```
iex> add.(1, 2)
3
```

Anonymous functions are closures and can access variables that were in scope when they were defined.

Variables assigned inside a function do *not* affect the surrounding environment, though:

```
iex> x = 42
42
iex> (fn -> x = 0 end).()
0
iex> x
42
```

10.7 Lists

Literal lists are written with square brackets. Values can be a mix of any types:

```
iex> length [1, 2, true, 3]
4
```

Lists are concatenated using ++/2 and can be “subtracted” using --/2:

```
iex> [1, 2, 3] ++ [4, 5, 6]
[1, 2, 3, 4, 5, 6]
iex> [1, true, 2, false, 3, true] -- [true, false]
[1, 2, 3, true]
```

The “head” of a list is like Lisp’s car but is accessed using the hd/1 function. Similarly, the “tail” is the cdr and you get it with tl/1:

```
iex> list = [1,2,3]
iex> hd(list)
1
iex> tl(list)
[2, 3]
```

You can add a new head to a list with |:

```
iex> [1 | [2, 3]]
[1, 2, 3]
```

Getting the head or the tail of an empty list is an error:

```
iex> hd []
** (ArgumentError) argument error
```

A list of small integers is printed by Elixir as a single-quoted “string” - but it’s not a string, it’s a list of chars:

```
iex> [11, 12, 13]
'\v\f\r'
iex> [104, 101, 108, 108, 111]
'hello'
```

10.8 Introspection

Use i/1 to get information about a value:

```
iex(2)> i 'hello'
Term
  'hello'
Data type
  List
Description
  This is a list of integers that is printed as a sequence of characters
  delimited by single quotes because all the integers in it represent valid
  ASCII characters. Conventionally, such lists of integers are referred to as
  "char lists" (more precisely, a char list is a list of Unicode codepoints,
```

(continues on next page)

(continued from previous page)

```

    and ASCII is a subset of Unicode).
Raw representation
  [104, 101, 108, 108, 111]
Reference modules
  List
iex(3)> i "hello"
Term
  "hello"
Data type
  BitString
Byte size
  5
Description
  This is a string: a UTF-8 encoded binary. It's printed surrounded by
  "double quotes" because all UTF-8 encoded codepoints in it are printable.
Raw representation
  <<104, 101, 108, 108, 111>>
Reference modules
  String, :binary

```

10.9 Tuples

Literal tuples are written with curly brackets {1, :ok, true}. Access any element with elem/2 using 0-indexing, get the length with tuple_size/1, and return a new tuple with an element changed using put_elem/3:

```

iex> elem({:ok, "hello"})
"hello"
iex> tuple_size({:ok, "hello"})
2
iex> put_elem({:ok, "hello"}, 1, "world")
{:ok, "world"}

```


11.1 Disaster recovery

`git reflog` will list all the recent commits, whether they're reachable from any branch or tag or not. Find the one you want, check it out by its commit, and then you can make that a branch with `git branch <new-name>` (I think).>

11.2 Undoing things

If you've committed some changes, then for some reason decide you didn't want to commit them yet - but still want the changes present in your local working directory - there are several options.

To get rid of the actual commit but keep all those changes staged:

```
$ git reset --soft HEAD~
```

To get rid of the actual commit and keep the changes, but not staged:

```
$ git reset HEAD~
```

And if you didn't want those changes at all - *WARNING* this will lose changes - gone:

```
$ git reset --hard HEAD~
```

11.3 Fetching

Update all local remote tracking branches from all remotes:

```
git fetch --all
```

Update all local remote tracking branches from origin:

```
git fetch origin
```

Update/create local branch origin/master from remote origin's branch master with default configuration:

```
git fetch origin master
```

Update/create local branch 'tmp' from remote origin's branch master (but only updates if fast-forward is possible):

```
git fetch origin master:tmp
```

Peek at an arbitrary remote's branch by pulling it into a (temporary) local branch, then check its log. The temporary local branch will eventually be garbage collected:

```
git fetch git://git.kernel.org/pub/scm/git/git.git maint
git log FETCH_HEAD
```

11.4 Branches and checkouts

Check out an existing branch:

```
git checkout <branch>
```

Create new branch:

```
git branch <branchname> [<start point>]
```

Create new branch and check it out in one command:

```
git checkout -b <newbranch> [<start point>]
```

11.5 Import one repo into another with history

<http://stackoverflow.com/questions/1683531/how-to-import-existing-git-repository-into-another>

11.6 Cleaning

Delete untracked files (be careful!):

```
git clean -fdx
```

Prune branches that have been merged and are no longer upstream:

```
http://devblog.springest.com/a-script-to-remove-old-git-branches
```

Prune branches that track remote branches that no longer exist (<http://kparal.wordpress.com/2011/04/15/git-tip-of-the-day-pruning-stale-remote-tracking-branches/>):

```
$ git remote prune origin --dry-run
$ git remote prune origin
```


11.7 Pulls

Easier access to pull requests on Github. Add to config:

```
# This will make pull requests visible in your local repo
# with branch names like 'origin/pr/NNN'
# WARNING: This also breaks adding a new remote called "origin" manually
# because git thinks there already is one. Comment this out temporarily
# in that case, unless you can think of a better solution.
[remote "pulls"]
    fetch = +refs/pull/*/head:refs/remotes/origin/pr/*
```

11.8 Aliases

Handy aliases for config:

```
[alias]
lg = log --oneline --graph --date-order
lgd = log --oneline --graph --date-order --format=format:"%ai %d %s\"

cb = checkout -b
cd = checkout develop
co = checkout

gd = !git fetch origin && git checkout develop && git pull origin develop
gm = !git fetch origin && git checkout master && git pull origin master

# push -u the current branch
pu = "!CURRENT=$(git symbolic-ref --short HEAD) && git push -u origin $CURRENT"

# push -f
pf = push -f

# Find the common ancestor of HEAD and develop and show a diff
# from that to HEAD
dd = "!git diff $(git merge-base develop HEAD)"
# Find the common ancestor of HEAD and master and show a diff
# from that to HEAD
dm = "!git diff $(git merge-base master HEAD)"

# These need 'hub' installed.
# Create pull request against develop. Must pass issue number.
#pr = pull-request -b develop -i
# Create pull request against develop, not passing issue number:
pr = pull-request -b develop

# Checkout pull request
# Assume origin/pr/NN is pull request NN
# Need a bash function because we need to concatenate something to $1
#cpr = "!f() {set -x;git checkout origin/pr/$1; };f"
cpr = "!gitcpr"

# Undo any uncommitted changes
abort = checkout -- .
```

11.9 Submodules

This will typically fix things:

```
git submodule update --init --recursive
```

(and yes, you need `--init` every time)

Add a new submodule [<http://git-scm.com/book/en/Git-Tools-Submodules>]

```
$ git submodule add git@github.com:mozilla/basket-client basket-client
```

11.10 Combining feature branches

Suppose you have branch A and branch B, which branched off of master at various times, and you want to create a branch C that contains the changes from both A & B.

According to Calvin: checkout the first branch, then `git checkout -b BRANDNEWBRANCH`. then rebase it on the second.

(SEE DIAGRAMS BELOW)

Example:

```
# Start from master
git checkout master
git pull [--rebase]

# Create the new branch from tip
git checkout -b C

# rebase A on master
git checkout A
git rebase -i master
# merge A into C
git checkout C
git merge A

# rebase B
git checkout B
git rebase -i master
# merge B into C
git checkout C
git merge B

# I think???
# Review before using, and verify the result
```

Combining git branches diagrams

Start:

```
o - o - o - o <--- master
 \   \
  \   o - o - o <--- A
   o - o - o <--- B
```


CHAPTER 12

Google APIs

Google APIs and Docs

Google Apps

Google Apps APIs

Google Drive API

Google Apps Script

12.1 Google libraries

These are all for Python.

12.2 Older libraries (Google Data)

12.3 Newer libraries

OAuth 2.0

Google APIs client library for Python (beta)

Because gulp tasks are actually from another package, [orchestrator](#), the gulp people don't feel the need to document how gulp tasks work very much. So...

13.1 Asynchronous

Gulp tasks are asynchronous. Here are some patterns to use.

13.1.1 Just exit

THIS IS ALMOST ALWAYS WRONG

Gulp calls the task function, which returns when it's done.

Even if everything your task is doing is synchronous, this is *still wrong* because Gulp is allowed to kick it off in the background and then immediately assume it's done and start running other tasks that depend on it, which might or might not work. A real bear to debug. But here's how it looks:

```
# DO NOT DO THIS:
gulp.task('sync_task', function() {
  stuff that runs immediately;
})
```

13.1.2 Callback

Gulp passes a callback. The tasks calls it when done:

```
gulp.task('foo', function(cb) {
  stuff...;
  cb();
})
```

13.1.3 Return a promise or an event stream

Gulp pipes, and lots of gulp plugins, return a promise or even stream. We can just return the thing they gave us, and gulp will use that to tell when the work is done:

```
gulp.task('make_promise', function() {  
  return some_gulp_plugin.pipe(stuff).pipe(more_stuff);  
})
```

13.2 Watching

Call a task whenever any of a group of files changes:

```
gulp.watch('filepattern', ['task1', 'task2'])
```

After calling this, anytime a file that existed when this was called and matched the filepattern changes, task1 and task2 will be called.

This *will result* in the gulp command not exiting, because it's hanging around to watch for changes!

(But the watch() function itself does return, no worries there.)

TBD: What does gulp.watch return? Anything?

Make sure to use the docs corresponding to the version you are using.

1.5: <https://cbonte.github.io/haproxy-dconv/1.5/configuration.html>

14.1 Pass SSL thru

Use proxy “mode tcp”. E.g:

```
listen sectionname
    bind :443
    mode tcp
    server server1 10.0.0.1:443
    default_backend sslserver
backend sslserver
    mode tcp
    server servername 1.2.3.4:443
```

14.2 Route based on SNI

This works even if haproxy is not terminating the SSL connection:

```
acl site_b req_ssl_sni -i site_b.com
use_backend site_b_backend if site_b
backend site_b
    mode tcp
    server b1 10.0.0.1:443
    server b2 10.0.0.2:443
```

14.3 Route based on Host request header

Use an ACL to check the header and then pick a backend:

```
acl site_a hdr(host) -i site_a.com
use_backend site_a_backend if site_a
backend site_a_backend
    mode http
    server a1 10.0.0.1:80
    server a2 10.0.0.2:80
```

By analogy to Python...

15.1 Imports

Python	Haskell	Notes
<code>from module import *</code>	<code>import module</code>	Import everything from the module directly into the namespace
<code>from module import a,b</code>	<code>import module (a,b)</code>	Import selected items from the module directly into the namespace
<code>import module</code>	<code>import qualified module</code>	Make all <code>module.NAME</code> available in the name space
<code>import module M = module del module</code>	<code>import qualified module as M</code>	Give module a shorter alias
<ul style="list-style-type: none">•	<code>import module hiding (c)</code>	Import everything from the module directly into the namespace, except selected items

15.2 Math

Python	Haskell	Notes
a+b, a*b, a-b	a+b, a*b, a-b	
	a/b	Integer division giving a float
a and b, a or b, not a	a && b, a b, not a	
a == b, a != b	a==b, a /= b	
min(a,b), max(a,b)	min a b, max a b	
min([a,b,c]), max([a,b,c]), sum([a,b,c])	minimum [a,b,c], maximum [a,b,c], sum [a,b,c]	

15.3 Expressions

Python	Haskell	Notes
b if a else c	if a then b else c	Both are expressions, not statements

15.4 Lists

Python	Haskell	Notes
[1,2,3,4]	[1,2,3,4]	
[1, 'a', "foo"]	•	Haskell lists must be homogeneous
[1,2] + [3,4]	[1,2] ++ [3,4]	
[1] + [2,3]	1:[2,3]	Haskell has more syntax for list expressions
['a','b','c'][1]	['a', 'b', 'c'] !! 1	Extract a list element
['a','b','c'][:3]	take 3 ['a','b','c']	First N elements
['a','b','c'][3:]	drop 3 ['a','b','c']	Drop N, take rest
['a','b','c'][1:2]	take 2 . drop 1 ['a','b','c']	Slices are uglier in Haskell
['a','b','c'][0]	head ['a','b','c']	First elt
['a','b','c'][1:]	tail ['a','b','c']	All but first elt
['a','b','c'][-1]	last ['a','b','c']	
['a','b','c'][:-1]	init ['a','b','c']	All but last elt
len(['a','b','c'])	length ['a','b','c']	
not len(['a','b','c'])	null ['a','b','c']	is list empty (but would probably just do bool(list) in Python)
reverse(['a','b','c'])	reverse ['a','b','c']	
'a' in ['a','b','c']	'a' elem ['a','b','c']	
range(1,21)	[1..20]	Python counts to the last value before the 2nd parm, Haskell includes it
set(x) - set(y)	x \ y	Set difference for lists x and y - except Haskell isn't really, it just removes one value 'z' from x for each occurrence of 'z' in y. See Data.Set for real sets.
set(x) + set(y)	x union y	Haskell adds one occurrence of each element of y to x if x doesn't already have one

15.5 Functions

Python	Haskell	Notes
def doubleMe(x): return x + x	doubleMe x = x + x	
def doubleUs(x, y): return x*2 + y*2	doubleUs x y = x*2 + y*2	

My i3 bindings

“**W**” is the “windows” key

Change what we see:

```

W-<NUMBER>: switch to workspace NUMBER on whatever monitor it's
             attached to.
W-Control-1: Only use built-in laptop display
W-Control-2: Use built-in laptop display, and external display
              positioned to its right
W-Control-3: Use built-in laptop display, and external display
              positioned above it
W-<n>:       Switch to workspace <n> (need not already exist)
             (if workspace <n> is on another screen, it'll switch
              that screen to workspace <n>, not your current screen)
W-<n> where <n> is the current workspace: Switch back to previous
workspace (So you can just do W-1 (look at screen) W-1
and be back where you started)

```

Focus:

```

W-j, W-<left>   left
W-k, W-<down>   down
W-l, W-<up>     up
W-;, W-<right>  right

```

Move things:

```

W-Control-<ARROW>: Move current workspace to another monitor.
W-Shift-Number:  Move current window to another workspace
                  (need not already exist)
Shift-<FOCUS COMMAND>: Move current window within workspace

```

Layouts:

⌘-e	default (splith/splitv), repeat to toggle splith/splitv
⌘-s	stacked
⌘-w	tabbed
⌘-f	fullscreen (toggle)
⌘-S-<spc>	float (toggle)
⌘-mouse1-drag	move floating
⌘-h	Make the current window/container a horizontal split container. New windows opened when this container is focused will be created by splitting this container horizontally (side-by-side)
⌘-v	Like ⌘-h, but vertical (one above another)
⌘-e	toggle between defaulting to horizontal and defaulting to vertical

Start/end things:

⌘- return :	open new terminal
⌘-D:	open dmenu at top to enter a command (output invisible, use to start new graphical programs)
⌘-S-q	kill window

Control I3:

⌘-S-c	reload I3 config
⌘-S-r	restart I3
⌘-S-e	kill I3 (logout)

Resizing:

⌘-mouse2-drag	stretch or shrink window
---------------	---------------------------------

Screen capture:

<Printscreen>	- capture whole screen
Shift-<Printscreen>	- select a rectangle or window (?)
Control-<Printscreen>	- capture currently focused window

17.1 Localhost

Localhost is `::1`

17.2 Any addr

Any addr (equivalent of 0.0.0.0) is `::`

17.3 With port

Adding “:<portnumber>” to an IPv6 address would be ambiguous. The solution is to put “[]” around the address part, e.g. `[::1]:8000`.

17.4 URLs with IPv6 addresses

You need the `[]` here too, at least if you’re using a hexadecimal IPv6 address. (Even without a port number.):

```
http://[fe80::b746:6473:e65f:5dd4]/foo/bar  
http://[fe80::b746:6473:e65f:5dd4]:8000/foo/bar
```

17.5 Ping

Use `ping6`:

```
ping6 ::1
```

17.6 Django

To run a local dev server listening on any IPv6 address:

```
python manage.py runserver --ipv6 "[::]:8000"
```

It does NOT appear possible to have the dev server listen on both IPv4 and IPv6, at least not in Django 1.8. (But I'm sure you could put nginx in front to do that for you.)

17.7 Private IPv6 network addresses

Try <http://simplifiedns.com/private-ipv6.aspx> to get a random private address range.

Example output:

```
Here is a unique private IPv6 address range generated just for you (refresh page to  
↪ get another one):
```

```
Prefix/L:      fd  
Global ID:     442da008f4  
Subnet ID:     cf4f  
Combined/CID:  fd44:2da0:08f4:cf4f::/64  
IPv6 addresses: fd44:2da0:08f4:cf4f:xxxx:xxxx:xxxx:xxxx
```

```
If you have multiple locations/sites/networks, you should assign each one a different  
↪ "Subnet ID", but use the same "Global" ID for all of them.
```

To add an address to your loopback interface:

```
sudo ifconfig lo add fd44:2da0:08f4:cf4f::1/64
```

17.8 Find out your ipv6 address

In theory, run “ifconfig” and look for “inet6 addr”. E.g.:

```
enp0s25  Link encap:Ethernet  HWaddr 54:ee:75:8b:03:99  
         inet addr:172.20.1.110  Bcast:172.20.3.255  Mask:255.255.252.0  
         inet6 addr: fe80::b746:6473:e65f:5dd4/64 Scope:Link  
         UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
```

Try to find an “inet6 addr” that isn’t “Scope:Link” or “Scope:Host”; you want “Scope:Global”.

Better way - if you think you have real IPv6 connectivity to the internet, go to google.com and search for “what’s my IP”. If you’re connecting to Google over IPv6, Google will tell you your IPv6 address (which might be a NATting router, I suppose - I don’t know if IPv6 still does that?) But that won’t help if you’re just running IPv6 locally for testing.

Contents:

18.1 Javascript Builtins

18.1.1 Arrays

Traditional literals, append, extend:

```
>> a = [1, 2, 3]                                # literal
>> console.log(a)
[1, 2, 3]
>> b = [4, 5]
>> Array.prototype.push.apply(a, b);           # 'extend'
>> console.log(a)
[1, 2, 3, 4, 5]
>> a.push(6)                                     # 'append'
>> console.log(a)
[1, 2, 3, 4, 5, 6]
```

ES2015 literals, append, extend (Spread operator):

```
>> a = [1, 2, 3]
>> b = [0, ...a]                                # literal incorporating another array
>> console.log(b)
[0, 1, 2, 3]
>> a.push(4)                                    # append
>> console.log(a)
[1, 2, 3, 4]
>> console.log(b)
[0, 1, 2, 3]
>> a.push(...b)                                # extend
```

(continues on next page)

(continued from previous page)

```
>> console.log(a)
[1, 2, 3, 4, 0, 1, 2, 3]
```

18.1.2 Iterate over an array

Run code for each element (ES5+) ([forEach](#)):

```
arr.forEach(function(value, index, arr){}, thisArg)
```

See if none, some, or all elements pass a test (ES5+) ([every](#), [some](#)):

```
all_passed = arr.every(function(value, index, arr){ return bool;}, thisArg)
some_passed = arr.some(function(value, index, arr){ return bool;}, thisArg)
none_passed = !arr.some(...)
```

Create a new array by applying a function to each element of an existing array (ES5+) ([map](#)):

```
arr2 = arr1.map(function(value, index, arr){}, thisArg)
```

Create a new array only including elements that pass a test (ES5+ ([MDN](#)):

```
arr2 = arr1.filter(function(value, index, arr){}, thisArg)
```

18.1.3 Iterate over a “dictionary” (JS object)

This will include enumerable properties of prototype classes (traditional) ([for...in](#)):

```
for (var key in object) {
  value = object[key]
  # ... do stuff ...
}
```

Note: For objects that you’re just using as dictionaries, *for...in* is perfectly good. Just be careful not to use it on instances of anything more complicated.

For more complicated objects, here’s how you get just properties that are directly on the object (not inherited from a prototype). Traditional:

```
for (var key in object) {
  if (object.hasOwnProperty(key)) {
    value = object[key]
    # ... do stuff ...
  }
}
```

Just properties directly on the object, with ES2015+ ([keys](#)):

```
var keys = object.keys();
for (var i = 0; i < keys.length; i++) {
  value = object[keys[i]];
  # ... do stuff ...
}
```

18.1.4 Is a key in a dictionary?

```
dict.hasOwnProperty(key)
```

18.1.5 Does an object have a key (possibly inherited)?

```
key in object
```

18.1.6 Remove key from dictionary

```
delete dict["key"] delete dict.key
```

18.1.7 String operations

String, endsWith, includes, indexOf, join, match, replace, search, split, startsWith, substr, substring

```

arr = str.split(sep=(str|regex)[, limit])
str = arr.join(sep)
index = str.indexOf(substr[, startIndex]) # returns -1 if not found
sub = str.substr(firstIndex[, length]) # firstIndex can be negative to count back_
↳from end
sub = str.substring(firstIndex[, lastIndex+1])
str2 = str.replace(regex|substr, newSubStr|function)
bool = str.startsWith(str2)
bool = str.includes(str2)
bool = str.endsWith(str2)
[matchstr, groups...] = str.match(regex) # returns null if doesn't match
[matchstr, groups...] = str.search(regex) # returns null if doesn't match entire_
↳string

```

Contains: haystack.indexOf(needle) !== -1

18.1.8 Timer

setTimeout:

```
window.setTimeout(func, delay, param1, param2, ...);
```

All but func is optional. *delay* defaults to 0.

```
timerId = window.setTimeout(func, [delay, param1, param2, ...]); window.clearTimeout(timerId);
```

18.2 Javascript Events

18.2.1 Event object

- `currentTarget`: the element that has the handler attached that is currently running
- `target`: the element that received the event initially

18.3 NVM and NPM

Use another node version:

```
$ nvm use 8
$ nvm use 10
```

etc.

Back-level your npm:

```
$ which npm
```

to make sure you're running npm from an nvm version, then:

```
$ npm install -g npm@5.10.0
$ npm version
...
```

18.4 Promises

[MDN docs](#)

quoting from there:

A Promise is in one of these states:

- pending: initial state, not fulfilled or rejected.
- fulfilled: meaning that the operation completed successfully.
- rejected: meaning that the operation failed.

A promise that is not pending is “settled”.

You can create a Promise with `new`, passing it an executor function:

```
let p = new Promise((resolve, reject) => {
  if (things go well) {
    resolve(resulting_value)
  } else {
    reject(error)
  }
})
```

At any time after the promise is created, you can call `then` on it and pass in a success handler, a failure handler, or both:

```
p.then(on_success) p.then(undefined, on_failure) p.then(on_success, on_failure)
```

If the promise is still pending, the appropriate handler will be called once it is settled. If it's settled already, the handler will be called immediately (or ASAP, anyway).

The `on_success` handler is called with the resolved value from the promise, while the `on_failure` handler is called with the error value from the promise.

`then()` returns a *new* promise, so you can chain calls:

```
p.then(on_success).then(another_on_success).then(a_third_on_success)
```

If the `on_success` handler returns a new promise, *that* promise will be the return value from `then()` (or an equivalent promise, anyway).

The handlers will be called in order. If one fails, then the promise returned by that `then()` call will be rejected. “Fail” can mean raises an exception.

Many async functions nowadays return promises.

18.4.1 Pausing

Here's a way of doing something after a delay:

```
let p = new Promise((resolve) => {
  setTimeout(resolve, 100)
}).then(()=>{
  do stuff
})
```

18.4.2 Some notes I made a while ago

Here are some examples:

```
// Return a new promise.
return new Promise(function(resolve, reject) {
  // do stuff.  if it works:
  resolve(result);
  // but if it fails:
  reject(Error(error text));
});

// USE a promise
promise.then(function(result) {
  // use result here
}, function(error) {
  // do something with error here
});

// CHAINING
// Note that 'then' returns a new promise

promise.then(function(val1) {
  return val1 + 2;
}).then(function(val2) {
  val2 is here val1 + 2!!!
});

// MORE CHAINING
promise.then(function(result){
  // do something, then return a new promise
  return promise2;
}).then(function(result2) {
  // this isn't invoked until promise2 resolves, and it gets promise2's result.
```

(continues on next page)

(continued from previous page)

```
...  
}
```

18.5 Javascript Syntax

18.5.1 Spread Operator

Spread operator

Given function and args:

```
function myFunction(x, y, z) { }  
var args = [0, 1, 2];
```

Traditional:

```
myFunction.apply(null, args);
```

ES2015:

```
myFunction(...args);  
myFunction(1, ...[2, 3]);
```

Caution: The ES2016 `...` operator is NOT the same as `*` in a Python function call. `...` basically splices the array it's applied to into the list at the place where it's used. It can be used repeatedly, and in any combination with other unnamed arguments. Python's `*` can only be used to extend the list of unnamed arguments at the end.

18.6 DOM operations with JavaScript

18.6.1 Finding elements by selector

<https://developer.mozilla.org/en-US/docs/Web/API/Element/querySelectorAll>

```
items = document.querySelectorAll('a.class-name')
```

Returns a `NodeList`.

18.6.2 Iterating over elements

<https://developer.mozilla.org/en-US/docs/Web/API/NodeList>

Note: Although `NodeList` is not an `Array`, it is possible to iterate on it using `forEach()`. Several older browsers have not implemented this method yet.

```
// Newer browsers
items.forEach(function(item){ do something with item })

// Old (how old?) browsers
Array.prototype.forEach.call(items, function(item) { do something with item })
```

Here's a polyfill for ES5+ browsers from that `NodeList.forEach` page linked above:

```
if (window.NodeList && !NodeList.prototype.forEach) {
  NodeList.prototype.forEach = function (callback, thisArg) {
    thisArg = thisArg || window;
    for (var i = 0; i < this.length; i++) {
      callback.call(thisArg, this[i], i, this);
    }
  };
}
```


19.1 Activating without connecting to kobo on the internet

- Start up the kobo. It'll ask if you want to setup over wifi or not. Say not.
- Plug the kobo into your computer using a USB cable.
- On my laptop, it mounts automatically at /media/usb0; I don't know if I set that up somehow in the past.
- Use sqlite3 to add a dummy user:

```
$ sqlite3 /media/usb0/.kobo/KoboReader.sqlite3
sqlite> INSERT INTO user (UserID, UserKey, UserDisplayName, UserEmail) values(
↪ "00000000-0000-0000-0000-000000000000", "00000000-0000-0000-0000-000000000000",
↪ "MyDummyUser@dummy.com", "MyDummyUser@dummy.com");
sqlite> .quit
```

Before rebooting, continue by manually updating the latest firmware, see next section.

19.2 Manually updating firmware

- Download a zipfile with the latest firmware for your Kobo device from [this thread](#)
- The Kobo should be mounted over USB as above.
- Change to the kobo's .kobo directory
- Unpack the zipfile containing the firmware update
- Properly unmount the Kobo
- Unplug the USB cable

Within a minute or so, the Kobo should notice the update, install it, and reboot.

19.3 Installing KSM (Kobo Start Menu)

- [instructions](#)
- [download v8 zipfile](#)
- follow the instructions, they're not bad
- One thing easy to miss: once it's starting okay, "From the home menu select "tools" > "activate" > "set runmenu settings.msh." There you will find the options "always," "once," etc. Choose "always" to make the Kobo Start Menu appear after each time you power the device on.)"

19.4 Installing Koreader

- Download the latest koreader release from [here](#)
- Follow the instructions from above

19.5 Sideloading books

Just mount on USB as above and copy epub files to the root directory of the kobo, or to any subdirectory (except subdirs starting with ".", which it won't look in).

20.1 Creating RPMS

- from ‘make install’: See [<http://asic-linux.com.mx/~izto/checkinstall/index.php> checkinstall]
- from CPAN packages: See [<http://perl.arix.com/cpan2rpm/> cpan2rpm]

20.2 Increase open file limit

Edit `/etc/security/limits.conf` and add or edit:

```
hard nofile 2048
soft nofile 2048
```

to change the default limits set at login. Also double-check `/etc/profile` etc for `ulimit` commands that might set it lower again.

20.3 Trace file operations

```
strace -F -e trace=file,process -o strace.out ‘command’
```

20.4 Changing fonts for GAIM, Sanity, etc.

Run `gnome-font-properties`

You *also* need to start `gnome-settings-daemon` at login time if you’re not running the Gnome desktop.

Q: How do you use anti-aliased fonts in gnome2?

A: You need to set the GDK_USE_XFT environment variable, if you are using bash type 'export GDK_USE_XFT=1'. You should not need to do this - your Linux distribution should do this automatically. Also AntiAliasingTroubleshoot-ing may help you find the problem with your set-up.

- __How do you make the entire session use anti-aliased fonts?__

Create a file, '.gnomerc', in your home directory and put in it

```
#!/bin/sh
export GDK_USE_XFT=1
exec gnome-session
```

Note

If you use a display manager such as gdm, you don't need the "exec gnome-session" line and it may in fact be counterproductive.

You may also have to edit the file /etc/X11/XftConfig. Look for lines with:

```
#
# Don't antialias small fonts
#
match
    any size < 14
    any size > 8
    edit antialias=false;
```

and comment them out.

Added 7-16-2002

NOTE /etc/X11/XftConfig is deprecated. Use /etc/fonts/fonts.conf now. (11-26-2002)

- __Is the Bitstream Vera font anywhere available for download?__

Yes, you can get them from <http://ftp.gnome.org/pub/GNOME/sources/ttf-bitstream-vera/>

In debian just type:

```
apt-get install ttf-bitstream-vera
```

In gentoo just type:

```
emerge ttf-bitstream-vera
```

Updated 2003-07-14

- __Where do I put new (downloaded) truetype fonts?__

You can just put them into \$HOME/.fonts directory

See <http://gnome-hacks.jodrell.net/hacks.html?id=42> for how to make GNOME settings apply elsewhere.

LIRC: linux infrared remote control

LIRC

21.1 Intro

This is about how to use an arbitrary infrared remote control (like for your TV) to control your own program running under Linux.

The end result will be that LIRC will receive signals from your remote and translate them to keypresses, and your program can connect to LIRC and get these keypresses as they happen, then do whatever it wants with them.

21.2 Difficulties

One problem I ran into getting started with most of the howtos and tutorials I found on the web was that almost anything not quite right in your configuration resulted in things not working with no indication of what the problem was. That is, you could press buttons on the remote all day and there was no indication that the PC was seeing any of it.

21.3 Why is this so hard

Here's what I *think* is going on, based on how these things seem to need to be configured etc.

Sending data via infra-red is messy. An infra-red receiver attached to your computer just continuously measures and reports the level of infra-red frequency light hitting it.

Then something we'll call a decoder program - e.g. LIRC - has to monitor that continuous stream of light levels and try to spot it when extra infra-red from a remote is hitting it, by the pattern of changes to the levels.

There's no standard on how IR remotes encode data into IR. So the only reasonable way to make this work is to tell the decoder what remote or remotes we're expecting to see commands from, so it can try to match up what it's seeing to commands from those specific remotes. Otherwise there are just too many possibilities to watch for them all.

The unfortunate consequence of this is that if we haven't told it the right remote, it just won't see any IR commands at all.

Luckily, there are some very smart people who have built tools that can look at input from a remote for a while and try to guess what remote it is, and we might have to resort to using them. But it's much easier to use a configuration someone else has already worked out for us, so we'll try that first.

21.4 How to tell if it's working

We need an easy way to tell if LIRC is seeing and correctly interpreting our remote commands. We'll use the *irw* tool.

Run `irw /var/run/lirc/lircd`, then start pressing buttons on the remote. If things are working, it should print out codes and key names. If not, you probably won't see anything.

When done, hit *Control-C* to stop *irw*.

21.5 The easy start

So, I'll start with the approach most often documented, that if it works, is simple. But if this doesn't work, don't waste a lot of time fiddling with it. Move on to the next section.

This approach will only work if LIRC has a configuration file for your remote - and you know which one it is.

1. Install LIRC. On Ubuntu, `sudo apt install lirc` should do it.
2. During the install, Ubuntu will prompt you to pick your remote so that it can configure LIRC for you. Feel free to try this.
3. If LIRC is already installed, or you want to try a different configuration, you can re-run the install-time configuration using `sudo dpkg-reconfigure lirc`.
4. Check if things are working as described above. If they are, you can skip down to the section on using LIRC input from a program. If not, you can try configurations for other remotes that seem likely.

21.6 If the easy approach doesn't work

LVM on Linux

Reference <http://tldp.org/HOWTO/LVM-HOWTO/index.html>

22.1 Terminology

PV Physical volume (e.g. a partition, RAID array, etc)

VG Volume Group - a collection of PV's that we can use the space from

LV Logical volume - a partition created from space in a VG

22.2 Physical volumes

- List physical volumes : `pvdisplay`
- Info about one PV : `pvdisplay <PV name>`
- Partition type for LVM: 8e
- Make a PV from a partition : `pvcreate <partition>`

22.3 Volume groups

- Create a volume group : `vgcreate <NewVGName> <PVname> [<PVname>...]`
- Add PV to VG : `vgextend <VGname> <PVname>`
- Remove PV from VG : `vgreduce <VG name> <PV name>`
- List volume groups : `vgdisplay`

22.4 Logical volumes

- List logical volumes : `lvdisplay`
- Create LV : `lvcreate -L<SIZE> -n<NewLVName> <VGname>` (SIZE=<num><units>, e.g. 1.47TiB) or `-l<EXTENTS>`
- Device name of the logical volume = `/dev/<VGname>/<LVname>`
- Enlarge LV : `lvextend -l+<extents> /dev/<VGname>/<LVname>`

22.5 Resize file system after enlarging LV

```
sudo ext2fs -f /dev/<VGname>/<LVname> sudo resize2fs /dev/<VGname>/<LVname>
```

Using LXDE desktop with i3 window manager.

There's a brief note [here](#) but this gives a little more depth.

- Install lxde:

```
sudo apt install lxde lxsession-logout
```

- Logout of the desktop
- Login again, this time choosing the LXDE desktop
- Create an executable shell script somewhere on your path, naming it “i3wm”. It should run “i3”. (I don’t know why it doesn’t work to just set the window manager to i3, but it doesn’t. Maybe someday I’ll take the time to debug that.)
- Edit `~/.config/lxsession/LXDE/desktop.conf`. In the `[Session]` section, change `windows_manager/command: windows_manager/command=i3wm`
- In `~/.config/lxsession/LXDE/autostart`, remove “@pcmanfm –desktop –profile LXDE”, it interferes with i3.
- Logout and login again.
- If you like, bind a key in i3 to “lxsession-logout” and use that to logout. Exiting i3 will *not* log you out with this configuration. Or just use the menus in lxpanel to log out.

Logitech Harmony

Programming it using Linux

I have the Harmony One model (no longer produced).

Some of this information comes from <http://ubuntuforums.org/showthread.php?t=781059>, but I'm not using the GUI tool (congruity), just command line.

Help *Harmony Devices* for finding devices.

24.1 Harmony Devices

```
<select class="devicetype_select" id="selectBoxVisible1" name="deviceType1"
↪onmousedown="openMenu(this);event.cancelBubble = true;">
<option value="none">-select device to add-</option>
```

TELEVISION:

```
<option value="1">TV</option>
<option value="7">Projector</option>
<option value="13">Monitor</option>
<option value="15">TvVcr</option>
<option value="37">TV DVD</option>
<option value="38">TV DVD VCR</option>
<option value="48">TV HDD</option>
```

AMPLIFIER:

```
<option value="19">Amplifier</option>
<option value="5">AV Receiver</option>
<option value="30">Audio/Video Switch</option>
<option value="36">Radio Tuner</option>
```

CABLE/SATELLITE BOX:

(continues on next page)

(continued from previous page)

```
<option value="14">Cable Box</option>
<option value="16">Satellite</option>
<option value="12">Digital Set Top Box</option>
```

VIDEO RECORDER:

```
<option value="2">VCR</option>
<option value="18">PVR</option>
<option value="25">DVD VCR</option>
<option value="38">TV DVD VCR</option>
<option value="15">TvVcr</option>
<option value="32">Mini System (DVD, VCR, Radio)</option>
<option value="47">DVDR VCR</option>
```

DVD:

```
<option value="4">DVD</option>
<option value="34">DVD Recorder</option>
<option value="25">DVD VCR</option>
<option value="37">TV DVD</option>
<option value="38">TV DVD VCR</option>
<option value="22">Mini System (DVD, CD, Radio)</option>
<option value="32">Mini System (DVD, VCR, Radio)</option>
<option value="33">Digital Music Server</option>
<option value="28">Laserdisc Player</option>
<option value="47">DVDR VCR</option>
```

MUSIC PLAYER:

```
<option value="3">CD Player</option>
<option value="11">CD Jukebox</option>
<option value="28">Laserdisc Player</option>
<option value="33">Digital Music Server</option>
<option value="36">Radio Tuner</option>
<option value="21">Mini System (CD, Radio, Cassette)</option>
<option value="6">Tape Deck</option>
<option value="29">Minidisc Player</option>
<option value="20">DAT</option>
```

GAME CONSOLE:

```
<option value="9">Game Console</option>
<option value="23">Game Console (With DVD)</option>
```

MINI SYSTEM:

```
<option value="22">Mini System (DVD, CD, Radio)</option>
<option value="21">Mini System (CD, Radio, Cassette)</option>
<option value="32">Mini System (DVD, VCR, Radio)</option>
```

COMPUTER:

```
<option value="8">Computer</option>
<option value="10">Laptop</option>
<option value="35">Media Center PC</option>
```

(continues on next page)

(continued from previous page)

HOME AUTOMATION:

```
<option value="24">Light Controller</option>
<option value="45">Climate Control</option>
<option value="44">Home Appliance</option>
```

MORE DEVICE TYPES:

```
<option value="6">Tape Deck</option>
<option value="29">Minidisc Player</option>
<option value="20">DAT</option></select>
```

UNDER COMPUTER (8): MANUFACTURERS

```
<select class="manufacturer_select" id="manufacturerDropDown1" name="manufacturer1"
  ↳ onblur="if (is_ie) changeManufacturer(this);" onchange="if(manufacturerLastAction_
  ↳ == 'mouse' || !is_ie) changeManufacturer(this);" onmousedown=
  ↳ "mousedownManufacturer(this);" onmouseup="manufacturerLastAction = 'mouse';"
  ↳ onkeypress="manufacturerLastAction = 'keyboard'; if(event.keyCode == 13)
  ↳ changeManufacturer(this) ;">
<option value="none">--select manufacturer--</option>
<option value="unavailable">--not listed--</option>
<option value="57283">@AlfaLine</option>
<option value="43116">@Xi Computer</option>
<option value="42270">10Moons</option>
<option value="15767">2PartsFusion</option>
<option value="18631">2Wire</option>
<option value="1096">3COM</option>
<option value="141159">3GO</option>
<option value="117481">3Q</option>
<option value="28619">3R System</option>
<option value="72513">4Geek</option>
<option value="147373">8level</option>
<option value="8394">Abit</option>
<option value="4594">ABS</option>
<option value="15631">Absolut Technology SA</option>
<option value="74609">AC Ryan</option>
<option value="3039">accessDTV</option>
<option value="35851">ACE</option>
<option value="227">Acer</option>
<option value="9418">Acesonic</option>
<option value="1296">Acoustic Research</option>
<option value="92895">Acrobak</option>
<option value="5666">Actiontec</option>
<option value="12962">Adaptec</option>
<option value="5745">Adesso</option>
<option value="42184">Adrenaline</option>
<option value="50556">ADS</option>
<option value="9184">Advance</option>
<option value="2199">Advanced Digital Broadcast</option>
<option value="10823">Advanced PC</option>
<option value="66439">Advantech</option>
<option value="894">Advent</option>
<option value="24911">AESO</option>
<option value="11838">AFK</option>
<option value="25068">AGE Computer</option>
```

(continues on next page)

(continued from previous page)

```
<option value="50973">AGFAPHOTO</option>
<option value="4065">Ahanix</option>
<option value="23905">AIC</option>
<option value="34025">Aigo</option>
<option value="15652">Aim</option>
<option value="161051">aios</option>
<option value="27289">AirLink</option>
<option value="75793">AirLive</option>
<option value="123523">AKASA</option>
<option value="199534">AKASO</option>
<option value="8657">AL Tech</option>
<option value="4079">Albatron</option>
<option value="201058">Alfawise</option>
<option value="11502">Ali</option>
<option value="4711">Alienware</option>
<option value="18651">Allen Bradley</option>
<option value="18441">ALLNET</option>
<option value="38786">Alpha Digital</option>
<option value="1263">Altec Lansing</option>
<option value="24938">Altex</option>
<option value="43902">Aluratek</option>
<option value="109710">Amazon</option>
<option value="245">AMC</option>
<option value="10654">AMD</option>
<option value="24204">AMI</option>
<option value="15014">Amisos</option>
<option value="7269">Amitech</option>
<option value="91359">AM-Logic</option>
<option value="3580">Amoi</option>
<option value="54025">Ampaq</option>
<option value="9141">AMS</option>
<option value="123">Amstrad</option>
<option value="301">AMX</option>
<option value="137244">Android</option>
<option value="4567">Antec</option>
<option value="24238">Anysee</option>
<option value="54131">Anyware</option>
<option value="2616">AOC</option>
<option value="7094">AOpen</option>
<option value="12799">Apacer</option>
<option value="52">Apex</option>
<option value="388">Apple</option>
<option value="20543">Aquado</option>
<option value="489">Archos</option>
<option value="116379">Arctic</option>
<option value="30792">Arcus</option>
<option value="85471">Ared</option>
<option value="21339">Argosy</option>
<option value="11235">Arisetec</option>
<option value="48963">Armada</option>
<option value="196687">ARNU BOX</option>
<option value="156741">AsiaBox</option>
<option value="131858">Asmax</option>
<option value="18605">ASRock</option>
<option value="48075">Astone</option>
<option value="270">Asus</option>
<option value="182185">asustor</option>
```

(continues on next page)

(continued from previous page)

```
<option value="2579">ATC</option>
<option value="19681">Atelco</option>
<option value="267">ATI</option>
<option value="11986">Atlas Computers</option>
<option value="35973">ATMT</option>
<option value="9695">Aura</option>
<option value="12613">Auvio</option>
<option value="191526">auxtek</option>
<option value="24339">AV Labs</option>
<option value="12149">Avalon Sound & Vision (ASV)</option>
<option value="102748">AVedia</option>
<option value="2534">AVerMedia</option>
<option value="121771">AVF</option>
<option value="8937">AVM</option>
<option value="152611">AVOV</option>
<option value="13264">Avox</option>
<option value="136433">Avtrol</option>
<option value="42685">AZbox</option>
<option value="14780">Aztech</option>
<option value="29">Bang & Olufsen</option>
<option value="8647">Bass Computers</option>
<option value="86839">Bauhn</option>
<option value="1319">BBK Electronics</option>
<option value="11901">BCOM</option>
<option value="188636">BeeLink</option>
<option value="26970">Beholder</option>
<option value="17406">Best Buy</option>
<option value="113346">Bevix</option>
<option value="111858">BeWan</option>
<option value="14568">Beyond</option>
<option value="43956">Beyonwiz</option>
<option value="3473">Biostar</option>
<option value="33799">Black Box</option>
<option value="13380">Black Gold</option>
<option value="11418">Blu Sens</option>
<option value="11947">Blue Maple Networks</option>
<option value="129611">Blue Times</option>
<option value="50443">Bluegear</option>
<option value="10688">BlueTinum</option>
<option value="51139">BMS</option>
<option value="20962">Bogobox</option>
<option value="26275">BOLData</option>
<option value="11771">Bow</option>
<option value="98564">Boxee</option>
<option value="199148">Bqeel</option>
<option value="23210">Brasse</option>
<option value="6111">Braun</option>
<option value="114081">Brite-View</option>
<option value="6992">BTC</option>
<option value="9810">Buffalo</option>
<option value="14968">Bullit</option>
<option value="47405">BuyCable</option>
<option value="1677">Calrad</option>
<option value="24450">CanonPC</option>
<option value="73">Canon</option>
<option value="3279">Canopus</option>
<option value="80907">Captiva</option>
```

(continues on next page)

(continued from previous page)

```
<option value="38548">Catronics</option>
<option value="191948">CaveTec</option>
<option value="24399">CDC</option>
<option value="59336">Ceconet</option>
<option value="60630">Celem</option>
<option value="20874">Cellar Cinemas</option>
<option value="16045">Centarea</option>
<option value="12856">Central Computer Systems</option>
<option value="1575">CentraLite</option>
<option value="37104">Certified Data</option>
<option value="150693">Ceton</option>
<option value="8474">Chaintech</option>
<option value="40682">Chicony</option>
<option value="12196">chiliGREEN</option>
<option value="81640">Chinavasion</option>
<option value="9167">CiBox</option>
<option value="10157">Cicero</option>
<option value="181616">CIK</option>
<option value="14587">Cinet</option>
<option value="114546">Cirago</option>
<option value="40715">Cisco</option>
<option value="11308">Cisnet</option>
<option value="30881">City Desk</option>
<option value="16692">Claritas</option>
<option value="41535">Clever Tech MCE</option>
<option value="181738">CloudMedia</option>
<option value="184342">CloudNetGo</option>
<option value="24647">CNB</option>
<option value="2641">Coby</option>
<option value="98502">Cocoon</option>
<option value="50372">Columbus Micro</option>
<option value="2380">Commodore</option>
<option value="216">Compaq</option>
<option value="40876">Compositor</option>
<option value="2890">Compro</option>
<option value="115433">CompuLab</option>
<option value="23528">Computer Connections</option>
<option value="55096">Comsis</option>
<option value="18013">Conceptronic</option>
<option value="56247">Connectland</option>
<option value="14887">Conrad</option>
<option value="10076">Control4</option>
<option value="196060">COOD-E</option>
<option value="198727">COOLEAD</option>
<option value="17256">Cooler Master</option>
<option value="50861">Cosk'in</option>
<option value="13440">Cowon</option>
<option value="420">Creative</option>
<option value="9875">Crestron</option>
<option value="16515">CTL</option>
<option value="2536">Cyberlink</option>
<option value="50389">Cybernet</option>
<option value="4496">CyberPower</option>
<option value="26069">Cyber-System</option>
<option value="23355">Cybertron</option>
<option value="78273">Cyclone</option>
<option value="59804">Cytem</option>
```

(continues on next page)

(continued from previous page)

```
<option value="26093">D1</option>
<option value="74740">Dane-Elec</option>
<option value="121445">Dark</option>
<option value="11590">DataFab</option>
<option value="16748">DataGo</option>
<option value="1715">Daytek</option>
<option value="4108">Dazzle Multimedia</option>
<option value="32402">DeBoxx</option>
<option value="151">Dell</option>
<option value="21933">Deltatronic</option>
<option value="22884">Demonite</option>
<option value="42">Denon</option>
<option value="239">Denver</option>
<option value="161197">DESIGNER HABITAT</option>
<option value="93381">Deutsche Telekom AG</option>
<option value="13078">DF Solutions</option>
<option value="25422">DFI</option>
<option value="863">DGTEC</option>
<option value="190475">DHG</option>
<option value="3155">Diamond</option>
<option value="32527">Dicota</option>
<option value="19149">Differo</option>
<option value="58827">Diframe</option>
<option value="30108">Digital Cowboy</option>
<option value="20958">Digital Cube</option>
<option value="68465">Digital Decor</option>
<option value="59576">Digital Perspective</option>
<option value="14290">Digital Rise</option>
<option value="35868">Digital Spectrum</option>
<option value="2279">Digital Stream</option>
<option value="16462">DigitalNow</option>
<option value="21704">Digitus</option>
<option value="189161">DigiXstream</option>
<option value="119269">Digma</option>
<option value="14744">Dign</option>
<option value="16777">Divx</option>
<option value="4190">D-Link</option>
<option value="34543">DNJ Technology</option>
<option value="21413">DNT</option>
<option value="48233">DNTV</option>
<option value="195796">docooler</option>
<option value="200608">dolamee</option>
<option value="81060">Dragon Tech</option>
<option value="24083">Dragonbox</option>
<option value="22280">Dreamsys</option>
<option value="185046">DroidBox</option>
<option value="189375">DROIDPLAYER</option>
<option value="2820">DScaler</option>
<option value="13516">DTS Infocom</option>
<option value="47325">Dueple</option>
<option value="44191">Dune</option>
<option value="124195">DUNE HD</option>
<option value="127586">Dutchtronics</option>
<option value="11698">DVB-Viewer</option>
<option value="3349">Dvico</option>
<option value="25400">D-Vision</option>
<option value="147566">DXtreme</option>
```

(continues on next page)

(continued from previous page)

```
<option value="15647">Dynex</option>
<option value="111187">Dyon</option>
<option value="118485">Eaget</option>
<option value="6448">E-Boda</option>
<option value="51146">EchoLogic</option>
<option value="24179">ECS</option>
<option value="5787">eCube</option>
<option value="70094">eData</option>
<option value="78369">eGreat</option>
<option value="22769">Elecom</option>
<option value="89256">Elektron</option>
<option value="199683">ELEMENT</option>
<option value="35233">Element Electronics</option>
<option value="21924">Elettrodata</option>
<option value="2377">Elgato</option>
<option value="7958">Elitegroup Computer Systems</option>
<option value="8055">Ellion</option>
<option value="41664">Elmak</option>
<option value="6755">Elonex</option>
<option value="3872">E-Machines</option>
<option value="76315">Emgeton</option>
<option value="68368">Eminent</option>
<option value="194722">EMISH</option>
<option value="31448">Emtec</option>
<option value="15621">Energy Sistem</option>
<option value="23781">Enermax</option>
<option value="19846">Enspire</option>
<option value="6050">Entertainment PC</option>
<option value="131252">EnVivo</option>
<option value="187505">Enybox</option>
<option value="6148">EPC</option>
<option value="7196">EPoX</option>
<option value="86735">Epsilon</option>
<option value="11">Epson</option>
<option value="14444">Equator</option>
<option value="46237">Equinux</option>
<option value="24265">Equus</option>
<option value="51">Escient</option>
<option value="40799">Essedi</option>
<option value="49474">Essentiel B</option>
<option value="35392">Eureka</option>
<option value="18369">Euromatrix</option>
<option value="5591">Evation</option>
<option value="19790">Everex</option>
<option value="75981">EverTech</option>
<option value="9633">Evesham</option>
<option value="46001">EVGA</option>
<option value="40522">Evolve</option>
<option value="15036">Extended Systems</option>
<option value="31700">Extreme Solutions</option>
<option value="16524">Falcon Northwest</option>
<option value="59073">Fantec</option>
<option value="69333">FAVI</option>
<option value="1386">Ferguson</option>
<option value="55548">FIA</option>
<option value="24649">FIC</option>
<option value="190107">Filmspeler</option>
```

(continues on next page)

(continued from previous page)

```

<option value="22857">Flipr</option>
<option value="162634">Flirc</option>
<option value="10385">fly video</option>
<option value="67272">Foxconn</option>
<option value="10850">Freecom</option>
<option value="8835">Freevo</option>
<option value="55741">Frontier</option>
<option value="55443">Frost ACM</option>
<option value="9433">Fry's</option>
<option value="143">Fujitsu</option>
<option value="6606">Fujitsu-Siemens</option>
<option value="3457">Fukushima</option>
<option value="103">Funai</option>
<option value="4949">Fusion</option>
<option value="50870">Fusion Research</option>
<option value="904">Galaxy</option>
<option value="501">Gateway</option>
<option value="93981">GBOX</option>
<option value="7440">GB-PVR</option>
<option value="199">Gemini</option>
<option value="35779">Geniatech</option>
<option value="7785">Genius</option>
<option value="127975">Giada</option>
<option value="8908">GIEC</option>
<option value="1790">GigaByte</option>
<option value="53692">Gladiator</option>
<option value="200346">GLOBMALL</option>
<option value="67067">GMC</option>
<option value="131255">Gmini</option>
<option value="165094">gmyle</option>
<option value="61367">Goldlantern</option>
<option value="195893">Goobang Doo</option>
<option value="2089">Goodmans</option>
<option value="36442">Goodview</option>
<option value="117498">Google</option>
<option value="145">GoVideo</option>
<option value="37722">GQ</option>
<option value="11753">Granville</option>
<option value="199751">GreATeVeR</option>
<option value="11433">Griffin</option>
<option value="11555">Gyration</option>
<option value="4183">H&B</option>
<option value="72558">H.TV</option>
<option value="26636">Hallmark Computer</option>
<option value="8358">Hama</option>
<option value="70481">Hantech</option>
<option value="187721">Hardkernel</option>
<option value="14306">Hatch</option>
<option value="829">Hauppauge</option>
<option value="66754">hBox</option>
<option value="67736">hd box</option>
<option value="85526">HD Digitech</option>
<option value="100375">HDI</option>
<option value="144395">HDium</option>
<option value="187596">HDLand</option>
<option value="122653">HDMAX</option>
<option value="1216">Helios</option>

```

(continues on next page)

(continued from previous page)

```

<option value="220">Hewlett Packard</option>
<option value="23980">HFX</option>
<option value="16492">Hifidelio</option>
<option value="5468">Hi-Grade</option>
<option value="77876">HiMedia</option>
<option value="13653">Hiper</option>
<option value="2945">Hisense</option>
<option value="185628">HiSilikon</option>
<option value="45">Hitachi</option>
<option value="74796">HMB</option>
<option value="272">HomeVision</option>
<option value="73717">Honor</option>
<option value="25431">Hoojum</option>
<option value="112080">HornetTek</option>
<option value="67267">HTC</option>
<option value="30264">Huawei</option>
<option value="864">Humax</option>
<option value="3559">Hush</option>
<option value="27848">Hyrican Tango</option>
<option value="2447">Hyundai</option>
<option value="253">IBM</option>
<option value="4997">iBuyPower</option>
<option value="2199">i-CAN</option>
<option value="16745">Ichbinleise</option>
<option value="88642">ICIDU</option>
<option value="67027">iconBIT</option>
<option value="7224">ICube</option>
<option value="37238">Ider</option>
<option value="50062">IDKorea</option>
<option value="31977">iDream</option>
<option value="195352">IDROIDNATION</option>
<option value="12412">i-Friend</option>
<option value="47240">Igor</option>
<option value="34118">IKBENSTIL</option>
<option value="32332">Imagin</option>
<option value="109693">Imation</option>
<option value="177725">iMito</option>
<option value="17594">iMuse Electronics</option>
<option value="166697">Incredisonic</option>
<option value="4509">iNeXT</option>
<option value="46867">Infni</option>
<option value="221">Infocus</option>
<option value="13305">Inmatrix</option>
<option value="24736">INOi</option>
<option value="138751">InOutTV</option>
<option value="109024">Inphic</option>
<option value="15167">Insight</option>
<option value="6394">Insignia</option>
<option value="167">Integra</option>
<option value="14196">Intel</option>
<option value="87608">Intertech</option>
<option value="10898">Inteset</option>
<option value="25174">Inves</option>
<option value="34131">Invion</option>
<option value="3217">IO Data</option>
<option value="102443">ioBox</option>
<option value="5913">IOGear</option>

```

(continues on next page)

(continued from previous page)

```

<option value="13838">Iomega</option>
<option value="25434">iOne</option>
<option value="94599">Iospirit</option>
<option value="49809">IPEX</option>
<option value="9887">I-Point</option>
<option value="17484">IPS</option>
<option value="25545">IQ</option>
<option value="153330">IRF Media</option>
<option value="2688">iRiver</option>
<option value="65942">iSonic</option>
<option value="75037">iSTAR</option>
<option value="21941">ITV Media</option>
<option value="19519">iWILL</option>
<option value="1701">IX</option>
<option value="54355">Izzy</option>
<option value="3889">J. River</option>
<option value="97564">Jadoo</option>
<option value="108373">JadooTV</option>
<option value="12368">Jaycar</option>
<option value="14514">JB Media</option>
<option value="79583">JCMatthew</option>
<option value="191593">JetStreamBox</option>
<option value="25402">Jetta</option>
<option value="7308">Jetway</option>
<option value="197297">Jide</option>
<option value="197692">JOCOKA</option>
<option value="47315">Jusst</option>
<option value="173197">Jynxbox</option>
<option value="76774">Kaiboer</option>
<option value="50055">Kaiser Baas</option>
<option value="3260">Kaleidescape</option>
<option value="4424">Kanam</option>
<option value="93629">Kartina TV</option>
<option value="95744">Kaser</option>
<option value="38829">Kazimogo</option>
<option value="163351">KdLinks</option>
<option value="187649">Keedox</option>
<option value="31735">Keelai</option>
<option value="24216">KeeLai Tech</option>
<option value="200093">KEKILO</option>
<option value="200">>Keyspan</option>
<option value="14493">Kingpin</option>
<option value="25896">Kingsun</option>
<option value="28052">KingWin</option>
<option value="4854">Kinyo</option>
<option value="16035">Kiss</option>
<option value="2387">Klegg</option>
<option value="8972">Kloss</option>
<option value="288">Kodak</option>
<option value="191591">kodi</option>
<option value="34014">Kogan</option>
<option value="25066">Komplett</option>
<option value="21590">Konig</option>
<option value="43724">KoolVu</option>
<option value="1049">Kramer</option>
<option value="24822">KSI</option>
<option value="200086">kudoTV</option>

```

(continues on next page)

(continued from previous page)

```

<option value="1656">KWorld</option>
<option value="10210">LaCie</option>
<option value="61959">LASER</option>
<option value="7575">LC Power</option>
<option value="40947">LCD Memories</option>
<option value="51079">LCD Photo</option>
<option value="2508">Leadtek</option>
<option value="52649">Ledtech</option>
<option value="196987">Leelbox</option>
<option value="38557">Lenovo</option>
<option value="167246">LETV</option>
<option value="956">Leviton</option>
<option value="39">LG</option>
<option value="45795">Lian Li</option>
<option value="8576">LifeView</option>
<option value="7483">Lindemann</option>
<option value="1470">Linksys</option>
<option value="38403">Linx</option>
<option value="2250">LiteOn</option>
<option value="945">Logitech</option>
<option value="26261">Longshine</option>
<option value="126117">LPINTE</option>
<option value="79318">m1</option>
<option value="37140">Mach Speed</option>
<option value="47718">Macom</option>
<option value="43458">Macpower</option>
<option value="106994">Macron</option>
<option value="19913">Macrosystem</option>
<option value="18">Magnavox</option>
<option value="62672">Magnetox</option>
<option value="11835">Magnex</option>
<option value="132232">Manli</option>
<option value="11763">Manta</option>
<option value="15547">Maqma</option>
<option value="12365">Marshall Electronics</option>
<option value="159002">Masora AG</option>
<option value="81074">Masscool</option>
<option value="177022">Matricom</option>
<option value="38874">Matsunichi</option>
<option value="13274">Maxdata</option>
<option value="47739">Maxdome</option>
<option value="45387">Maxian</option>
<option value="31155">MaxInPower</option>
<option value="47729">MBOX</option>
<option value="85772">MBX</option>
<option value="164">McIntosh</option>
<option value="17231">mCubed</option>
<option value="7635">MDG</option>
<option value="55209">Me2</option>
<option value="108070">Measy</option>
<option value="11085">Mecer</option>
<option value="200054">MECOOL</option>
<option value="119054">MEDE8ER</option>
<option value="9261">Media Portal</option>
<option value="27916">MediaCom</option>
<option value="22530">Mediaman</option>
<option value="16523">Mediasonic</option>

```

(continues on next page)

(continued from previous page)

```

<option value="707">Mediatech</option>
<option value="1085">Medion</option>
<option value="5589">Meedio</option>
<option value="66701">Mele</option>
<option value="182978">memoBOX</option>
<option value="26015">Memory Express</option>
<option value="27177">Memup</option>
<option value="32380">Merlin</option>
<option value="11213">MESH</option>
<option value="80895">Mesiro</option>
<option value="153246">Meteorit</option>
<option value="16276">Mevis</option>
<option value="133306">Micca</option>
<option value="11458">Micro Innovations</option>
<option value="6013">MicroByte</option>
<option value="310">Micron</option>
<option value="115">Microsoft</option>
<option value="12651">Microstar</option>
<option value="25594">Microtel</option>
<option value="53470">Midte</option>
<option value="8014">Miglia</option>
<option value="25487">Mind</option>
<option value="199805">Mini TV BOX</option>
<option value="18910">Minix</option>
<option value="17910">MiTAC</option>
<option value="176011">MK808</option>
<option value="32074">ML Arvutid</option>
<option value="19334">MM-Vision</option>
<option value="49767">Modecom</option>
<option value="13232">Monarch Computer Systems</option>
<option value="34646">Moneual</option>
<option value="141668">Monsoon Multimedia</option>
<option value="188511">MoonBox</option>
<option value="7725">Motion Computing</option>
<option value="49">Motorola</option>
<option value="74365">MPMAN</option>
<option value="24323">MPX</option>
<option value="3157">MSI</option>
<option value="42970">mStation</option>
<option value="13700">MS-Tech</option>
<option value="68417">MT Visions</option>
<option value="913">Music Mountain</option>
<option value="962">Mustek</option>
<option value="29398">Muvid</option>
<option value="15481">Mvix</option>
<option value="193428">MXQ</option>
<option value="123241">MXV</option>
<option value="80196">MyGica</option>
<option value="186226">MySku</option>
<option value="2438">MythTV</option>
<option value="213">Naim</option>
<option value="152173">Natec</option>
<option value="35236">NComputing</option>
<option value="57318">Nebula Media Solutions</option>
<option value="12">NEC</option>
<option value="140705">Neewer</option>
<option value="27900">NeoDigits</option>

```

(continues on next page)

(continued from previous page)

```

<option value="51108">Neonumeric</option>
<option value="15187">Net Digital</option>
<option value="4018">Netgear</option>
<option value="36687">Netkia</option>
<option value="26435">Network</option>
<option value="17214">Neuros</option>
<option value="3563">Neuston</option>
<option value="46115">Neutron</option>
<option value="57261">Newron</option>
<option value="198510">NEXBOX</option>
<option value="201113">NEXSMART</option>
<option value="28577">NHJ</option>
<option value="14658">NiceTracker</option>
<option value="5599">Niveus Media</option>
<option value="137839">Nixeus</option>
<option value="17020">NMediaPC</option>
<option value="29954">Noontec</option>
<option value="112638">Normus</option>
<option value="55514">Northern Micro</option>
<option value="41583">Novac</option>
<option value="82581">Novatech</option>
<option value="20393">Novatron</option>
<option value="31606">Novita</option>
<option value="66565">NOX</option>
<option value="51663">NZXT</option>
<option value="86564">O2Media</option>
<option value="50747">O3B</option>
<option value="44588">Odsonic</option>
<option value="17959">ODYS</option>
<option value="37261">Okoro</option>
<option value="35961">Okoro Media Systems</option>
<option value="12432">Olidata</option>
<option value="201177">OMG FREETV</option>
<option value="9193">OnCinema</option>
<option value="40">Onkyo</option>
<option value="13847">Open Source Factory</option>
<option value="156846">openelec</option>
<option value="192538">OpenHour</option>
<option value="13058">Optima</option>
<option value="38">Optimus</option>
<option value="18215">Orange</option>
<option value="188862">Orbsmart</option>
<option value="35204">Ordi</option>
<option value="13986">Origen AE Technology</option>
<option value="423">Orion</option>
<option value="69839">ORtek</option>
<option value="194279">OSMC</option>
<option value="189948">OTT TV BOX</option>
<option value="18619">Overdrive PC</option>
<option value="197403">ovomedia</option>
<option value="41658">PAC</option>
<option value="4067">Packard Bell</option>
<option value="13152">Palmbutler</option>
<option value="13">Panasonic</option>
<option value="35848">PanDigital</option>
<option value="51446">Pantiac</option>
<option value="9082">Paradigit</option>

```

(continues on next page)

(continued from previous page)

```
<option value="3595">Parex</option>
<option value="53562">Patriot</option>
<option value="14764">PC Mak</option>
<option value="11551">PC Plaza</option>
<option value="12465">PC Sleuth</option>
<option value="34201">Pcalchemy</option>
<option value="17547">PCCity</option>
<option value="13406">PCzapper</option>
<option value="12112">Peach Tron Systems</option>
<option value="43337">Pearl</option>
<option value="3785">Peekton</option>
<option value="87179">PeerTV</option>
<option value="14981">Penguin Computing</option>
<option value="191786">PeQ</option>
<option value="51900">Perfection</option>
<option value="27">Philips</option>
<option value="53661">PINE</option>
<option value="1205">Pinnacle</option>
<option value="4">Pioneer</option>
<option value="87995">Pipi</option>
<option value="67727">Pirelli</option>
<option value="160980">Pivos</option>
<option value="8916">Pixel Magic Systems</option>
<option value="8133">PixelView</option>
<option value="19217">Platinum</option>
<option value="38108">Playtime</option>
<option value="130339">Playtronics</option>
<option value="87601">Plex</option>
<option value="8577">Plextor</option>
<option value="506">Polaroid</option>
<option value="3367">Polycom</option>
<option value="11243">Polywell</option>
<option value="10699">POMi</option>
<option value="136750">PopBox</option>
<option value="98519">Popcorn Hour</option>
<option value="105572">Popstar</option>
<option value="39701">Portable USA</option>
<option value="12275">Positivo</option>
<option value="94355">Power Zest</option>
<option value="12119">PowerColor</option>
<option value="8678">PowerSpec</option>
<option value="38912">PQI</option>
<option value="76415">PremiumBlue</option>
<option value="10700">Prestigio</option>
<option value="24373">PrimeDTV</option>
<option value="15623">Princeton</option>
<option value="117103">Probox</option>
<option value="191728">Probox2</option>
<option value="124073">Prodigi</option>
<option value="1134">ProSAT</option>
<option value="23032">Psychone</option>
<option value="32606">Pyrogate</option>
<option value="200336">QcoQce</option>
<option value="24678">QCS</option>
<option value="192892">Qianxun</option>
<option value="61834">QMedia</option>
<option value="19527">Q-Motion</option>
```

(continues on next page)

(continued from previous page)

```
<option value="94246">Qnap</option>
<option value="6089">QPVision</option>
<option value="14852">Quality PC</option>
<option value="15741">Quartek</option>
<option value="17178">Quibus</option>
<option value="198550">Quick Play</option>
<option value="4246">Quixun</option>
<option value="95603">Radiogears</option>
<option value="443">RadioShack</option>
<option value="58220">RaidSonic</option>
<option value="66209">Rapsody</option>
<option value="172343">Raspberry</option>
<option value="7">RCA</option>
<option value="80763">RCKING</option>
<option value="65951">Red Sky Global</option>
<option value="12347">Redbell</option>
<option value="12378">Reel Multimedia</option>
<option value="14706">Remotec</option>
<option value="28132">Replay Plus</option>
<option value="371">ReQuest</option>
<option value="192140">revez</option>
<option value="179">Revox</option>
<option value="14896">Reycom</option>
<option value="14412">RicaVision</option>
<option value="172678">Rikomagic</option>
<option value="8325">Rimax</option>
<option value="16278">Riscom</option>
<option value="77146">Ritmo</option>
<option value="37418">Rivertech</option>
<option value="28566">Rixid</option>
<option value="47796">Rock</option>
<option value="165884">Rockchip</option>
<option value="199853">Rometak</option>
<option value="117528">Rosen Aviation</option>
<option value="57775">Rosewill</option>
<option value="90031">Roxcore</option>
<option value="15049">Royal</option>
<option value="14308">R-Style</option>
<option value="62551">RTI</option>
<option value="147">Russound</option>
<option value="12726">S1Digital</option>
<option value="12802">SACA Technologies</option>
<option value="8864">Sage</option>
<option value="12183">SageTV</option>
<option value="2086">Salora</option>
<option value="118935">Sama</option>
<option value="2">Samsung</option>
<option value="44128">San Hawk</option>
<option value="1959">SanDisk</option>
<option value="17052">Sansun</option>
<option value="11253">Sapphire</option>
<option value="11908">Sarotech</option>
<option value="12323">Savit Micro</option>
<option value="10717">Sceneo</option>
<option value="190242">SCISHION</option>
<option value="26485">Sdtec</option>
<option value="72767">Seagate</option>
```

(continues on next page)

(continued from previous page)

```

<option value="10744">Seanix</option>
<option value="5910">Sencor</option>
<option value="19890">Sharkoon</option>
<option value="8">Sharp</option>
<option value="52509">Shenzhen Newland</option>
<option value="53887">Shenzhen Soma</option>
<option value="48241">SHG</option>
<option value="104362">Shintaro</option>
<option value="797">Shuttle</option>
<option value="13284">Sicuro</option>
<option value="3306">Siemens</option>
<option value="2844">Sigma Designs</option>
<option value="1502">Sigmatek</option>
<option value="25501">SigmaTel</option>
<option value="12578">Silentmaxx</option>
<option value="8960">SilverCrest</option>
<option value="12330">Silverleaf</option>
<option value="8123">Silverstone</option>
<option value="61342">SIMEREC</option>
<option value="13022">Sitecom</option>
<option value="27874">SkipJam</option>
<option value="12553">Sky</option>
<option value="61553">Skystar</option>
<option value="191494">Skystream</option>
<option value="3519">SkyWorth</option>
<option value="6496">SleekLine</option>
<option value="42268">Sling Media</option>
<option value="4613">SmartDisk</option>
<option value="28148">SmartParts</option>
<option value="190095">SmartThings</option>
<option value="1135">SMC</option>
<option value="1300">SnapStream</option>
<option value="10166">SnaZio</option>
<option value="43395">Snogard</option>
<option value="46183">Solid Year</option>
<option value="165345">SolidRun</option>
<option value="23706">Solidtek</option>
<option value="16103">Soltek</option>
<option value="10776">Sonavis</option>
<option value="42804">Sonbook</option>
<option value="30444">Sonic Impact</option>
<option value="1">Sony</option>
<option value="17599">SorensonVRS</option>
<option value="10294">SoundGraph</option>
<option value="12287">SOYO</option>
<option value="12255">Spectra</option>
<option value="81773">Speed</option>
<option value="17607">Speed-Link</option>
<option value="14620">ST Lab</option>
<option value="1650">Stands Unique</option>
<option value="12523">Stein</option>
<option value="37926">STiNO</option>
<option value="15776">Storex</option>
<option value="159725">Streacom</option>
<option value="200967">STREAM TEAM MEDIA</option>
<option value="196331">STREAMAXTV</option>
<option value="193260">StreamSmart</option>

```

(continues on next page)

(continued from previous page)

```
<option value="1071">Streamzap</option>
<option value="47851">Sumicom</option>
<option value="7188">Sumvision</option>
<option value="25750">Sun Microsystems</option>
<option value="198096">Sunvell</option>
<option value="50097">SuperMicro</option>
<option value="50009">Superna</option>
<option value="53413">SuperPower</option>
<option value="116962">Sveon</option>
<option value="34182">Sweex</option>
<option value="10362">Syabas</option>
<option value="20578">Syba</option>
<option value="61311">Synology</option>
<option value="4171">Syntax-Brilliant</option>
<option value="35636">Sysmaster</option>
<option value="9717">Systemax</option>
<option value="94196">syvio</option>
<option value="4923">Tagar Systems</option>
<option value="2056">Tandberg</option>
<option value="199761">Tanix</option>
<option value="6337">Targa</option>
<option value="23631">Target</option>
<option value="21874">Tarox</option>
<option value="13073">TCI</option>
<option value="50571">TCR</option>
<option value="36">Teac</option>
<option value="67282">TechGear</option>
<option value="4693">Technaxx</option>
<option value="7251">Technika</option>
<option value="1329">TechniSat</option>
<option value="7371">TechnoTrend</option>
<option value="18681">Techsolo</option>
<option value="132344">TechStudioTV</option>
<option value="75571">Tekkeon</option>
<option value="28819">Tekram</option>
<option value="31399">Telecom Italia</option>
<option value="15495">Telekom</option>
<option value="12075">Tempest Microsystems</option>
<option value="1118">Terratec</option>
<option value="5825">Teufel</option>
<option value="91367">TeVii</option>
<option value="5808">Texas Instruments</option>
<option value="36483">Texet</option>
<option value="7531">Textorm</option>
<option value="184224">The Little Black Box</option>
<option value="13599">TheBox</option>
<option value="85794">Thecus</option>
<option value="16114">Thermaltake</option>
<option value="126">Thomson</option>
<option value="196639">TICTID</option>
<option value="4842">Tilgin</option>
<option value="249">TiVo</option>
<option value="163223">TizzBird</option>
<option value="52301">Tomacro</option>
<option value="196411">Tonbux</option>
<option value="62865">Topke</option>
<option value="124655">Topseed</option>
```

(continues on next page)

(continued from previous page)

```

<option value="199799">Topsion</option>
<option value="9883">Torspoal</option>
<option value="5">Toshiba</option>
<option value="9101">Touch Systems</option>
<option value="36195">Tranquil PC</option>
<option value="11883">Transcend</option>
<option value="8961">Traxdata</option>
<option value="48659">Trekstor</option>
<option value="38804">Tricod</option>
<option value="31646">Triline</option>
<option value="21864">Trinix</option>
<option value="192589">Tronfy</option>
<option value="199331">Trongle</option>
<option value="179524">Tronsmart</option>
<option value="9246">Trust</option>
<option value="22371">Tsunami Dream</option>
<option value="23580">Tundra</option>
<option value="55707">Turbo-X</option>
<option value="191819">TVB Anywhere</option>
<option value="37199">TVonics</option>
<option value="4293">TwinHan</option>
<option value="76981">Twisted Melon</option>
<option value="41737">TwonkyMedia</option>
<option value="10871">Tyan</option>
<option value="9554">Typhoon</option>
<option value="186933">U2C</option>
<option value="166823">UBOX</option>
<option value="139766">Uebo</option>
<option value="175913">UGOOS</option>
<option value="7534">UIRT</option>
<option value="13348">Umax</option>
<option value="57978">Umbra</option>
<option value="192076">Unblock Tech</option>
<option value="12940">Uneed</option>
<option value="6532">UniBrain</option>
<option value="14204">Unicorn</option>
<option value="12751">Unika</option>
<option value="12504">Uvem</option>
<option value="32264">Vadim</option>
<option value="291">Vantage</option>
<option value="31405">VEI</option>
<option value="19091">Velocity Micro</option>
<option value="162528">VENZ</option>
<option value="89397">Verbatim</option>
<option value="44280">Verkkokauppa</option>
<option value="14642">Via</option>
<option value="34618">Vibe</option>
<option value="41761">Vidabox</option>
<option value="193849">Videostrong</option>
<option value="179094">VidOn.me</option>
<option value="217">ViewSonic</option>
<option value="192387">VIGICA</option>
<option value="29890">Viglen</option>
<option value="11851">Viscom</option>
<option value="14590">VisionPlus</option>
<option value="1156">Vizio</option>
<option value="21634">VLSystem</option>

```

(continues on next page)

(continued from previous page)

```
<option value="21512">Vobis</option>
<option value="184069">VONHAUS</option>
<option value="12218">Voodoo PC</option>
<option value="8463">VPR Matrix</option>
<option value="54604">Vudu, Inc.</option>
<option value="43858">We Digital</option>
<option value="16573">Webster</option>
<option value="19795">Welland</option>
<option value="7748">Wellton Way</option>
<option value="73793">Western Digital</option>
<option value="193008">WeTek</option>
<option value="12161">White Box Computers</option>
<option value="4020">WinBook</option>
<option value="51240">Wincomm</option>
<option value="8381">Woxter</option>
<option value="118559">Wyplayer</option>
<option value="56670">Wyse</option>
<option value="58">X10 Wireless</option>
<option value="15111">X4-Tech</option>
<option value="118279">Xenta</option>
<option value="195958">XGODY</option>
<option value="181054">Xiaomi</option>
<option value="28767">Ximeta</option>
<option value="4499">Xitel</option>
<option value="49736">X-ONE</option>
<option value="5100">Xoro</option>
<option value="58753">Xparity</option>
<option value="113618">Xtreameer</option>
<option value="3432">Xtreme</option>
<option value="5499">Yakumo</option>
<option value="25">Yamaha</option>
<option value="196901">Yixu</option>
<option value="200561">YokaTV</option>
<option value="25205">Yuan</option>
<option value="197047">Yuntab</option>
<option value="131996">YuppTV</option>
<option value="23925">Zaapa</option>
<option value="136988">ZAAPT</option>
<option value="17847">Zalman</option>
<option value="179654">Zappiti</option>
<option value="32609">Zelfbouw</option>
<option value="22010">Zepto</option>
<option value="194333">Zidoo</option>
<option value="17658">Ziga</option>
<option value="1000">Zinwell</option>
<option value="31936">Zio</option>
<option value="23594">Zioncom</option>
<option value="43123">Ziova</option>
<option value="15555">Zitech</option>
<option value="186475">Zoomtak</option>
<option value="110163">Zotac</option>
<option value="6054">ZT Group</option>
<option value="28514">ZyXEL</option></select>
```


24.2 Install Concordance

Install the `concordance` tool.

I'm using v1.0 on Ubuntu 15.10 64-bit.

Briefly:

```
sudo apt-get install libusb-dev libzip-dev
wget http://downloads.sourceforge.net/project/concordance/concordance/1.0/concordance-
→1.0.tar.bz2
tar xjf concordance-1.0.tar.bz2
cd concordance-1.0/libconcord
./configure
make
sudo make install
sudo ldconfig
cd ../concordance
./configure
make
sudo make install
```

24.3 Arrange to run without needing sudo

- Run 'lsusb' to see what devices are already attached to your computer.
- Plug in your remote
- Run 'lsusb', looking for the device that wasn't there before. E.g. my Harmony One produced this:

```
Bus 001 Device 021: ID 046d:c121 Logitech, Inc.
```

- Unplug the remote
- Now (using sudo as needed) create a new file, `/etc/udev/rules.d/custom-concordance.rules`, substituting in the values for your own remote:

```
ATTRS{idVendor}=="046d", ATTRS{idProduct}=="c121", MODE="666"
```

- Test
 - Plug the remote in
 - Run "concordance -i". It should print out information about the attached remote

24.4 Programming the remote

- Make a directory to hold firmware and config backups, e.g. `~/Documents/LogitechHarmonyOne`.
- `cd ~/Documents/LogitechHarmonyOne`
- Backup firmware if you haven't previously:

```
concordance --dump-firmware <filename> (default: firmware.EZUp)
```

- Backup config if you haven't previously:

```
concordance --dump-config <filename> (default: config.EZHex)
```

- Go to <http://members.harmonyremote.com/EasyZapper/> and log in (ignore the message about upgrading your software). If you don't already have an account, create one.
- Update your remote's configuration using the web site.
- When ready to update your remote:
 - Click "Update your remote". It'll prompt to connect your remote.
 - Connect the remote to the computer via USB.
 - Click "Next" on the web page.
 - Your browser will download a file named "Connectivity.EZHex". Save it to your LogitechHarmonyOne directory.
 - Go to your shell.
 - Run:

```
concordance Connectivity.EZHex
```

- Go back to your browser.
- Click "Next".
- Wait... could be several minutes.
- Eventually your browser will download a file named "Update.EZHex". Save it to your LogitechHarmonyOne directory.
- Go to your shell.
- Run this command. This will take quite a while (5 minutes?), but it will print progress status as it goes:

```
concordance Update.EZHex
```

- When that's done, disconnect your remote and try it out.

25.1 Fixing on Ubuntu

On Ubuntu, mpd is installed broken. You'll have to go read [this page](#) and do some work to fix things so mpd will actually work on Ubuntu. Sigh.

Important: To play *anything*, you must first get it into the current playlist. Even if you just want to play one thing.

25.2 Playlist commands

Commands that change the playlist:

clear Empty the playlist (stops playing if anything was playing)

crop Delete all playlist entries *except* the one currently playing

del <number> Delete one item from the playlist; 0 means the currently playing item, otherwise items are numbered starting at 1. (To see playlist with numbers, try `mpc playlist | cat -n`. There's probably a better way.)

add <file> Add an item from the database to the current playlist, at the end. <file> should be a path to the item, as shown by "mpc ls".

insert <file> Like add, only it puts the item immediately after the currently playing item, so it'll play next. If random mode is enabled, it'll still be played next.

mv|move <from> <to> Move item at position <from> to be at position <to>.

save <name> Save current playlist as database playlist with name <name>.

load <name> Add contents of database playlist named <name> to the current playlist.

rm <name> Delete database playlist named <name> from database.

25.3 Playing things

Status:

playlist [-f <format>] List songs in current playlist. See “man mpc” for format string syntax.

current Show what's playing right now

Control:

play [<position>] Start playing the item at <position> in the playlist. Default is number 1.

pause Pause playing

stop Stop playing

next Stop playing the current entry from the current playlist and start playing the next one.

prev Reverse of next.

25.4 Playing a playlist from the music database with mpc

Suppose you have a playlist in the database already - e.g. a file “/var/lib/mpd/playlists/WCPE.m3u” that you've created earlier.

Now you want to play that playlist.

```
$ mpc clear
$ mpc lsplaylists
WUNC
WCPE
Favorites
$ mpc load WCPE
loading: WCPE
$ mpc play
volume: 96%   repeat: on    random: off   single: off   consume: off
loading: WCPE
http://audio-ogg.ibiblio.org:8000/wcpe.ogg
[playing] #1/1   0:00/0:00 (0%)
volume: 96%   repeat: on    random: off   single: off   consume: off
```

MySQL with Django

Ubuntu: need to install (to use MySQL with Django):

```
sudo apt-get install mysql-client mysql-server libmysqlclient-dev
```

Django:

```
pip install mysqlclient

DATABASES['default'] = {
    'ENGINE': 'django.db.backends.mysql',
    'NAME': 'dbname',
    'USER': 'username',
}
```

26.1 Using the client

Starting the client:

```
$ mysql --user=username [database]      # if user has no password
$ mysql --user=username --password [database]  # to be prompted for password
```

To do things that require mysql root:

```
$ mysql -u root      # If root has no password and older Debian
$ mysql -u root -p   # if root has password and older Debian
$ sudo mysql -u root # On more recent Debian, no need for root password but must be
↳root user
```

26.2 Users and permissions

In the client:

```
mysql> SELECT user, host from mysql.user;                                # List existing_
↪users
mysql> CREATE USER 'username' IDENTIFIED BY 'plaintextpassword';          # Create user_
↪with password
mysql> CREATE USER 'username'@'localhost';    # no password, can only connect locally
mysql> SHOW DATABASES;
mysql> CREATE DATABASE databasename;
mysql> GRANT ALL ON databasename.* TO "username"@"hostname" IDENTIFIED BY "password";
mysql> FLUSH PRIVILEGES;
mysql> DROP DATABASE databasename;
mysql> DROP USER username;
mysql> EXIT
Bye
```

26.3 Change user password

Note: default host is ‘%’ which will not let you connect via unix socket, must set password for host ‘localhost’ to allow that:

```
mysql> update mysql.user set password=password('foo'),host='localhost' where user=
↪'poirier_wordpres';    # On older MySQL
mysql> set password for 'dpoirier'@'localhost' = 'plainpass';    # More recent MySQL
mysql> flush privileges;
```

26.4 Recover lost password

<http://dev.mysql.com/doc/refman/5.5/en/resetting-permissions.html>

C.5.4.1.3. Resetting the Root Password: Generic Instructions On any platform, you can set the new password using the mysql client:

```
Stop mysqld
Restart it with the --skip-grant-tables option. This enables anyone to connect_
↪without a password and with all privileges. Because this is insecure, you might_
↪want to use --skip-grant-tables in conjunction with --skip-networking to prevent_
↪remote clients from connecting.

$ mysql
mysql> UPDATE mysql.user SET Password=PASSWORD('MyNewPass') WHERE User='root';
mysql> FLUSH PRIVILEGES;
mysql> EXIT

Stop the server
Restart it normally (without the --skip-grant-tables and --skip-networking options).
```

26.5 Dumps

Make a dump:

```
mysqldump --single-transaction _dbname_ > dumpfile.sql  
mysqldump --result-file=dumpfile.sql --single-transaction _dbname_
```

(Use `--single-transaction` to avoid locking the DB during the dump.)

Restore a dump:

```
mysql dbname < dumpfile.sql
```

26.6 Create a new MySQL database

Step by step:

```
$ mysql -u root -p  
<ENTER MYSQL ROOT PASSWORD>  
mysql> create user 'ctsv2_TR'@'localhost';  
mysql> create database ctsv2_TR;  
mysql> grant all on ctsv2_TR.* to 'ctsv2_TR'@'localhost';  
mysql> flush privileges;  
mysql> exit  
Bye
```


[Nginx docs](#) are here but good luck finding anything there if you don't already where it is.

27.1 Redirect non-SSL to SSL

From <https://serverfault.com/questions/250476/how-to-force-or-redirect-to-ssl-in-nginx>:

```
server {
    listen 80;
    listen [::]:80;
    listen 443 default_server ssl;

    server_name www.example.com;

    ssl_certificate /path/to/my/cert;
    ssl_certificate_key /path/to/my/key;

    if ($scheme = http) {
        return 301 https://$server_name$request_uri;
    }
}
```

27.2 Most useful variables

\$host in this order of precedence: host name from the request line, or host name from the “Host” request header field, or the server name matching a request

\$http_host Value of the “Host:” header in the request (same as all \$http_<headername> variables)

\$https “on” if connection operates in SSL mode, or an empty string otherwise

\$request_method request method, usually “GET” or “POST”

\$request_uri full original request URI (with arguments)

\$scheme request scheme, e.g. "http" or "https"

\$server_name name of the server which accepted a request

\$server_port port of the server which accepted a request

27.3 Variables in configuration files

See above for "variables" that get set automatically for each request (and that we cannot modify).

The ability to set variables at runtime and control logic flow based on them is part of the [rewrite module](#) and *not* a general feature of nginx.

You can [set](#) a variable:

```
Syntax:      set $variable value;
Default:     --
Context:     server, location, if
```

"The value can contain text, variables, and their combination." – but I have not yet found the documentation on how these can be "combined".

Then use [if](#) etc.:

```
Syntax:      if (condition) { rewrite directives... }
Default:     --
Context:     server, location
```

Conditions can include:

```
* a variable name; false if the value of a variable is an empty string or "0";
* comparison of a variable with a string using the "=" and "!=" operators;
* matching of a variable against a regular expression using the "~" (for case-
  ↳sensitive matching) and "~*" (for case-insensitive matching) operators. Regular
  ↳expressions can contain captures that are made available for later reuse in the $1..
  ↳$9 variables. Negative operators "!~" and "!~*" are also available. If a regular
  ↳expression includes the "\" or ";" characters, the whole expressions should be
  ↳enclosed in single or double quotes.
* checking of a file existence with the "-f" and "!-f" operators;
* checking of a directory existence with the "-d" and "!-d" operators;
* checking of a file, directory, or symbolic link existence with the "-e" and "!-e"
  ↳operators;
* checking for an executable file with the "-x" and "!-x" operators.
```

Examples:

```
if ($http_user_agent ~ MSIE) {
    rewrite ^(.*)$ /msie/$1 break;
}

if ($http_cookie ~* "id=([^;]+)(?:;|$)") {
    set $id $1;
}

if ($request_method = POST) {
    return 405;
}
```

(continues on next page)

(continued from previous page)

```
}  
  
if ($slow) {  
    limit_rate 10k;  
}  
  
if ($invalid_referer) {  
    return 403;  
}
```

Warning: You *CANNOT* put any directive you want inside the `if`, only rewrite directives like `set`, `rewrite`, `return`, etc.

Warning: The values of variables you set this way can *ONLY* be used in `if` conditions, and maybe rewrite directives; don't try to use them elsewhere.

CHAPTER 28

NPM

Make *npm install* less noisy:

```
npm config set loglevel warn
```

or add this to `~/.npmrc`:

```
loglevel=warn
```

[source](#).

Some of this from <http://www.coresecuritypatterns.com/blogs/?p=763> and <http://www.bogpeople.com/networking/openssl.shtml>.

29.1 End-user Functions

29.1.1 Create key

Create a 2048-bit key pair:

```
openssl genrsa 2048 > myRSA-key.pem
openssl genrsa -out blah.key.pem
openssl genrsa -out blah.key.pem 2048
```

Create a password-protected 2048-bit key pair:

```
openssl genrsa 2048 -aes256 -out myRSA-key.pem
```

OpenSSL will prompt for the password to use. Algorithms: AES (aes128, aes192, aes256), DES/3DES (des, des3).

Remove passphrase from a key:

```
openssl rsa -in server.key -out server-without-passphrase.key
```

Extract public key:

```
openssl rsa -in blah.key.pem -out public.key -pubout
```

29.1.2 Getting Certificates

Create Certificate Request and Unsigned Key:

```
openssl req -nodes -new -keyout blah.key.pem -out blah.csr.pem
```

More thorough example:

```
openssl req -new rsa:1024 -node -out myCSR.pem \  
-keyout myPrivCertkey.pem \  
-subj "/C=US/ST=MA/L=Burlington/CN=myHost.domain.com/emailAddress=user@example.com"  
↪ "
```

Create a self-signed certificate:

```
openssl req -nodes -x509 -newkey rsa:1024 -days 365 \  
-out mySelfSignedCert.pem -set_serial 01 \  
-keyout myPrivServerKey.pem \  
-subj "/C=US/ST=MA/L=Burlington/CN=myHost.domain.com/emailAddress=user@example.com"  
↪ "
```

`-x509` identifies it as a self-signed certificate and `-set_serial` sets the serial number for the server certificate.

Create a single file that contains both private key and the self-signed certificate:

```
openssl req -x509 -nodes -days 365 -newkey rsa:1024 \  
-keyout myServerCert.pem -out myServerCert.pem \  
-subj "/C=US/ST=MA/L=Burlington/CN=myHost.domain.com/emailAddress=user@example.com"  
↪ "
```

Fingerprint for Unsigned Certificate:

```
openssl x509 -subject -dates -fingerprint -in blah.key.pem
```

Display Certificate Information:

```
openssl x509 -in blah.crt.pem -noout -text
```

Creating a PEM File for Servers:

```
cat blah.key.pem blah.crt.pem blah.dhp.pem > blah.pem
```

Download some server's certificate:

```
openssl s_client -connect www.example.com:443
```

(then hit `^C` out of the interactive shell)

29.1.3 Viewing Certificate Contents

X.509 certificates are usually stored in one of two formats. Most applications understand one or the other, some understand both:

- DER which is raw binary data.
- PEM which is a text-encoded format based on the Privacy-Enhanced Mail standard (see RFC1421). PEM-format certificates look something like this:

```
-----BEGIN CERTIFICATE-----  
MIIBrjCCAWwCAQswCQYFKw4DAhsFADBTMQswCQYDVQQGEwJBVTETMBEGA1UECBMK
```

(continues on next page)

(continued from previous page)

```
U29tZS1TdGF0ZTEhMB8GA1UEChMYSW50ZXJuZXQgV2lkZ2l0cyBQdHkgTHRkMQww
:
:
MQAwLgIvAJ4wtQsANPxHo7Q4IQZYsL12SKdbAhUAjJ9n38zxT+iai2l64xS+LIfa
C1Q=
-----END CERTIFICATE-----
```

OpenSSL uses the PEM `format` by default, but you can tell it to process DER `format`.
 ↳certificates...you just need to know which you are dealing `with`.

The command to view an X.509 certificate is:

```
openssl x509 -in filename.cer -inform der -text
```

You can specify `-inform pem` if you want to look at a PEM-format certificate

29.1.4 Convert Between Formats

If you have a PEM-format certificate which you want to convert into DER-format, you can use the command:

```
openssl x509 -in filename.pem -inform pem -out filename.cer -outform der
```

29.1.5 PKCS12 files

PKCS12 files are a standard way of storing multiple keys and certificates in a single file. Think of it like a zip file for keys & certificates, which includes options to password protect etc.

Don't worry about this unless you need it because some application requires a PKCS12 file or you're given one that you need to get stuff out of.

Viewing PKCS12 Keystore Contents:

```
openssl pkcs12 -in filename.p12 -info
```

If you have two separate files containing your certificate and private key, both in PEM format, you can combine these into a single PKCS12 file using the command:

```
openssl pkcs12 -in cert.pem -inkey key.pem -export -out filename.p12
```

29.1.6 Encrypting and signing things

Signing E-mails:

```
openssl smime -sign -in msg.txt -text -out msg.encrypted -signer blah.crt.pem -inkey_
↳blah.key.pem
```

Sign some text:

```
openssl dgst -sign private.key -out signature.asc
```

Verify signature:

```
if openssl dgst -verify public.key -signature signature.asc ; then echo GOOD; else_  
↪echo BAD; fi
```

Encrypt and decrypt a single file:

```
openssl aes-128-cbc -salt -in file -out file.aes  
openssl aes-128-cbc -d -salt -in file.aes -out file
```

Simple file encryption:

```
openssl enc -bf -A -in file_to_encrypt.txt
```

(password will be prompted)

Simple file decryption:

```
openssl enc -bf -d -A -in file_to_encrypt.txt
```

tar and encrypt a whole directory:

```
tar -cf - directory | openssl aes-128-cbc -salt -out directory.tar.aes  
openssl aes-128-cbc -d -salt -in directory.tar.aes | tar -x
```

tar zip and encrypt a whole directory:

```
tar -zcf - directory | openssl aes-128-cbc -salt -out directory.tgz.aes  
openssl aes-128-cbc -d -salt -in directory.tgz.aes | tar -xz
```

29.2 Certificate Authority Functions

When setting up a new CA on a system, make sure index.txt and serial exist (empty and set to 01, respectively), and create directories private and newcert.

Edit openssl.cnf - change default_days, certificate and private_key, possibly key size (1024, 1280, 1536, 2048) to whatever is desired.

Create CA Certificate:

```
openssl req -new -x509 -keyout private/something-CA.key.pem \  
-out ./something-CA.crt.pem -days 3650
```

Export CA Certificate in DER Format:

```
openssl x509 -in something-CA.crt.pem -outform der \  
-out something-CA.crt
```

Revoke Certificate:

```
openssl ca -revoke blah.crt.pem
```

Generate Certificate Revocation List:

```
openssl ca -gencrl -out crl/example.com-CA.crl
```

Sign Certificate Request:

```
openssl ca -out blah.crt.pem -in blah.req.pem
```

Create Diffie-Hoffman Parameters for Current CA:

```
openssl dhparam -out example.com-CA.dhp.pem 1536
```

Creating Self-Signed Certificate from Generated Key:

```
openssl req -new -x509 -key blah.key.pem -out blah.crt.pem
```

Use only when you've no CA and will only be generating one key/certificate (useless for anything that requires signed certificates on both ends)

CHAPTER 30

Org mode (Emacs)

<http://orgmode.org/org.html> See also <http://orgmode.org/orgcard.pdf>

Binding	Operation
<i>Tasks</i>	
M-S-Ret	Add another item at same level
C-c C-t	Change TODO state
C-c / t	Show only uncompleted todos
<i>Agenda</i>	
C-c C-a n	View schedule and unscheduled tasks
b	Move backward (previous day)
f	Move forward (next day)
<i>Scheduling</i>	
C-c C-s	Schedule a task (set a date and optional time to do it)
C-u C-c C-s	Unschedule a task (remove schedule date/time)

Keys outside org-mode:

Key	What
C-c g	my gtd file
C-c c t	Create task
C-c a X	Agenda view X
C-c l	org-store-link
C-c c	org-capture
C-c b	org-iswitchb (?)

Keys in org-mode file:

C-c C-x p	org-set-property
M-Return	org-meta- return - start new line with new heading at same level
M-S-right arrow	move current heading one deeper

(continues on next page)

(continued from previous page)

```
C-c C-s      schedule task
C-c C-d      set task deadline
C-c C-q      org-set-tags-command add task tag - USE FOR CONTEXT
C-c / d      org-check-deadlines - spared tree with deadlines that are past-due or_
→soon to be

<TAB>        org-cycle
S-<TAB>       org-global-cycle
```

Keys in agenda views:

TBD

CHAPTER 31

Postfix

Retry sending all queued mail:

```
postfix flush
```

Delete all queued mail:

```
postsuper -d ALL deferred
```


32.1 Snippets

- In psql:

```
# \d *pattern*
```

Display definition of table, view, or other relations with name matching the pattern

See <http://jacobian.org/writing/pg-encoding-ubuntu/> to fix postgres to default to UTF-8 on ancient Ubuntu. (Destroys data.)

To set up your user on ubuntu to auth to local postgres automatically <https://help.ubuntu.com/community/PostgreSQL>

Check replication:

```
ON master, select * from pg_stat_replication;  
(from http://www.dansketcher.com/2013/01/27/monitoring-postgresql-streaming-replication/)
```

32.2 Environment Variables

The following environment variables can be used to select default connection parameter values, which will be used if no value is directly specified. These are useful to avoid hard-coding database connection information into simple client applications, for example.

PGHOST behaves the same as host connection parameter.

PGHOSTADDR behaves the same as hostaddr connection parameter. This can be set instead of or in addition to PGHOST to avoid DNS lookup overhead.

PGPORT behaves the same as port connection parameter.

PGDATABASE behaves the same as dbname connection parameter.

PGUSER behaves the same as user connection parameter.

PGPASSWORD behaves the same as password connection parameter. Use of this environment variable is not recommended for security reasons (some operating systems allow non-root users to see process environment variables via ps); instead consider using the ~/.pgpass file (see Section 30.14).

PGPASSFILE specifies the name of the password file to use for lookups. If not set, it defaults to ~/.pgpass (see Section 30.14).

32.3 Dump Postgres table to a .CSV file

Started with this: <http://stackoverflow.com/questions/1120109/export-postgres-table-to-csv-file-with-headings>

Using COPY requires superuser but the error message helpfully tells you that you can use copy instead :-)

Using cactus' django template, something like:

```
$ fab -udpoirier production manage_run:dbshell
[huellavaliente.com:2222] out: venezuelasms_production=> \copy message_log_message to
↳ '/tmp/messages.csv' csv header
[huellavaliente.com:2222] out: venezuelasms_production=> \q

$ sftp -o Port=2222 dpoirier@huellavaliente.com
Connected to huellavaliente.com.
sftp> cd /tmp
sftp> ls
messages.csv
sftp> get messages.csv
Fetching /tmp/messages.csv to messages.csv
/tmp/messages.csv
↳                               100% 1776KB 888.0KB/s   00:02
sftp> ^D
```

32.4 Postgres with non-privileged users

How do we do things on Postgres without giving superuser to the user that actually uses the database every day? The following assumes a Postgres superuser named 'master'. (Or the RDS 'master' user, who has most superuser privileges.)

In the examples below, for readability I'm omitting most of the common arguments to specify where the postgres server is, what the database name is, etc. You can set some environment variables to use as defaults for things:

```
export PGDATABASE=dbname
export PGHOST=xxxxxxxxx
export PGUSER=master
export PGPASSWORD=xxxxxxxxx
```

32.5 Create user

This is pretty standard. To create user \$username with plain text password \$password:

```
export PGUSER=master
export PGDATABASE=postgres
createuser -DERS $username
psql -c "ALTER USER $username WITH PASSWORD '$password';"
```

Yes, none of the options in `-DERS` are strictly required, but if you don't mention them explicitly, `createuser` asks you about them one at a time.

If not on RDS, for the user to actually do something useful like connect to postgres, you might also have to edit `pg_hba.conf` and add a line like:

```
local    <dbname>    <rolename>                                md5
```

to let it connect using `host=` (unix domain socket) and provide a password to access `<dbname>`. You could also put "all" there to let it access any password it otherwise has auth for. E.g. to allow local connections via both unix socket and tcp connections to localhost:

```
local    all          all                                     md5
host     all          all          127.0.0.1/32             md5
```

32.6 Create database

If you need a database owned by `$project_user`, you can:

- Create it as `$project_user` if that user has `CREATEDB`:

```
export PGUSER=$project_user
createdb --template=template0 $dbname
```

- Create it as a superuser and specify that the owner should be `$project_user`:

```
export PGUSER=postgres
createdb --template=template0 --owner=$project_user $dbname
```

- Create it as any other user, so long as the other user is a member, direct or indirect, of the `$project_user` role. That suggests that we could add `master` to that role... need to research that. I think we could do:

```
export PGUSER=master
psql -c "grant $project_user to master;" postgres
createdb --template=template0 --owner=$project_user $dbname
```

The question would be: Does `master` have enough privileges to grant itself membership in another role?

- Finally, you could create it as `master` when `master` is not a member of the `project_user` role. To do that, you'll need to create it as `master` and then modify the ownership and permissions:

```
export PGUSER=master
createdb --template=template0 $dbname
psql -c "revoke all on database $dbname from public;"
psql -c "grant all on database $dbname to master;"
psql -c "grant all on database $dbname to $project_user;"
```

If you need to enable extensions etc, do that now (see below). When done, then:

```
psql -c "alter database $dbname owner to $project_user;"
```

A superuser could create the database already owned by a specific user, but RDS's master user cannot.

32.7 PostGIS

To enable PostGIS, as the master user:

```
export PGUSER=master
psql -c "create extension postgis;"
psql -c "alter table spatial_ref_sys OWNER TO $project_user;"
```

where `$project_user` is the postgres user who will be using the database.

(Outside of RDS, only a superuser can use `create extension`; RDS has special handling for a whitelist of extensions.)

32.8 Hstore

Hstore is simpler, but you still have to use the master user:

```
export PGUSER=master
psql -c "create extension hstore;"
```

32.9 Grant read-only access to a database

Only let *readonly_user* do reads:

```
$ psql -c "GRANT CONNECT ON DATABASE $dbname TO $readonly_user;"
$ psql -c "GRANT SELECT ON ALL TABLES IN SCHEMA PUBLIC TO $readonly_user;" $dbname
```

32.10 Restore a dump to a new database

Create the database as above, including changing ownership to the project user, and enabling any needed extensions. Then as the project user:

```
export PGUSER=$project_user
pg_restore --no-owner --no-acl --dbname=$dbname file.dump
```

Note that you might get some errors during the restore if it tries to create extensions that already exist and that kind of thing, but those are harmless. It does mean you can't use `--one-transaction` or `--exit-on-error` for the restore though, because they abort on the first error.

32.11 Dump the database

This is pretty standard and can be done by the project user:

```
export PGUSER=$project_user
pg_dump --file=output.dump --format=custom $dbname
```

32.12 Drop database

When it comes time to drop a database, only master has the permission, but master can only drop databases it owns, so it takes two steps. Also, you can't drop the database you're connected to, so you need to connect to a different database for the dropdb. The postgres database is as good as any:

```
export PGUSER=master PGDATABASE=postgres
psql -c "alter database $dbname owner to master;"
psql -c "drop database if exists $dbname;"
```

(Outside of RDS, a superuser can drop any database. A superuser still has to be connected to some other database when doing it, though.)

32.13 Drop user

This is standard too. Just beware that you cannot drop a user if anything they own still exists, including things like permissions on databases.:

```
$ export PGUSER=master
$ dropuser $user
```

32.14 Postgres on RDS

- Add `django-extensions` to the requirements and `django_extensions` to the `INSTALLED_APPS` so we can use the `[sqldsn]`(<http://django-extensions.readthedocs.org/en/latest/sqldsn.html>) management command to get the exact Postgres settings we need to access the database from outside of Django. Here's how it works:

```
manage.py [--settings=xxxx] sqldsn
```


Contents:

33.1 Asyncio

33.1.1 What is it

`asyncio` is a library included in Python 3.5 that supports a programming model where sometimes, operations that would normally block the thread until some other event happened (like getting a response from a network connection) instead allow other code to run on that thread while waiting.

`asyncio` takes a very, very explicit approach to asynchronous programming: only code written in methods flagged as `async` can call any code in an asynchronous way.

Which creates a chicken/egg problem: your *async* methods can only be called by other *async* methods, so how do you call the first one?

The answer: you don't. What you have to do instead is turn over control of the thread to an event loop, after arranging for the loop to (sooner or later) invoke your *async* code.

Then once you start the loop running, it can invoke the *async* code.

33.1.2 What good is it

Note first that you can use threads to accomplish the same things as `asyncio` in most cases, with better performance. So what good is `asyncio`?

For one thing, it leads to more straightforward code than managing multiple threads, protecting data structures from concurrent access, etc. There's only one thread and no preemptive multitasking.

If you want to play with *async* programming in Python, `asyncio` looks easier to work with and understand than Twisted, but that's not a very practical reason.

More significantly, threads won't scale as well if you need to wait for many, many things at the same time - asyncio might be somewhat slower, but might be the only way that some tasks can be run at all. Each thread can take 50K of memory, while a coroutine might take only 3K.

33.1.3 Event loops

Async code can only run inside an *event loop*. The event loop is the driver code that manages the cooperative multi-tasking.

(I think) a typical pattern would be to get or create an event loop, set up some things to be run by it, then start the event loop running and have it run until the program is finished.

If it's useful for some reason, you can create multiple threads and run different event loops in each of them. For example, Django uses the main thread to wait for incoming requests, so we can't run an asyncio event loop there, but we can start a separate worker thread for our event loop.

33.1.4 Coroutines

coroutines

- Python distinguishes between a *coroutine function* and a *coroutine object*
- Write a coroutine function by putting `async` in front of the `def`.
- Only a coroutine function can use `await`, non-coroutine functions cannot.
- Calling a *coroutine function* does *not* execute it, but rather returns a *coroutine object*. (This is analogous to generator functions - calling them doesn't execute the function, it returns a generator object, which we then use later.)
- To execute a *coroutine object*, either:
 - use it in an expression with `await` in front of it, or
 - schedule it with `ensure_future()` or `create_task()`.

Example with `await`:

```
async def coro_function():
    return 2 + 2

coro = coro_function()
# not executed yet; coro is a coroutine, not 4

print(await coro)
# prints "4"
```

Example of scheduling it:

```
async def coro_function(hostname):
    conn = await .... connect async to hostname somehow...

coro = coro_function("example.com")
asyncio.ensure_future(coro)
```

Of course, usually you wouldn't split it onto two lines with a temp variable:

```
asyncio.ensure_future(coro_function("example.com"))
```


or:

```
asyncio.get_event_loop().create_task(coro_function("example.com"))
```

33.1.5 Futures

A *future* is an object that represents something uncompleted. It makes it easy for code in one place to indicate when the work is done, and optionally what the result was, and for code elsewhere that was interested in it to find out about it.

In other words, you can use future objects to manage synchronization more explicitly.

Create one on the fly by calling `loop.create_future()`:

```
future = loop.create_future()
```

Arrange for something to be called when the future becomes done:

```
future.add_done_callback(fn)
```

You can add lots of callbacks. They'll all be called (one at a time).

The callback receives the future object as an argument. Use `functools.partial` as usual if you want to pass other arguments.

When the future is done, mark it done and set its result:

```
future.set_result(value)
```

The callbacks can call `future.result()` to find out what the result was if they care.

33.1.6 Tasks

A Task is a way to arrange for a coroutine to be executed by an event loop, while also providing the caller a way to find out what the result was.

A task is automatically scheduled for execution when it is created.

There are two ways to do this, which seem equivalent as far as I can tell:

```
future = loop.create_task(coroutine)
future = asyncio.ensure_future(coroutine[, loop=loop])
```

Now you can add callbacks if you want:

```
future.add_done_callback(fn1)
```

Also, if the loop isn't already running and you just want to run the loop for this one thing, you can now:

```
loop.run_until_complete(future)
```

33.1.7 Awaitables

Coroutine *objects* and future *objects* are called *awaitables* - either can be used with `await`.

Note: You can only invoke an awaitable *once*; after that, it's completed, done, it runs no more.

33.1.8 Event loops

Creating/getting one

- To get the current thread's default event loop object, call `asyncio.get_event_loop()`
- `get_event_loop` will *not* create an event loop object unless you're on the main thread, and otherwise will raise an exception if the current thread doesn't have a default loop set.
- To create a new event loop: `new_event_loop()`
- To make a loop the default loop for the current thread: `set_event_loop(loop)`

So, to use an event loop in the main thread, you can just do:

```
loop = asyncio.get_event_loop()
# use loop....
```

But to run an event loop in another thread, you would do something like:

```
loop = asyncio.new_event_loop()
asyncio.set_event_loop(loop)
# use loop...
```

You don't have to set your loop as the thread's default, though, if you're willing to pass your loop object to all the APIs that otherwise use the default loop. But that's a pain.

Running a loop

If you want a long-running loop that keeps responding to events until it's told to stop, use `loop.run_forever()`.

If you want to compute some finite work using coroutines and then stop, use `loop.run_until_complete(<future or coroutine>)`.

Stopping a loop

Use `loop.stop()`.

Getting a loop to call a synchronous callable

By a *synchronous callable*, I mean a callable that is *not* an *awaitable* as described above.

This is more like Javascript's callback-style async programming than in the spirit of Python's coroutines, but sometimes you need it.

To call the callable as soon as possible, use `loop.call_soon(callback)`. If you want to pass args to the callable, use `functools.partial`:

```
loop.call_soon(functools.partial(callable, arg1, arg2))
```

To delay for *N* seconds before calling it, use `loop.call_later(delay, callable)`.

To schedule a callback from a different thread, the `AbstractEventLoop.call_soon_threadsafe()` method should be used. Example:

```
loop.call_soon_threadsafe(callback, *args)
```

Getting a loop to call an awaitable

Use `asyncio.ensure_future(awaitable, *, loop=None)`.

Or `loop.run_until_complete`, but as noted above, that just runs the loop as long as it takes to complete the awaitable.

If you're doing this from another thread, then you need to use a different method, `asyncio.run_coroutine_threadsafe(coro, loop)`:

```
future = asyncio.run_coroutine_threadsafe(coroutine, loop)
```

Running blocking code in another thread

If you need to call some blocking code from a coroutine, and don't want to block the whole thread, you can make it run in another thread using `coroutine AbstractEventLoop.run_in_executor(executor, func, *args)`:

```
fn = functools.partial(method, *args)
result = await loop.run_in_executor(None, fn)
```

Sleep

Calling `asyncio.sleep(seconds)` does not sleep; it returns a *coroutine object*. When you *execute* it by invoking it with `await` etc, it will complete after `<seconds>` seconds. So, mostly you'd do:

```
await asyncio.sleep(10) # pause 10 seconds
```

33.2 How to Fix your Python Code's Style

(A formatted version of this draft is at https://cheat.readthedocs.io/en/latest/python/fixing_style.html.)

Sometimes we inherit code that doesn't follow the style guidelines we prefer when we're writing new code. We could just run `flake8` on the whole codebase and fix everything before we continue, but that's not necessarily the best use of our time.

Another approach is to update the styling of files when we need to make other changes to them. To do that, it's helpful to be able to run a code style checker on just the files we're changing. I've written tools to do that for various source control systems and languages over the years. Here's the one I'm currently using for Python and `flake8`.

I call this script *flake*. I have a key in my IDE bound to run it and show the output so I can click on each line to go to the code that has the problem, which makes it pretty easy to fix things.

It can run in two modes. By default, it checks any files that have uncommitted changes. Or I can pass it the name of a git branch, and it checks all files that have changes compared to that branch. That works well when I'm working on a feature branch that is several commits downstream from *develop* and I want to be sure all the files I've changed while working on the feature are now styled properly.

The script is in Python, of course.

33.2.1 Work from the repository root

Since we're going to work with file paths output from git commands, it's simplest if we first make sure we're in the root directory of the repository.

```
#!/usr/bin/env python3
import os
import os.path
import subprocess

if not os.path.isdir('.git'):
    print("Working dir: %s" % os.getcwd())
    result = subprocess.run(['git', 'rev-parse', '--show-toplevel'],
        stdout=subprocess.PIPE)
    dir = result.stdout.rstrip(b'\n')
    os.chdir(dir)
    print("Changed to %s" % dir)
```

We use `git rev-parse --show-toplevel` to find out what the top directory in the repository working tree is, then change to it. But first we check for a `.git` directory, which tells us we don't need to change directories.

33.2.2 Find files changed from a branch

If a branch name is passed on the command line, we want to identify the Python files that have changed compared to that branch.

```
import sys
...
if len(sys.argv) > 1:
    # Run against files that are different from *branch_name*
    branch_name = sys.argv[1]
    cmd = ["git", "diff", "--name-status", branch_name, "--"]
    out = subprocess.check_output(cmd).decode('utf-8')
    changed = [
        # "M\tfilename"
        line[2:]
        for line in out.splitlines()
        if line.endswith(".py") and "migrations" not in line and line[0] != 'D'
    ]
```

We use `git diff --name-status <branch-name>` – to list the changes compared to the branch. We skip file deletions — that means we no longer have a file to check — and migrations, which never seem to quite be PEP-8 compliant and which I've decided aren't worth trying to fix. (You may decide differently, of course.)

33.2.3 Find files with uncommitted changes

Alternatively, we just look at the files that have uncommitted changes.

```
else:
    # See what files have uncommitted changes
    cmd = ["git", "status", "--porcelain", "--untracked=no"]
    out = subprocess.check_output(cmd).decode('utf-8')
    changed = []
    for line in out.splitlines():
        if "migrations" in line:
            # Auto-generated migrations are rarely PEP-8 compliant. It's a losing
            # battle to always fix them.
            continue
        if line.endswith('.py'):
```

(continues on next page)

(continued from previous page)

```

if '->' in line:
    # A file was renamed. Consider the new name changed.
    parts = line.split(' -> ')
    changed.append(parts[1])
elif line[0] == 'M' or line[1] != ' ':
    changed.append(line[3:])

```

Here we take advantage of *git porcelain* to ensure the output won't change from one git version to the next, and it's fairly easy to parse in a script. (Maybe I should investigate using *porcelain* with the other git commands in the script, but what I have now works well enough.)

33.2.4 Run flake8 on the changed files

Either way, *changed* now has a list of the files we want to run flake8 on.

```

cmd = ['flake8'] + changed
rc = subprocess.call(cmd)
if rc:
    print("Flake8 checking failed")
    sys.exit(rc)

```

Running *flake8* with *subprocess.call* this way sends the output to stdout so we can see it. *flake8* will exit with a non-zero status if there are problems; we print a message and also exit with a non-zero status.

33.2.5 Wrapping up

I might have once written a script like this in Shell or Perl, but Python turns out to work quite well once you get a handle on the *subprocess* module.

The resulting script is useful for me. I hope you'll find parts of it useful too, or at least see something you can steal for your own scripts.

33.3 Mock

- <http://www.voidspace.org.uk/python/mock/mock.html>
- <https://docs.python.org/3/library/unittest.mock.html>

33.3.1 Mock a function being called in another module

Function being called from another module (medley/photos/tasks/galleryimport.py):

```

from urllib2 import urlopen

```

our test.py:

```

with mock.patch('medley.photos.tasks.galleryimport.urlopen') as urlopen:
    urlopen.side_effect = ValueError("Deliberate exception for testing")

```

33.3.2 Mock a method on a class

Use `mock.patch('python.packagepath.ClassName.method_name:`

```
# medley/photos/models.py
class MedleyPhotoManager(...):
    def build_from_url(...):
```

Test code:

```
with mock.patch('medley.photos.models.MedleyPhotoManager.build_from_url') as build_
    from_url:
        build_from_url.return_value = None
```

33.3.3 Replace something with an existing object or literal

Use `mock.patch` and pass `new=`:

```
with mock.patch("medley.photos.tasks.galleryimport.cache", new=get_cache('locmem://
    ↳')):
    ...
with mock.patch(target='medley.photos.tasks.galleryimport.MAX_IMPORTS', new=2):
    ...
```

33.3.4 Mock an attribute on an object we have a reference to

Use `mock.patch.object(obj, 'attrname')`:

with `mock.patch.object(obj, 'attrname')` as `foo`: ...

33.3.5 Data on the mock

Attributes on mock objects:

```
obj.call_count    # number of times it was called
obj.called == obj.call_count > 0
obj.call_args_list # a list of (args,kwargs), one for each call
obj.call_args     # obj.call_args_list[-1] (args,kwargs from last call)
obj.return_value  # set to what it should return
obj.side_effect   # set to an exception class or instance that should be raised when
    ↳ its called
obj.assert_called() # doesn't work with autospec=True? just assert obj.called
obj.assert_called_with(*args, **kwargs) # last call was with (*args, **kwargs)
```

33.4 Python Packaging

The authoritative docs.

Example `setup.py`:

```
# Always prefer setuptools over distutils
from setuptools import setup, find_packages
# To use a consistent encoding
from codecs import open
from os import path

here = path.abspath(path.dirname(__file__))

setup(
    name='ursonos',
    version='0.0.1',
    packages=find_packages(),
    url='',
    license='',
    author='poirier',
    author_email='dan@poirier.us',
    description='Urwid application to control Sonos',
    install_requires=[
        'soco',
        'urwid'
    ]
)
```

To include non-Python files in the packaging, create a **MANIFEST.in** file. Example:

```
include path/to/*.conf
```

adds the files matching `path/to/*.conf`. Another:

```
recursive-include subdir/path *.txt *.rst
```

adds all files matching `*.txt` or `*.rst` that are anywhere under `subdir/path`. Finally:

```
prune examples/sample?/build
```

should be obvious.

33.4.1 Source package

To build a source package:

```
python setup.py sdist
```

33.4.2 Wheels

To build a Universal Wheel:

```
python setup.py bdist_wheel --universal
```

You can also permanently set the `--universal` flag in “`setup.cfg`” (e.g., see [sampleproject/setup.cfg](#))

```
[bdist_wheel]
universal=1
```

Only use the `--universal` setting, if:

1. Your project runs on Python 2 and 3 with no changes (i.e. it does not require 2to3).
2. Your project does not have any C extensions.

33.4.3 Upload

The docs recommend `twine`

```
twine upload dist/*
```

33.5 Pip

Based on [A Better Pip Workflow](#), from a [hackernews](#) comment by [blaze33](#):

To keep your file structure with commands and nicely separate your dependencies from your dependencies' dependencies:

```
pip freeze -r requirements-to-freeze.txt > requirements.txt
```

instead of just:

```
pip freeze > requirements.txt
```

Danger: beware of git urls being replaced by egg names in the process.

33.6 Pipenv

[Pipenv docs](#)

33.6.1 Converting from a requirements file

Just run “`pipenv install [-r requirementsfile]`” and it'll see that there's no Pipfile but a requirements file, and will generate a new Pipfile and `.lock` file for you. Then edit the Pipfile to clean it up.

33.6.2 Creating a requirements file

Do this:

```
pipenv lock --requirements >non-dev-requirements.txt
pipenv lock --requirements --dev >only-dev-requirements.txt
```

33.6.3 Keeping dev-only packages out of production

1. Add dev-only packages using `pipenv install --dev <packages>`
2. For development, install using `pipenv install --dev`
3. In production, leave off the `--dev`

33.7 Better Python dependency management with pip-tools

I recently looked into whether I could use [pip-tools](#) to improve my workflow around projects' Python dependencies. My conclusion was that pip-tools would help on some projects, but it wouldn't do everything I wanted, and I couldn't use it everywhere. (I tried pip-tools version 2.0.2 in August 2018. If there are newer versions, they might fix some of the things I ran into when trying pip-tools.)

33.7.1 My problems

What were the problems I wanted to find solutions for, that just pip wasn't handling? Software engineer Kenneth Reitz explains them pretty well [in his post](#), but I'll summarize here.

Let me start by briefly describing the environments I'm concerned with. First is my development environment, where I want to manage the dependencies. Second is the test environment, where I want to know exactly what packages and versions we test with, because then we come to the deployed environment, where I want to use exactly the same Python packages and versions as I've used in development and testing, to be sure no problems are introduced by an unexpected package upgrade.

The way we often handle that is to have a requirements file with every package and its version specified. We might start by installing the packages we know that we need, then saving the output of *pip freeze* to record all the dependencies that also got installed and their versions. Installing into an empty virtual environment using that requirements file gets us the same packages and versions.

But there are several problems with that approach.

First, we no longer know which packages in that file we originally wanted, and which were pulled in as dependencies. For example, maybe we needed `Celery`, but installing it pulled in a half-dozen other packages. Later we might decide we don't need Celery anymore and remove it from the requirements file, but we don't know which other packages we can also safely also remove.

Second, it gets very complicated if we want to upgrade some of the packages, for the same reasons.

Third, having to do a complete install of all the packages into an empty virtual environment can be slow, which is especially aggravating when we know little or nothing has changed, but that's the only way to be sure we have exactly what we want.

33.7.2 Requirements

To list my requirements more concisely:

- Distinguish direct dependencies and versions from incidental
- Freeze a set of exact packages and versions that we know work
- Have one command to efficiently update a virtual environment to have exactly the frozen packages at the frozen versions and no other packages
- Make it reasonably easy to update packages
- Work with both installing from PyPI, and installing from Git repositories
- Take advantage of pip's hash checking to give a little more confidence that packages haven't been modified
- Support multiple sets of dependencies (e.g. dev vs. prod, where prod is not necessarily a subset of dev)
- Perform reasonably well
- Be stable

That's a lot of requirements. It turned out that I could meet more of them with `pip-tools` than just `pip`, but not all of them, and not for all projects.

Here's what I tried, using `pip`, `virtualenv`, and `pip-tools`.

33.7.3 How to set it up

1. I put the top-level requirements in `requirements.in/*.txt`.

To manage multiple sets of dependencies, we can include “-r file.txt”, where “file.txt” is another file in `requirements.in`, as many times as we want. So we might have a `base.txt`, a `dev.txt` that starts with `-r base.txt` and then adds `django-debug-toolbar` etc, and a `deploy.txt` that starts with `-r base.txt` and then adds `unicorn`.

There's one annoyance that seems minor at this point, but turns out to be a bigger problem: `pip-tools` only supports URLs in these requirements files if they're marked editable with `-e`.

```
# base.txt
Django<2.0
-e git+https://github.com/cactus/django-scribbler@v0.8.0#egg=django-scribbler

# dev.txt
-r base.txt
django-debug-toolbar

# deploy.txt
-r base.txt
unicorn
```

2. Install `pip-tools` in the relevant virtual environment:

```
$ <venv>/bin/pip install pip-tools
```

3. Compile the requirements as follows:

```
$ <venv>/bin/pip-compile --output-file requirements/dev.txt requirements.in/dev.txt
```

This looks only at the requirements file(s) we tell it to look at, and *not* at what's currently installed in the virtual environment. So one unexpected benefit is that `pip-compile` is faster and simpler than installing everything and then running `pip freeze`.

The output is a new requirements file at `requirements/dev.txt`.

`pip-compile` nicely puts a comment at the top of the output file to tell developers exactly how the file was generated and how to make a newer version of it.

```
#
# This file is autogenerated by pip-compile
# To update, run:
#
#   pip-compile --output-file requirements/dev.txt requirements.in/dev.txt
#
-e git+https://github.com/cactus/django-scribbler@v0.8.0#egg=django-scribbler
django-debug-toolbar==1.9.1
django==1.11.15
pytz==2018.5
sqlparse==0.2.4          # via django-debug-toolbar
````
```

4. Be sure `requirements`, `requirements.in`, and their contents are in version control.

### 33.7.4 How to make the current virtual environment have the same packages and versions

To update your virtual environment to match your requirements file, ensure pip-tools is installed in the desired virtual environment, then:

```
$ <venv>/bin/pip-sync requirements/dev.txt
```

And that's all. There's no need to create a new empty virtual environment to make sure only the listed requirements end up installed. If everything is already as we want it, no packages need to be installed at all. Otherwise only the necessary changes are made. And if there's anything installed that's no longer mentioned in our requirements, it gets removed.

Except ...

pip-sync doesn't seem to know how to uninstall the packages that we installed using `-e <URL>`. I get errors like this:

```
Can't uninstall 'pkgname1'. No files were found to uninstall.
Can't uninstall 'pkgname2'. No files were found to uninstall.
```

I don't really know, then, whether pip-sync is keeping those packages up to date. Maybe before running pip-sync, I could just

```
rm -rf $VIRTUAL_ENV/src
```

to delete any packages that were installed with `-e`? But that's ugly and would be easy to forget, so I don't want to do that.

### 33.7.5 How to update versions

1. Edit `requirements.in/dev.txt` if needed.
2. Run pip-compile again, exactly as before:

```
$ <venv>/bin/pip-compile--output-file requirements/dev.txt requirements.in/dev.txt
```

3. Update the requirements files in version control.

### 33.7.6 Hash checking

I'd like to use hash checking, but I can't yet. pip-compile can generate hashes for packages we will install from PyPI, but not for ones we install with `-e <URL>`. Also, pip-sync doesn't check hashes. *pip install* will check hashes, but if there are any hashes, then it will fail unless *all* packages have hashes. So if we have any `-e <URL>` packages, we have to turn off hash generation or we won't be able to *pip install* with the compiled requirements file. We could still use pip-sync with the requirements file, but since pip-sync doesn't check hashes, there's not much point in having them, even if we don't have any `-e` packages.

### 33.7.7 What about pipenv?

*pipenv* promises to solve many of these same problems. Unfortunately, it imposes other constraints on my workflow that I don't want. It's also changing too fast at the moment to rely on in production.

Pipenv solves several of the requirements I listed above, but fails on these: It only supports two sets of requirements: base, and base plus dev, not arbitrary sets as I'd like. It can be very slow. It's not (yet?) stable: the interface and behavior is changing constantly, sometimes [multiple times in the same day](#).

It also introduces some new constraints on my workflow. Primarily, it wants to control where the virtual environment is in the filesystem. That both prevents me from putting my virtual environment where I'd like it to be, and prevents me from using different virtual environments with the same working tree.

### 33.7.8 Shortcomings

pip-tools still has some shortcomings, in addition to the problems with checking hashes I've already mentioned.

Most concerning are the errors from pip-sync when packages have previously been installed using `-e <URL>`. I feel this is an unresolved issue that needs to be fixed.

Also, I'd prefer not to have to use `-e` at all when installing from a URL.

This workflow is more complicated than the one we're used to, though no more complicated than we'd have with pipenv, I don't think.

The number and age of open issues in the pip-tools git repository worry me. True, it's orders of magnitude fewer than some projects, but it still suggests to me that pip-tools isn't as well maintained as I might like if I'm going to rely on it in production.

### 33.7.9 Conclusions

I don't feel that I can trust pip-tools when I need to install packages from Git URLs.

But many projects don't need to install packages from Git URLs, and for those, I think adding pip-tools to my workflow might be a win. I'm going to try it with some real projects and see how that goes for a while.

## 33.8 Six

Writing Python code that works the same in Python 2 & 3, using six as necessary.

| Python 2                                                                                   | Python 3                                                                                   | Both                                                                                            |
|--------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------|
| <code>from urlparse import urlparse</code>                                                 | <code>from urllib.parse import urlparse</code>                                             | <code>from six.moves.urllib.parse import urlparse</code>                                        |
| <code>dict.iteritems()</code>                                                              | <code>dict.items()</code>                                                                  | <code>six.iteritems(dict)</code>                                                                |
| <code>dict.iterkeys()</code>                                                               | <code>dict.keys()</code>                                                                   | <code>six.iterkeys(dict)</code>                                                                 |
| <code>dict.keys()</code>                                                                   | <code>list(dict.keys())</code>                                                             | <code>list(dict.keys())</code>                                                                  |
| <code>Exception.message</code>                                                             | <code>Exception.args[0]</code>                                                             | <code>Exception.args[0]</code>                                                                  |
| <code>isinstance(x, str)</code>                                                            | <code>isinstance(x, bytes)</code>                                                          | <code>isinstance(x, six.binary_type)</code>                                                     |
| <code>filter(condition, iterable)</code>                                                   | <code>[x for x in iterable if condition(x)]</code>                                         | <code>[x for x in iterable if condition(x)]</code>                                              |
| <code>map(function, iterable)</code>                                                       | <code>[function(x) for x in iterable]</code>                                               | <code>[function(x) for x in iterable]</code>                                                    |
| <code>self.assertItemsEqual(<br/>    random_order(items), random_order(more_items))</code> | <code>self.assertCountEqual(<br/>    random_order(items), random_order(more_items))</code> | <code>six.assertCountEqual( self,<br/>    random_order(items), random_order(other_items)</code> |
| <code>string.lowercase</code>                                                              | <code>string.ascii_lowercase</code>                                                        | <code>string.ascii_lowercase</code>                                                             |
| <code>class Foo(BaseClass):<br/>    __metaclass__ = MetaClass</code>                       | <code>class Foo(BaseClass,<br/>    metaclass=MetaClass):</code>                            | <code>@six.add_metaclass(MetaClass)<br/>class Foo(BaseClass)</code>                             |
| <code>str</code>                                                                           | <code>bytes</code>                                                                         | <code>six.binary_type</code>                                                                    |
| <code>unicode</code>                                                                       | <code>str</code>                                                                           | <code>six.text_type</code>                                                                      |
| <code>[int, long]</code>                                                                   | <code>[int]</code>                                                                         | <code>six.integer_types</code>                                                                  |
| <code>iterator.next()</code>                                                               | <code>next(iterator)</code>                                                                | <code>six.next(iterator)</code>                                                                 |
| <code>zip(a, b).__iter__()</code>                                                          | <code>zip(a, b)</code>                                                                     |                                                                                                 |
| <code>zip(a, b)</code>                                                                     | <code>list(zip(a, b))</code>                                                               | <code>list(zip(a, b))</code>                                                                    |

`six.BytesIO` This is a fake file object for binary data. In Python 2, it's an alias for `StringIO.StringIO`, but in Python 3, it's an alias for `io.BytesIO`.

`six.StringIO` This is an fake file object for textual data. It's an alias for `StringIO.StringIO` in Python 2 and `io.StringIO` in Python 3.

`six.moves.reduce` `reduce` `reduce()` `functools.reduce()`

`six.moves.range`

`range` `xrange()` `range`

`six.reraise(exc_type, exc_value, exc_traceback=None)` Reraise an exception, possibly with a different traceback. In the simple case, `reraise(*sys.exc_info())` with an active exception (in an except block) reraises the current exception with the last traceback. A different traceback can be specified with the `exc_traceback` parameter. Note that since the exception reraising is done within the `reraise()` function, Python will attach the call frame of `reraise()` to whatever traceback is raised.

2to3 does `raise E, V, T` to `raise E(V).with_traceback(T)`

## 33.9 Timezones in Python

### 33.9.1 Key points

- Install the pytz package to provide actual time zones. Python doesn't come with them.
- There are two kinds of datetime objects in Python, and you need to always know which you're working with:
  - naive - has no timezone info. (datetime.tzinfo is None)
  - aware - has timezone info (datetime.tzinfo is not None)
- There will always be some things you want to do with datetimes that are just inherently ambiguous. Get used to it.

### 33.9.2 Some use cases

Start by importing useful modules:

```
import datetime, time, pytz
```

#### Given a formatted time string with timezone, end up with a datetime object

Suppose we have RFC 2822 format:

```
s = "Tue, 3 July 2012 14:11:03 -0400"
```

It would be nice if `strptime` could just parse this and give you an aware datetime object, but no such luck:

```
>>> fmt = "%a, %d %B %Y %H:%M:%S %Z"
>>> datetime.datetime.strptime(s, fmt)
Traceback (most recent call last):
 File "<stdin>", line 1, in <module>
 File "/System/Library/Frameworks/Python.framework/Versions/2.7/lib/python2.7/_
↳strptime.py", line 325, in _strptime
 (data_string, format))
ValueError: time data 'Tue, 3 July 2012 14:11:03 -0400' does not match format '%a, %d
↳%B %Y %H:%M:%S %Z'
>>> fmt = "%a, %d %B %Y %H:%M:%S %z"
>>> datetime.datetime.strptime(s, fmt)
Traceback (most recent call last):
 File "<console>", line 1, in <module>
 File "/System/Library/Frameworks/Python.framework/Versions/2.6/lib/python2.6/_
↳strptime.py", line 317, in _strptime
 (bad_directive, format))
ValueError: 'z' is a bad directive in format '%a, %d %B %Y %H:%M:%S %z'
```

So, we have to parse it without the time zone:

```
>>> fmt = "%a, %d %B %Y %H:%M:%S"
>>> dt = datetime.datetime.strptime(s[:-6], fmt)
>>> dt
datetime.datetime(2012, 7, 3, 14, 11, 3)
```

That is assuming we know exactly how long the timezone string was, but we might not. Try again:

```
>>> last_space = s.rindex(' ')
>>> last_space
25
>>> datetime.datetime.strptime(s[:last_space], fmt)
```

Now, we need to figure out what that timezone string means. Pick it out:

```
>>> tzs = s[last_space+1:]
>>> tzs
'-0400'
```

We could have a timezone name or offset, but let's assume the offset for now. RFC 2282 says this is in the format `[+-]HHMM`:

```
>>> sign = 1
>>> if tzs.startswith("-"):
... sign = -1
... tzs = tzs[1:]
... elif tzs.startswith("+"):
... tzs = tzs[1:]
...
>>> tzs
'0400'
>>> sign
-1
```

Now compute the offset:

```
>>> minutes = int(tzs[0:2])*60 + int(tzs[2:4])
>>> minutes *= sign
>>> minutes
-240
```

Unfortunately, we can't just plug that offset into our datetime. To create an aware object, Python wants a `tzinfo` object that has more information about the timezone than just the offset from UTC at this particular moment.

So here's one of the problems - we don't KNOW what timezone this date/time is from, we only know the current offset from UTC.

So, the best we can do is to figure out the corresponding time in UTC, then create an aware object in UTC. We know this time is 240 minutes less than the corresponding UTC time, so:

```
>>> import time
>>> time_seconds = time.mktime(dt.timetuple())
>>> time_seconds -= 60*minutes
>>> utc_time = datetime.datetime.fromtimestamp(time_seconds, pytz.utc)
>>> utc_time
datetime.datetime(2012, 7, 3, 22, 11, 3, tzinfo=<UTC>)
```

And there we have it, an aware datetime object for that moment in time.

## 33.10 Tox

### 33.10.1 Command line

Some useful options:

- `tox -l` - list environments
- `tox -e envname[, envname...]` - run tests for just those environments

### 33.10.2 Configuration

Links:

- [Configuration](#)

### 33.10.3 Substitutions

E.g. `{toxindir}` expands to the directory where `tox.ini` is, and `{{env:ENV_NAME}}` expands to the value of `ENV_NAME` from the environment:

- [Substitutions](#)

### 33.10.4 Environment variables

Set env vars for the test:

```
[testenv]
setenv =
 VAR1 = value1
 VAR2 = value2
```

**Warning:** By default tox strips most values from the environment.

You can override with [passenv](#). I generally just put:

```
passenv = *
```

in every `tox.ini` file.

- [passing in environment variables](#)
- [setting environment variables](#)

## 33.11 Virtualenv

Using Python virtual environments

Assumes `virtualenv` and [virtualenvwrapper](#) is installed

pip install into site.USER\_BASE (kind of a special virtualenv):

```
PIP_REQUIRE_VIRTUALENV= pip install --user ...
```

<http://www.doughellmann.com/docs/virtualenvwrapper>:

```
sudo /usr/bin/easy_install pip
sudo /usr/local/bin/pip install virtualenvwrapper
```



Then add to .bashrc:

```
export WORKON_HOME=$HOME/.virtualenvs
source /usr/local/bin/virtualenvwrapper.sh
```

creates and activates a new env, //envname//:

```
mkvirtualenv //envname//
```

switch to //envname2//:

```
workon //envname2//
```

no longer working with a virtual env:

```
deactivate
```

List all of the environments:

```
lsvirtualenv
```

Show the details for a single virtualenv:

```
showvirtualenv
```

delete a virtual env (must deactivate first):

```
rmvirtualenv
```

## 33.12 XML in Python

### 33.12.1 Formatting an etree Element

Like this:

```
def format_element(element):
 """
 Given an etree Element object, return a string with the contents
 formatted as an XML file.
 """
 tree = ElementTree(element)
 bytes = StringIO()
 tree.write(bytes, encoding='UTF-8')
 return bytes.getvalue().decode('utf-8')
```

## 33.13 Most minimal logging

Python doesn't provide any logging handlers by default, resulting in not seeing anything but an error from the logging package itself. ... Add a handler to the root logger so we can see the actual errors.:

```
import logging
logging.getLogger('').addHandler(logging.StreamHandler())
```

## 33.14 Binary data to file-like object (readable)

```
f = io.BytesIO(binary_data) # Python 2.7 and 3
```

## 33.15 Join list items that are not None

Special case use of *filter*:

```
join(', ', filter(None, the_list))
```

## 33.16 Declare file encoding

Top of .py file:

```
-*- coding: UTF-8 -*-
vim:fileencoding=UTF-8
```

## 33.17 Quick 'n' easy static web server

- Change to the top level directory of the static files
- Run `python -m http.server`

## 33.18 Python snippet copy logging to stdout

add to top:

```
import logging

handler = logging.StreamHandler()
root_logger = logging.getLogger('')
root_logger.setLevel(logging.DEBUG)
root_logger.addHandler(handler)
```

## 33.19 Warnings

Hiding python warnings. e.g. Deprecations:

```
python -Wignore::DeprecationWarning ...
```

## 33.20 Pylint

Running it:

```
python /usr/bin/pylint --rcfile=.pylintrc -f colored aremind.apps.patients | less -r
```

Disabling a warning in code:

```
pylint: disable-msg=E1101
```



## 34.1 USB drive with Raspberry Pi

Here are detailed steps to get a Raspberry PI running entirely off a USB drive (after booting from its SD card).

Some sources:

- **Move root to USB drive:** <http://magnatecha.com/using-a-usb-drive-as-os-root-on-a-raspberry-pi/>  
[http://elinux.org/Transfer\\_system\\_disk\\_from\\_SD\\_card\\_to\\_hard\\_disk](http://elinux.org/Transfer_system_disk_from_SD_card_to_hard_disk)
- **UUIDs:** <http://liquidat.wordpress.com/2013/03/13/uuids-and-linux-everything-you-ever-need-to-know/>

Detailed log:

- Start with 4 GB SD card
- Copy 2013-07-26-wheezy-raspbian.img onto it:  

```
sudo dd if=path/2013-07-26-wheezy-raspbian.img of=/dev/mmcblk0 bs=1M
```
- Partition USB drive:
  - Partition 1: system root: 16 GB
  - Partition 2: Swap: 1 GB
  - Partition 3: /usr/local: rest of drive
- **Format the partitions too (ext4)**
  - `sudo mkfs.ext3 -q -m 0 -L ROOT /dev/sdb1`
  - `sudo mkswap -q -L SWAP120 /dev/sdb2`
  - `sudo mkfs.ext3 -q -m 0 -L LOCAL /dev/sdb3`
- Boot off 4 GB SD card
- ssh to its IP address (get from router or whatever):

```
ssh pi@[IP ADDRESS]
password: raspberry
```

- Copy current root partition to USB drive (see blog post mentioned above to make sure you're using the right partitions):

```
sudo dd if=/dev/mmcblk0p2 of=/dev/sda1 bs=4M
```

- Resize:

```
sudo e2fsck -f /dev/sda1
sudo resize2fs /dev/sda1
```

- See what UUID it got:

```
$ sudo blkid /dev/sda1
/dev/sda1: UUID="9c7e2035-df9b-490b-977b-d60f2170889d" TYPE="ext4"
```

- Mount:

```
sudo mount /dev/sda1 /mnt
```

- Edit config files and change current root partition to the new root UUID in fstab, and /dev/sda1 in cmdline.txt (cmdline.txt doesn't support UUID, darn)

- vi /mnt/etc/fstab:

```
UUID=9c7e2035-df9b-490b-977b-d60f2170889d / ext4 defaults,noatime,
↪async 0 1
```

- vi /boot/cmdline.txt ([http://elinux.org/RPi\\_cmdline.txt](http://elinux.org/RPi_cmdline.txt))

- Umount /mnt:

```
sudo umount /mnt
```

- Reboot and check things out

Swap:

- Format swap partition:

```
$ sudo mkswap -L swappart /dev/sda2
Setting up swapspace version 1, size = 1048572 KiB
LABEL=swappart, UUID=a471af01-938b-4ad0-8653-dafe211cdfba
```

- Make sure it'll work:

```
sudo swapon -U a471af01-938b-4ad0-8653-dafe211cdfba
free -h
```

- Edit /etc/fstab, add at end:

```
UUID=a471af01-938b-4ad0-8653-dafe211cdfba swap swap defaults 0 0
```

- Remove the default 100M swap file:

```
sudo apt-get purge dphys-swapfile
```

- reboot and check swap space again, should be 1 G (not 1.1 G)

Now move /usr/local to the USB drive:

- Format partition:

```
sudo mkfs.ext4 -L usr.local /dev/sda3
```

- Find out its UUID:

```
$ blkid /dev/sda3
/dev/sda3: LABEL="usr.local" UUID="3c6e0024-d0e4-412e-a4ab-35d7c9027070" TYPE=
↪ "ext4"
```

- Mount it temporarily on /mnt:

```
sudo mount UUID="3c6e0024-d0e4-412e-a4ab-35d7c9027070" /mnt
```

- Copy the current /usr/local over there:

```
(cd /usr/local;sudo tar cf - .) | (cd /mnt;sudo tar xf -)
```

- Umount:

```
sudo umount /mnt
```

- Remove files from /usr/local:

```
sudo rm -rf /usr/local/*
```

- Edit /etc/fstab to mount /dev/sda3 on /usr/local at boot:

```
UUID=3c6e0024-d0e4-412e-a4ab-35d7c9027070 /usr/local ext4 defaults,
↪noatime 0 1
```

- See if that works:

```
sudo mount -a
df -h
```

- reboot and make sure it works again





How to release

This is a checklist for releasing a new version of *something* if we're using the Git Flow <<http://nvie.com/posts/a-successful-git-branching-model/>> process for development.

### 35.1 Git branches

Since we use Git Flow, the two branches of concern at release time are:

- **master** - always has the most recently released code. Each release is tagged *X.X.X*.
- **develop** - contains the code under development for the next release

So technically, what makes a new release is merging *develop* to *master* and tagging it. Of course, we don't want to do that until we're ready.

We also use the git flow <<https://github.com/nvie/gitflow>> tool to help with the Git Flow branching model, especially for releases.

### 35.2 Release steps

Take these steps to release the new version:

- Make a fresh clone of the repo (to make sure we're working off the same code that's on github):

```
git clone git@github.com:org/project.git cd project
```

- Set up git flow in this new repo:

```
git checkout master
git flow init -d
```

- Start release branch using git flow:

```
git flow release start <VERSION>
```

e.g.

```
git flow release start '0.0.5'
```

Do **not** include `v` on the front of the version number - there's nothing wrong with it, we're just not using it for our version numbers here and want to be consistent.

- Run the tests locally. The tests must pass before proceeding. Fix any problems and commit the changes.
- Set `VERSION` in `project/__init__.py` to the same version, e.g. `VERSION = '0.0.5'`.
- Start a new section in `RELEASE_NOTES.rst` for the new release. Always put the new release section above the previous ones.
- Review `git log` and add major new features and incompatibilities to the release notes.
- Commit changes. Be sure to include the new version number in the commit message first line, e.g. "Bump version for 0.0.5".
- Push the release branch.
- Open a pull request from the release branch to the master branch.
- When pull request has been reviewed, use **git flow commands** to make the release:

```
git flow release finish '0.0.5'
```

You'll be prompted for a commit message for the merge to master. The default is fine (*Merge branch 'release/0.0.5'*).

You'll be prompted for a tag message. Make it "Tag for v0.0.5" or whatever the version is.

You'll be prompted for a commit message for the merge back to develop. The default is fine.

- Push the merged master and develop branch and tag to github:

```
git push origin master --tags git push origin develop --tags
```

- Verify that CI tests have passed for the pushed master
- Email the release announcement

REST reStructuredText notes (see also Sphinx notes)

SEE CHEATSHEET AT END

- [rst tutorial](#)
- [rst primer](#)
- [rst markup](#)
- [autodoc](#)

Example:

```
file1.rst

Contents:

.. toctree::
 :maxdepth: 2

 keys
 api

Go read :ref:`morestuff` for more. Read about python package :py:mod:`package.name`
and python function :py:func:`package.name.function`. The python class
:py:class:`package.name.ClassName` might also be useful, and its method
:py:meth:`package.name.ClassName.method`. Not to mention the attribute
:py:attr:`package.name.ClassName.attrname`. The code might throw the exception
:py:exc:`package.name.MyException`.

file2.rst

```

(continues on next page)

(continued from previous page)

```
.. _morestuff:

More stuff
=====

To show code::

 Indent this 4 spaces.
 You can indent more or less.
 And keep going.

 Even include blank lines

It ends after a blank line and returning to the original indentation.

You can also markup ``short inline code`` like this.
```

Automatically include the doc for a class like this:

```
.. _api:

.. autoclass:: healthvaultlib.healthvault.HealthVaultConn
 :members:
```

And document them:

```
class MyClassName(object):
 """
 Description. Can refer to :py:meth:`.method1` here or anywhere in the file.

 :param string parm1: His name
 :param long parm2: His number
 """

 def __init__(self, parm1, parm2):
 pass

 def method1(self, arg1, arg2):
 """
 Description

 :param unicode arg1: something
 :param object arg2: something else
 """
```

<http://techblog.ironfroggy.com/2012/06/how-to-use-sphinx-autodoc-on.html>

Various code blocks:

```
.. code-block:: bash|python|text
 :linenos:
 :emphasize-lines: 1,3-5

 # Hi
 # there
 # all
 # you
```

(continues on next page)

(continued from previous page)

```
coders
rejoice
```

```
1 # Hi
2 # there
3 # all
4 # you
5 # coders
6 # rejoice
```

You can include an ``HTML link`_` like this and the definition can go nearby or at the bottom of the page:

```
.. _HTML link: http://foo.bar.com/
```

Or you can just write ``HTML link <http://foo.bar.com>`_` all in one place.

<http://sphinx-doc.org/markup/inline.html#ref-role>

Link to a filename in this set of docs using `:doc:`Any text you want </path/to/page>`` or just `:doc:`path``.

Don't include the ".rst" on the end of the filename. Relative filenames work too. But it's better to use `:ref:`, see next.

You can define an anchor point, which Sphinx calls a label. Put this above a section header:

```
.. _my-label:

My Section

```

Now from somewhere else, you can write `:ref:`my-label`` and it'll be rendered as "My Section" and will link to the section. If you want some other text in the link, you can write `:ref:`any text <my-label>`` instead.

## 36.1 Cheatsheets

Copied from <http://docs.sphinxdocs.com/en/latest/cheatsheet.html> - thanks to Read The Docs.

BUGS:

- codeblock should be code-block

# reStructuredText & Sphinx Cheatsheet

## Styles

```
Bold text
Italic text
``Inline literal/code``
:sup:`super`\ Script
:sub:`sub`\ Script
```

## Bullet Lists

```
* Unordered item
* Unordered item

1. Nested ordered item
2. Nested ordered item

 a. Nested ordered item
* Unordered item
```

## Definition Lists

```
First term
 Definition of first term,
 it can span multiple lines.

Second term
 Definition of second term.

 Indent additional paragraphs.
```

## Literal Code Blocks

```
Here is a literal block::

 Its contents are indented.

::

 The :: marker is omitted here
```

The :: marker will insert a ":" in the output in the above example. To omit the ":", precede the :: marker with white space, or use the marker on a line of its own.

## Section Headings

```
=====
Level 1 Heading
=====

Level 2 Heading

```

```
Level 1 Heading
^^^^^^^^^^^^^^

Level 2 Heading
+++++
```

- Heading structure is determined only by occurrence order.
- Heading overline is optional.
- Under/overlines use the following characters:

Recommended: = - ` : . ' " ~ ^ \_ \* + #  
! \$ % & ( ) , / ; < > ? @ [ \ ] { | }

## Targets and Links

```
Anchor target .. _anchorbyref:
External target .. _External link by text:
Footnote target .. [1] A footnote
Citation target .. [cit1] A global citation

External links `External link <http://example.com>`_
`External link name`_ or `Example <External link name>`_

Internal links `Anchor link by text`_ or `Anchor <Anchor link by text>`_
`Anchor link by ref <anchorbyref>`_
:ref:`anchorbyref`

Footnote Reference a footnote [1]_
Citation or a global citation [cit1]_

Section link Section Heading

`Link <Section Heading>`_
```

## Tables

```
=====
Simple table Header 2 Header 3
=====
Column 1 Column 2 Column 3
Horizontal column span ...

... ..
=====

+-----+-----+-----+
| Grid | Header 2 | Header 3 |
| table | | |
+-----+-----+-----+
| Column 1 | Column 2 | Vertical |
+-----+-----+ column |
| Horizontal span | span |
+-----+-----+-----+
```

Simple tables and grid tables can be replaced with external CSV files, using the csv-table directive.

## Images and Figures

```
.. image:: image.png
 :height: 100px
 :width: 100px
 :align: bottom
 :target: target_
```

```
.. figure:: image.png
 :height: 100px
```

Figures are images with captions. They support all image options.

## Comments

```
.. This is a single line comment, comments can span multiple lines as well.
```

## Directives

reStructuredText directives consist of a **directive type**, **arguments**, and any number of **options**. Some directives expect a block of indented **content**.

```
.. type:: arguments
 :option: option value

 Directive content
```

```
.. container:: [container class]
.. csv-table:: [table title]
 :header: CSV data for headers
 :widths: number [, number]
 :file: filename
 :encoding: encoding
 :header-rows: number
 :delim: character, "tab", or "space"
 :quote: character
 :escape: character
.. rubric:: title
```

## Table of Contents

```
.. toctree::
 :maxdepth: number of title levels
 :glob:
 :hidden:
 :numbered:
 :caption: caption text
 :titlesonly:
```

[Include paths]

## Python Domain

Domain directives expect a domain **directive type**, the **construct name** as a directive argument, and optionally, **domain directive options**.

```
.. py:method:: MyClass.foo(args)
 :param args: Arguments
 :type args: list
```

Link to `py:meth:'MyClass.foo'`.

### Python Domain Directives

```
.. py:module:: module
.. py:class:: signature
.. py:function:: signature
.. py:data:: name
.. py:exception:: name
.. py:attribute:: name
.. py:method:: signature
.. py:staticmethod:: signature
.. py:classmethod:: signature
.. py:decorator:: signature
.. py:currentmodule:: module
```

### Python Domain References

```
:py:mod:`module`
:py:class:`signature`
:py:func:`signature`
:py:data:`name`
:py:exc:`name`
:py:attr:`name`
:py:meth:`signature`
```

### Python Domain Directive Options

```
:param [type] name: description
:type name: type
:raises class: description
:var name: description
:vartype name: type
:returns description:
:rtype type:
```

```
:py:const:`name`
```

There are a number of other Sphinx language domains for representing code constructs in reference documentation: rst, c, cpp, js, ruby, php, dotnet, scala, go, lisp, coffee, and others.

## Content Block Directives

```
.. topic:: [title]
.. sidebar:: [title]
 :subtitle: subtitle
.. admonition:: title
.. attention::
.. caution::
.. danger::
.. error::
.. hint::
.. important::
.. note::
.. tip::
.. warning::
.. seealso::
.. deprecated:: [version]
.. versionadded:: [version]
.. versionchanged:: [version]
.. math::
.. raw:: output format
```

```
.. topic:: Examples
```

A topic block

```
.. note::
 This is a note.
```

```
.. versionchanged:: 1.3
 Something changed.
```

```
.. math::
 a_1 = b_1 + c_1
```

```
.. raw:: html
```

```
HTML output
```

## Code Examples

```
.. highlight:: language
 :linenothreshold: number
.. codeblock:: [language]
 :linenos:
 :emphasize-lines: number [, number]
 :caption: caption text
 :name: block target name
```

[Code example, indented]

```
.. literalinclude:: filename
 :language: language
 :emphasize-lines: number [, number]
 :linenos:
 :encoding: encoding
 :diff: filename
 :dedent: number
```

Some of the language lexers supported by the code example language option: none, python, js, php, ruby, perl, c, cpp, csharp, go, scala, lisp, coffee, dart, julia, lua, html, css, sass, json, yaml, diff, sql, bash, and shell-session

## References

```
:role:`title <target>` Link to target with link text title
:role:`!title` Don't create link or reference
:role:`~module.Object` Use last element, Object, for link text
```

### Cross-reference roles

```
:any:`name`
:doc:`document name`
:download:`filename`
:envvar:`name`
:keyword:`python keyword`
:option:`CLI option`
```

### Other inline markup roles

```
:abbr:`long (abbreviation)`
:command:`name`
:kbd:`keystrokes`
:program:`program`
:pep:`number`
:rfc:`number`
```

Brought to you by **Read the Docs**.

We offer `private hosting` <http://readthedocs.com> and `open source hosting` <http://readthedocs.org> for *Sphinx* documentation projects.





<https://docs.saltstack.com/en/latest/topics/tutorials/pillar.html>

## 37.1 Fetching pillar data in templates

In the Jinja2 context, `pillar` is just a dictionary, so you can use the usual Python dictionary methods, e.g.:

```
{% for user, uid in pillar.get('users', {}).items() %}
{{user}}:
 user.present:
 - uid: {{uid}}
{% endfor %}
```

If you have nested data, it can be easier to use `salt['pillar.get']`, which accepts a single key parameter like `key1:key2`. E.g. if the pillar data looks like:

```
apache:
 user: fred
```

you could access the username as:

```
{{ salt['pillar.get']['apache:user'] }}
```

instead of:

```
{{ pillar['apache']['user'] }}
```

though that would work too.



## CHAPTER 38

---

### Socat

---

<http://stackoverflow.com/questions/2149564/redirecting-tcp-traffic-to-a-unix-domain-socket-under-linux>

Turns out socat can be used to achieve this:

```
socat TCP-LISTEN:1234,reuseaddr,fork UNIX-CLIENT:/tmp/foo
```

And with a bit of added security:

```
socat TCP-LISTEN:1234,bind=127.0.0.1,reuseaddr,fork,su=nobody,range=127.0.0.0/8 UNIX-
↪CLIENT:/tmp/foo
```



## CHAPTER 39

---

### Tmux

---

My tmux char: C-g (control-g)

session: a whole set of processes. Your tmux client is attached to one session. window:

Bindings:

```
Open prompt: :

Client:
 detach current: d
 detach any: D
 suspend: C-z

Sessions:
 rename $
 select s
 last L

Windows:
 Info: i
 Create: c
 Kill: &
 Switch to: '
 Choose: w
 Rename: ,
 Switch to 0-9: 0-9
 Previously selected: l
 Next: n
 Previous: p
 Next window with bell M-n

Pane:
 Kill: x
 Previous: ;
 swap with previous {
```

(continues on next page)

(continued from previous page)

```
swap with next }
```

Panes:

Split current pane into two:

Top and bottom: -

Left and right: |

## CHAPTER 40

---

### Travis CI

---

Magic syntax to run tox tests with multiple Python versions (as of Nov 1, 2017 anyway):

```
Use travis container-based build system for speed
sudo: false

Ubuntu trusty (14.04) - latest that Travis offers
dist: trusty

Make sure all the python versions we need are pre-installed
(apt-get is not available in the container-based build system)
addons:
 apt:
 sources:
 - deadsnakes
 packages:
 - python2.7
 - python3.4
 - python3.5
 - python3.6

language: python

The version of Python that'll be used to invoke tox. Has no effect
on what version of Python tox uses to run each set of tests.
python:
 - "3.5"

Test a sampling of combinations
env:
 - TOXENV=py27-1.7
 - TOXENV=py27-1.8
 - TOXENV=py27-1.10
 - TOXENV=py27-1.11
 - TOXENV=py34-1.7
```

(continues on next page)

(continued from previous page)

```
- TOXENV=py34-1.8
- TOXENV=py34-1.11
- TOXENV=py35-1.9
- TOXENV=py35-1.10
- TOXENV=py35-1.11
- TOXENV=py36-1.8
- TOXENV=py36-1.10
- TOXENV=py36-1.11

install:
 - pip install tox

script:
 - tox
#
#matrix:
allow_failures:
- env: TOXENV=py27-trunk,py33-trunk
```



Video tools to use on Linux

### 41.1 avidemux

GUI tool that's good for chopping up longer videos into shorter pieces without re-encoding when you need to watch the video and scan back and forth to find the places you want to split.

### 41.2 makemkv

CLI good for ripping DVDs and Blu-Rays to mkv.

(Has GUI too.)

### 41.3 mkvmerge

CLI only.

Join mkv streams end-to-end without re-encoding. (Inverse of avidemux.)

Split mkv streams at timestamps or chapters without re-encoding.



### 42.1 Startup

Looking around on Ubuntu

/etc/X11 is a good place to start

Xsession is invoked by the display manager.

See ‘man 5 Xsession’

Xsession redirects stdout & stderr to \$HOME/.xsession-errors almost first thing

Xsession.d/\* has a bunch of scripts that are sourced in order by Xsession:

20x11-common\_process-args: sets STARTUP if arg passed to Xsession 40x11-common\_xsessionrc: will source \$HOME/.xsessionrc if it is readable 50x11-common\_determine-startup: will set STARTUP to ~/.xsession or ~/.Xsession but ONLY if the display manager has not already set a startup command by passing it to Xsession - or so the comments say. . .

Also look at /etc/kde4/kdm:

kdmrc: LOTS of useful config stuff - read the comments

Which directories kdm looks in for different desktop sessions - It might just use the first of these where it finds anything?????

/usr/share/xsessions /var/lib/menu-xdg/xsessions (no such dir) /etc/kde4/kdm/sessions (empty)  
/usr/share/kde4/apps/kdm/sessions (a zillion files, I suspect it doesn't look here since others come first)



## CHAPTER 43

---

### Virtualbox

---

List running VMs:

```
VBoxManage list runningvms
```

Stop a running VM:

```
VBoxManage controlvm <UUID|VMNAME> pause|resume|reset|poweroff etc.
```

Delete a VM:

```
VboxManage unregistervm <UUID|VMNAME> [--delete]
```



## 44.1 Reactivity

The following applies to the store's state, anything in a component's *data*, and other things that get pulled into the reactivity system.

When an object is added to Vue's reactivity system, Vue replaces all its properties with getters and setters under the covers, so that if you fetch the value of a property, or assign a new value to it, Vue is aware and can react. (<https://vuejs.org/v2/guide/reactivity.html>)

However, for technical reasons, Vue cannot detect when a property is added to or removed from an object. (<https://vuejs.org/v2/guide/reactivity.html#Change-Detection-Caveats>)

**The implications are:**

- When updating the store, it's fine to assign a new value to a property of the *state*.
- When updating component data, it's fine to assign a new value to a property of the component.
- Don't try to use `Object.assign` or equivalent to update properties of objects in-place in the store???? It doesn't seem to work.

## 44.2 Vuex (the store)

### 44.2.1 Getters

Getters doc

Getters provide computed values based on the state. Their results are cached until the state they depend on changes.

Getters are accessed as *properties* not *methods*.

They are passed as a second arg an object with all the store's getters, in case they want to use them.

```
const store = new Vuex.Store({
 ...
 getters: {
 totalCost: (state, othergetters) => {
 return some_computation_on_state
 }
 }

 // component...

 computed: {
 the_total_cost () {
 return store.getters.totalCost // No parens, not called like a method
 }
 }
})
```

### 44.2.2 Mutations

[Mutations doc](#)

Mutations *must be synchronous*.

They cannot be called. They must be invoked using `commit`.

They receive a state and optional arguments, and can change the state.

When the state changes, other Vue components observing the state will update automatically.

Any value returned by a mutation is *not* passed back to the caller of `commit`.

### 44.2.3 Actions

[Actions doc](#)

Actions can contain asynchronous code. They receive a `context` object that has methods like `commit` and properties like `state` and `getters`.

Actions cannot be called. They must be invoked using `dispatch`.

Any value returned by an action is passed back to the caller of `dispatch`, by way of resolving the promise that `dispatch` returns to that value.

Actions always return Promises (via `dispatch`).

Example:

```
const store = new Vuex.Store({
 state: {
 count: 0
 },
 mutations: {
 increment (state) {
 state.count++
 }
 },
 actions: {
 increment (context) {
 context.commit('increment')
 },
 },
})
```

(continues on next page)



(continued from previous page)

```
checkout ({ commit, state }, products) {
 // save the items currently in the cart
 const savedCartItems = [...state.cart.added]
 // send out checkout request, and optimistically
 // clear the cart
 commit(types.CHECKOUT_REQUEST)
 // the shop API accepts a success callback and a failure callback
 shop.buyProducts(
 products,
 // handle success
 () => commit(types.CHECKOUT_SUCCESS),
 // handle failure
 () => commit(types.CHECKOUT_FAILURE, savedCartItems)
)
},
async actionA ({ commit }) {
 commit('gotData', await getData())
},
async actionB ({ dispatch, commit }) {
 await dispatch('actionA') // wait for `actionA` to finish
 commit('gotOtherData', await getOtherData())
}
}
})
```



## 45.1 YAML syntax

A value:

```
value
```

A value named “foo”:

```
foo: value
```

A list:

```
- 1
- 2
- 'a string'
```

A list named “bar”:

```
bar:
 - 1
 - 2
 - 'a string'
```

Alternate list syntax (“Flow style”):

```
bar: [1, 2, 'a string']
```

A dictionary:

```
key1: val1
key2: val2
key3: val3
```

A dictionary named “joe”:

```
joe:
 key1: val1
 key2: val2
 key3: val3
```

Dictionary in flow style:

```
joe: {key1: val1, key2: val2, key3: val3}
```

A list of dictionaries:

```
children:
 - name: Jimmy Smith
 age: 15
 - name: Jenny Smith
 age: 12
```

## CHAPTER 46

---

### Indices and tables

---

- `genindex`
- `modindex`
- `search`



## R

`redirect()` (built-in function), [98](#)